

# 1-train\_\_test\_\_split

October 20, 2024

## 1 Train Test Split

The train-test split is a critical concept in machine learning for evaluating the performance of a model. It involves dividing the dataset into two distinct parts: the training set and the test set. This method ensures that we can assess how well the model generalizes to unseen data, which is the key to making accurate predictions in real-world applications.

**Purpose of Train-Test Split:** The main goal is to evaluate the model's ability to generalize to new, unseen data by testing it on a separate test set that the model has never encountered during training. If a model performs well on both the training data and the test data, it's likely to perform well on new, real-world data.

### Key Terminology:

1. **Training Set:** The subset of the dataset used to train the model. The model learns patterns, relationships, and trends from this data.
2. **Test Set:** The subset used to evaluate the model after training. It provides an unbiased estimate of the model's performance on unseen data.

### Why Train-Test Split is Important:

- **Avoid Overfitting:** Without a train-test split, the model might memorize the training data (overfitting) rather than learning generalizable patterns. This would result in poor performance on new data.
- **Unbiased Evaluation:** It provides an unbiased measure of how well the model generalizes beyond the training data.
- **Generalization Performance:** It allows us to assess the model's generalization error, i.e., how it will perform on future data.

### Train-Test Split Ratio:

- **80:20 Split:** Commonly used when the dataset is sufficiently large. This balance ensures that there is enough data for the model to learn (80%) while still keeping a good portion (20%) for evaluating the model's performance.
- **70:30 Split:** Used when you want more data for testing, but still have enough data for training. This might be more appropriate when the dataset is relatively smaller, as having more test data can give a clearer picture of generalization.

- 90:10 Split: Used when the dataset is very large. A small test set is enough to provide a good estimate of model performance since the model has a huge amount of data to train on.

```
[23]: import pandas as pd
      from sklearn.model_selection import train_test_split
```

```
[24]: data = pd.read_csv('500hits.csv', encoding='latin-1')
      data.head()
```

```
[24]:
```

	PLAYER	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	\
0	Ty Cobb	24	3035	11434	2246	4189	724	295	117	726	1249	
1	Stan Musial	22	3026	10972	1949	3630	725	177	475	1951	1599	
2	Tris Speaker	22	2789	10195	1882	3514	792	222	117	724	1381	
3	Derek Jeter	20	2747	11195	1923	3465	544	66	260	1311	1082	
4	Honus Wagner	21	2792	10430	1736	3430	640	252	101	0	963	

  

	SO	SB	CS	BA	HOF
0	357	892	178	0.366	1
1	696	78	31	0.331	1
2	220	432	129	0.345	1
3	1840	358	97	0.310	1
4	327	722	15	0.329	1

```
[25]: X = data.drop(columns=['PLAYER', 'HOF'])
      y = data['HOF']

      X.shape
      y.shape
```

```
[25]: (465,)
```

```
[26]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=11,
      ↪test_size=0.2)
      print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(372, 14)
```

```
(93, 14)
```

```
(372,)
```

```
(93,)
```

```
[27]: print(X_train.describe().round(3))
```

	YRS	G	AB	R	H	2B	3B	\
count	372.000	372.000	372.000	372.000	372.000	372.000	372.000	
mean	17.011	2046.522	7526.078	1154.126	2177.995	382.505	78.094	
std	2.662	351.233	1302.406	291.308	426.615	97.173	48.798	

min	11.000	1331.000	4981.000	651.000	1660.000	177.000	3.000
25%	15.000	1797.500	6507.500	936.000	1838.000	312.000	41.000
50%	17.000	1992.000	7237.000	1099.000	2080.500	367.000	67.000
75%	19.000	2245.500	8198.250	1305.000	2383.750	436.250	108.000
max	26.000	3308.000	12364.000	2295.000	4189.000	792.000	309.000

	HR	RBI	BB	SO	SB	CS	BA
count	372.000	372.000	372.000	372.000	372.000	372.000	372.000
mean	202.642	901.073	780.105	850.323	196.927	58.987	0.289
std	141.726	484.370	327.453	472.918	185.586	49.322	0.021
min	9.000	0.000	239.000	0.000	7.000	0.000	0.246
25%	79.750	645.000	536.500	448.000	64.500	23.000	0.274
50%	185.500	977.500	719.000	844.000	141.000	52.000	0.288
75%	293.250	1218.500	961.250	1234.250	285.500	84.000	0.300
max	755.000	2297.000	2190.000	1936.000	1406.000	335.000	0.366

```
[28]: print(X_test.describe().round(3))
```

	YRS	G	AB	R	H	2B	3B \
count	93.000	93.000	93.000	93.000	93.000	93.000	93.000
mean	17.204	2057.409	7452.968	1135.065	2139.258	374.742	80.398
std	3.154	368.580	1265.371	283.877	415.165	93.929	51.792
min	11.000	1399.000	5472.000	601.000	1660.000	206.000	14.000
25%	15.000	1820.000	6622.000	935.000	1818.000	310.000	45.000
50%	17.000	1997.000	7359.000	1108.000	2054.000	361.000	68.000
75%	19.000	2282.000	8096.000	1283.000	2256.000	432.000	99.000
max	25.000	2850.000	10876.000	1859.000	3430.000	668.000	252.000

	HR	RBI	BB	SO	SB	CS	BA
count	93.000	93.000	93.000	93.000	93.000	93.000	93.000
mean	194.677	867.011	797.387	836.065	191.817	54.473	0.287
std	151.600	495.127	328.755	552.309	166.926	42.508	0.022
min	15.000	0.000	266.000	15.000	8.000	0.000	0.248
25%	78.000	618.000	527.000	359.000	61.000	15.000	0.272
50%	151.000	926.000	750.000	745.000	137.000	50.000	0.285
75%	291.000	1138.000	937.000	1179.000	271.000	83.000	0.299
max	612.000	1922.000	1747.000	2597.000	744.000	173.000	0.340