

# 12-Logistical\_Regression

October 20, 2024

## 1 Logistic Regression

**Logistic Regression** is a supervised learning algorithm used for classification tasks, not regression, despite its name. It estimates the probability that a given input point belongs to a certain class by applying a logistic (sigmoid) function to a linear combination of input features. The output is a probability between 0 and 1, making it ideal for binary classification problems (two possible outcomes, like yes/no, spam/not spam).

**Logistic regression** predicts a binary outcome based on one or more predictor variables (features) and is one of the foundational techniques in machine learning and statistics.

**When to Use Logistic Regression?** Logistic Regression is commonly used when:

- **The dependent variable is binary:** For example, a classification problem where the output is either “1” or “0” (e.g., true/false, success/failure).
- **You need probabilistic interpretations:** Logistic regression outputs probabilities that a data point belongs to a certain class, which makes it useful in areas where you need both the prediction and the confidence of that prediction.
- **When interpretability is important:** It provides clear and interpretable coefficients, which can explain the impact of each feature on the prediction.
- **The relationship between features and the target is linear in the log-odds:** While logistic regression assumes a linear relationship between the features and the log-odds of the outcome, it can still handle non-linear relationships by transforming the input features (e.g., polynomial or interaction terms).

**Typical use cases include:**

- **Medical Diagnosis:** Predicting whether a patient has a disease or not.
- **Email Classification:** Determining whether an email is spam or not.
- **Customer Churn:** Predicting whether a customer will leave a service or remain.
- **Credit Scoring:** Estimating the probability of a customer defaulting on a loan.

**Who Should Use Logistic Regression?** Logistic Regression is ideal for:

- **Data scientists and machine learning practitioners** who need a simple, interpretable model for binary classification tasks.

- **Business analysts and researchers** looking for an explainable model where they can interpret the effect of each feature on the outcome.
- **Medical professionals** making decisions based on classification problems like disease detection.
- **Beginner machine learning students** since logistic regression is one of the simplest and easiest algorithms to learn and apply.

#### **Advantages of Logistic Regression:**

- **Simplicity and Interpretability:** It is easy to implement and interpret, making it a popular choice for binary classification problems.
- **Probabilistic Output:** Logistic regression provides probabilities for class membership, which can be useful for understanding the confidence of predictions.
- **Fast and Efficient:** It performs well on relatively small datasets and is computationally efficient.
- **No Feature Scaling Required:** While performance can be improved by scaling features, logistic regression does not require it.
- **Linear Decision Boundary:** The decision boundary is linear, which makes it suitable for linearly separable datasets.

#### **Disadvantages of Logistic Regression:**

- **Limited to Linear Boundaries:** Logistic regression assumes a linear relationship between the features and the log-odds, making it unsuitable for complex or highly non-linear problems.
- **Assumes Independence of Features:** Like Naive Bayes, logistic regression assumes that the features are independent of each other, which may not hold true in many cases.
- **Cannot Handle Multiclass Problems Directly:** Logistic regression is a binary classifier by default, but techniques like One-vs-Rest or Softmax Regression can extend it to multiclass problems.
- **Sensitive to Outliers:** Logistic regression can be sensitive to outliers, which can skew the decision boundary.

#### **Real-World Applications of Logistic Regression:**

1. **Credit Scoring:** Predicting the probability of a customer defaulting on a loan.
2. **Medical Diagnosis:** Predicting the likelihood of a disease based on symptoms.
3. **Customer Retention:** Estimating the probability that a customer will churn or stay with the company.
4. **Marketing Campaigns:** Predicting whether a customer will respond to an advertisement or not.
5. **Email Classification:** Classifying emails as spam or non-spam.

```
[1]: import pandas as pd

d = {
    'miles_per_week': [
        ↪[37,39,46,51,88,17,18,20,21,22,23,24,25,27,28,29,30,31,32,33,34,38,40,42,57,68,35,36,41,43,
        'completed_50m_ultra': [
        ↪['no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no',
        ]
    }

df = pd.DataFrame(data=d)
df
```

```
[1]:      miles_per_week  completed_50m_ultra
0              37              no
1              39              no
2              46              no
3              51              no
4              88              no
..            ...              ...
96             67              yes
97             74              yes
98             79              yes
99             90              yes
100            112              yes
```

[101 rows x 2 columns]

```
[2]: from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import OrdinalEncoder
from sklearn.linear_model import LogisticRegression
from matplotlib import pyplot as plt

import seaborn as sns

finished_race = ['no', 'yes']
enc = OrdinalEncoder(categories=[finished_race])
df['completed_50m_ultra'] = enc.fit_transform(df[['completed_50m_ultra']])

df
```

```
[2]:      miles_per_week  completed_50m_ultra
0              37              0.0
1              39              0.0
2              46              0.0
3              51              0.0
4              88              0.0
```

```


..          ...          ...
96          67          1.0
97          74          1.0
98          79          1.0
99          90          1.0
100         112          1.0

```

```
[101 rows x 2 columns]
```


```
[3]: plt.scatter(df['miles_per_week'], df['completed_50m_ultra'])
```

```
[3]: <matplotlib.collections.PathCollection at 0x1f0ce21e410>
```



```
[4]: sns.countplot(x='completed_50m_ultra', data=df)
```

```
[4]: <Axes: xlabel='completed_50m_ultra', ylabel='count'>
```



```
[5]: X = df.iloc[:,0:1]
y = df.iloc[:,1]

X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.8,
↳random_state=11)

model = LogisticRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
model.score(X_test, y_test)
```

```
[5]: 0.9047619047619048
```

```
[6]: print(confusion_matrix(y_test, y_pred))
```

```
[[ 5  1]
 [ 1 14]]
```

```
[7]: print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support
```

0.0	0.83	0.83	0.83	6
1.0	0.93	0.93	0.93	15
accuracy			0.90	21
macro avg	0.88	0.88	0.88	21
weighted avg	0.90	0.90	0.90	21