

27-Random_Forest_Regressor

October 20, 2024

1 Random Forest Regressor

A **Random Forest Regressor** is an ensemble learning method that combines the predictions of multiple decision trees to improve predictive performance for regression tasks. Unlike a single decision tree, which might overfit or have high variance, a random forest aggregates the outputs of many trees to make the final prediction, thus reducing overfitting and variance while improving generalization.

It works by constructing several decision trees during training and averaging their predictions. This process of averaging multiple trees helps to smooth out the predictions and gives more robust results.

1.0.1 Why is Random Forest Regressor Important?

- **Reduces Overfitting:** Individual decision trees tend to overfit, especially when the trees are deep. By averaging the predictions of many trees, Random Forest mitigates overfitting, making the model more generalizable to unseen data.
- **Handles Non-linear Relationships:** Decision trees naturally handle non-linear relationships. A random forest regressor benefits from this, enabling it to capture complex patterns in the data.
- **Robustness to Noise and Outliers:** Because Random Forest averages the predictions of many trees, it is less sensitive to outliers and noise than individual decision trees.
- **Feature Importance:** Random Forest provides feature importance scores, which can be useful in understanding which features have the most predictive power.

1.0.2 How does Random Forest Regressor work?

Random Forest Regressor is based on an ensemble of decision trees. Here's a step-by-step explanation of how it works:

1. Bagging (Bootstrap Aggregating): The first key idea behind Random Forest is bagging. Random subsets (bootstrapped samples) of the original training data are created. Each decision tree in the forest is trained on a different random sample of the data. This technique helps in reducing the variance and overfitting that can occur with individual trees.

2. Random Feature Selection:** During the construction of each tree, Random Forest also selects a random subset of features to consider at each split. This random selection of features ensures that the trees in the forest are more diverse, further reducing correlation between the trees.

3. Tree Construction: Each tree is built independently by splitting the data at different nodes based on feature values (as in any decision tree). The process is recursive, and the tree grows until a stopping criterion is met (e.g., a minimum number of samples per leaf or a maximum tree depth).

4. Prediction: After all the trees are trained, predictions are made by each tree, and the final prediction for regression is the average of the predictions from all the trees.

1.0.3 When should you use Random Forest Regressor?

- **Non-linear and complex datasets:** Random Forest works well when your dataset has complex, non-linear relationships between features and target values.
- **High-dimensional data:** It handles high-dimensional datasets with many features and can reduce the risk of overfitting by selecting random subsets of features.
- **When interpretability is not the primary goal:** Random Forest is less interpretable than simpler models like linear regression, so it's best used when predictive accuracy is more important than understanding the model's inner workings.
- **When there are many categorical and continuous variables:** It can handle mixed-type data without much preprocessing.

1.0.4 Who uses Random Forest Regressor?

- **Data Scientists:** Random Forest is widely used in industries like finance, healthcare, and marketing to solve regression problems where complex patterns need to be captured.
- **Business Analysts:** It is often used to predict continuous outcomes such as sales forecasts, customer retention rates, and other business metrics.
- **Researchers:** In research fields like genomics or environmental science, Random Forest is applied to make predictions based on numerous features, and it's useful for feature selection.

1.0.5 Key Points to Remember:

- **Ensemble of Decision Trees:** Random Forest is a collection of decision trees, each trained on a different bootstrap sample of the data.
- **Random Feature Selection:** At each split, a random subset of features is chosen, making the trees less correlated with one another.
- **Averaging Predictions:** For regression tasks, the final prediction is the average of the predictions made by all the trees.
- **Feature Importance:** Random Forest naturally provides feature importance scores, which can help identify the most relevant features.

1.0.6 Advantages of Random Forest Regressor:

- **Reduced Overfitting:** By averaging multiple trees, Random Forest smooths out the predictions and reduces overfitting compared to a single decision tree.
- **Handles Missing Data:** Random Forest can handle missing values relatively well by filling in missing data based on correlations.

- **Robust to Outliers:** Due to the averaging process, Random Forest is less sensitive to outliers in the data.
- **No need for feature scaling:** Unlike models like SVM or linear regression, Random Forest doesn't require feature scaling (e.g., normalization or standardization).

```
[1]: import pandas as pd
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

healthexp = sns.load_dataset('healthexp')
healthexp
```

```
[1]:
```

	Year	Country	Spending_USD	Life_Expectancy
0	1970	Germany	252.311	70.6
1	1970	France	192.143	72.2
2	1970	Great Britain	123.993	71.9
3	1970	Japan	150.437	72.0
4	1970	USA	326.961	70.9
..
269	2020	Germany	6938.983	81.1
270	2020	France	5468.418	82.3
271	2020	Great Britain	5018.700	80.4
272	2020	Japan	4665.641	84.7
273	2020	USA	11859.179	77.0

[274 rows x 4 columns]

```
[2]: healthexp = pd.get_dummies(healthexp)
healthexp
```

```
[2]:
```

	Year	Spending_USD	Life_Expectancy	Country_Canada	Country_France	\
0	1970	252.311	70.6	False	False	
1	1970	192.143	72.2	False	True	
2	1970	123.993	71.9	False	False	
3	1970	150.437	72.0	False	False	
4	1970	326.961	70.9	False	False	
..	
269	2020	6938.983	81.1	False	False	
270	2020	5468.418	82.3	False	True	
271	2020	5018.700	80.4	False	False	
272	2020	4665.641	84.7	False	False	
273	2020	11859.179	77.0	False	False	

	Country_Germany	Country_Great Britain	Country_Japan	Country_USA
0	True	False	False	False
1	False	False	False	False
2	False	True	False	False
3	False	False	True	False
4	False	False	False	True
..
269	True	False	False	False
270	False	False	False	False
271	False	True	False	False
272	False	False	True	False
273	False	False	False	True

[274 rows x 9 columns]

```
[3]: X = healthexp.drop(['Life_Expectancy'], axis=1)
y = healthexp['Life_Expectancy']

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
↳random_state=19)

rfr = RandomForestRegressor(random_state=13)
rfr.fit(X_train, y_train)
```

```
[3]: RandomForestRegressor(random_state=13)
```

```
[4]: y_pred = rfr.predict(X_test)
y_pred
```

```
[4]: array([79.705, 81.089, 74.833, 79.357, 71.152, 79.186, 82.417, 81.7 ,
78.368, 71.48 , 76.446, 75.79 , 82.694, 78.601, 76.848, 77.614,
71.838, 78.91 , 72.414, 82.033, 71.816, 78.415, 79.937, 78.513,
78.265, 72.639, 75.759, 81.337, 81.408, 82.269, 74.003, 74.395,
78.677, 74.98 , 71.777, 74.347, 74.581, 76.581, 81.142, 80.895,
74.979, 81.68 , 77.489, 78.219, 80.888, 83.572, 78.392, 73.976,
77.236, 80.114, 74.728, 76.782, 78.031, 78.778, 77.519])
```

```
[5]: mean_absolute_error(y_test, y_pred)
```

```
[5]: 0.25916363636361917
```

```
[6]: root_mean_squared_error(y_test, y_pred)
```

```
[6]: 0.31970520512155615
```

```
[7]: r2_score(y_test, y_pred)
```

```
[7]: 0.9910457602615238
```

```
[8]: param_grid = {
      'n_estimators': [100, 200, 300],
      'max_depth': [10,20,30],
      'min_samples_split': [2,5,10],
      'min_samples_leaf': [1,2,4]
    }

    rfr_cv = GridSearchCV(rfr, param_grid, cv=3, n_jobs=-1)
    rfr_cv.fit(X_train, y_train)
```

```
[8]: GridSearchCV(cv=3, estimator=RandomForestRegressor(random_state=13), n_jobs=-1,
                  param_grid={'max_depth': [10, 20, 30],
                              'min_samples_leaf': [1, 2, 4],
                              'min_samples_split': [2, 5, 10],
                              'n_estimators': [100, 200, 300]})
```

```
[9]: y_pred_cv = rfr_cv.predict(X_test)
      y_pred_cv
```

```
[9]: array([79.71933333, 81.09833333, 74.83733333, 79.40433333, 71.159      ,
           79.232      , 82.41166667, 81.686      , 78.43533333, 71.404      ,
           76.47266667, 75.904      , 82.69633333, 78.64833333, 76.843      ,
           77.58633333, 71.715      , 78.93433333, 72.42533333, 82.016      ,
           71.731      , 78.45666667, 80.         , 78.529      , 78.31966667,
           72.51466667, 75.78433333, 81.29033333, 81.392      , 82.281      ,
           73.972      , 74.404      , 78.67233333, 74.99466667, 71.68766667,
           74.34333333, 74.56433333, 76.604      , 81.14366667, 80.91      ,
           74.97633333, 81.66533333, 77.59366667, 78.17166667, 80.87033333,
           83.52633333, 78.43366667, 73.93433333, 77.149      , 80.104      ,
           74.718      , 76.741      , 78.136      , 78.70566667, 77.53133333])
```

```
[10]: mean_absolute_error(y_test, y_pred_cv)
```

```
[10]: 0.25089696969702713
```

```
[11]: root_mean_squared_error(y_test, y_pred_cv)
```

```
[11]: 0.31042865768991584
```

```
[12]: r2_score(y_test, y_pred_cv)
```

```
[12]: 0.9915578528520343
```