

0-Overfitting__Underfitting

October 20, 2024

1 Overfitting and Underfitting

Overfitting occurs when a machine learning model learns not only the underlying patterns in the training data but also the noise or random fluctuations in that data. As a result, the model performs very well on the training data but poorly on new, unseen data (test data or validation data), as it fails to generalize beyond the training set.

In essence, an overfitted model becomes too complex, capturing irrelevant details that don't contribute to the actual patterns in the data.

What is Underfitting? Underfitting occurs when a machine learning model is too simple to capture the underlying patterns in the data. This results in poor performance on both the training data and the test data. The model doesn't learn enough from the training data to make accurate predictions.

In essence, an underfitted model has high bias and cannot model the complexity of the data.

1.0.1 Why do Overfitting and Underfitting happen?

Overfitting happens when the model is too flexible or too complex, often due to:

- Too many features or irrelevant features in the model.
- A high degree of the polynomial function (in the case of regression).
- Too long training time, which allows the model to fit even random noise in the training data.

Underfitting happens when the model is too simplistic, often due to:

- Too few features or not enough relevant features.
- Too simple a model (e.g., using linear regression for data that has a non-linear relationship).
- Insufficient training, meaning the model hasn't been given enough time to learn from the data.

1.0.2 How to detect Overfitting and Underfitting?

Overfitting:

- Low error on the training data.
- High error on the test data or new data.
- The model seems to memorize the training data rather than generalize patterns.

Underfitting:

- High error on both the training data and the test data.
- The model cannot capture the relationship between the input features and target variable effectively.

1.0.3 How to prevent Overfitting and Underfitting?

For Overfitting:

- **Simplify the model:** Use fewer features or reduce the complexity of the model.
- **Regularization:** Apply techniques like L1 (Lasso), L2 (Ridge), or Elastic Net regularization, which penalize overly complex models.
- **Cross-validation:** Use techniques like k-fold cross-validation to assess the model on different subsets of data.
- **Prune decision trees:** For tree-based models, pruning helps limit the depth and complexity of the model.
- **Early stopping:** In algorithms like gradient boosting, stop the training process when performance on the validation set begins to deteriorate.
- **Increase the size of the dataset:** More training data helps the model generalize better.

For Underfitting:

- **Increase model complexity:** Use more powerful models, add more features, or use higher-degree polynomial functions in regression.
- **Remove constraints:** For models like decision trees, allow the tree to grow deeper.
- **Reduce regularization:** Lessen the penalization of model complexity to allow the model to capture more details in the data.
- **Train for more epochs:** Let the model train longer (for models like neural networks) to capture more patterns.

1.0.4 Examples in Classification and Regression

Classification Example: Decision Trees Let's start by illustrating overfitting and underfitting in classification using a Decision Tree.

1. Underfitting in Classification: If you limit the depth of the decision tree (e.g., a maximum depth of 1 or 2), the tree may not be complex enough to capture the relationships in the data. As a result, it underfits the data, meaning the model cannot make accurate predictions for both the training and test sets.

2. Overfitting in Classification: If you allow the decision tree to grow too deep without any constraints (e.g., depth = 20), it may fit the training data perfectly. However, the model will likely memorize the data, fitting noise and making poor predictions on new data (overfitting).

```
[1]: from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import accuracy_score

# Load dataset
data = load_iris()
X_train, X_test, y_train, y_test = train_test_split(data.data, data.target,
    ↳test_size=0.2, random_state=42)

# Underfitting: Shallow tree (max_depth=1)
clf_underfit = DecisionTreeClassifier(max_depth=1)
clf_underfit.fit(X_train, y_train)
y_pred_train_underfit = clf_underfit.predict(X_train)
y_pred_test_underfit = clf_underfit.predict(X_test)

# Overfitting: Deep tree (max_depth=20)
clf_overfit = DecisionTreeClassifier(max_depth=20)
clf_overfit.fit(X_train, y_train)
y_pred_train_overfit = clf_overfit.predict(X_train)
y_pred_test_overfit = clf_overfit.predict(X_test)

# Calculate accuracy
train_acc_underfit = accuracy_score(y_train, y_pred_train_underfit)
test_acc_underfit = accuracy_score(y_test, y_pred_test_underfit)

train_acc_overfit = accuracy_score(y_train, y_pred_train_overfit)
test_acc_overfit = accuracy_score(y_test, y_pred_test_overfit)

print(f"Underfitting - Train Accuracy: {train_acc_underfit}, Test Accuracy:
    ↳{test_acc_underfit}")
print(f"Overfitting - Train Accuracy: {train_acc_overfit}, Test Accuracy:
    ↳{test_acc_overfit}")

```

Underfitting - Train Accuracy: 0.675, Test Accuracy: 0.6333333333333333
 Overfitting - Train Accuracy: 1.0, Test Accuracy: 1.0

1.0.5 Regression Example: Polynomial Regression

1. Underfitting in Regression: Using a linear regression model for data that has a non-linear relationship (e.g., quadratic data) will cause underfitting. The model will be too simple to capture the patterns in the data.

2. Overfitting in Regression: If you use a high-degree polynomial regression (e.g., 9th or 10th degree) on data that has a simple relationship, the model will capture noise in the training data and won't generalize well to new data.

```

[2]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures

```

```

from sklearn.metrics import root_mean_squared_error

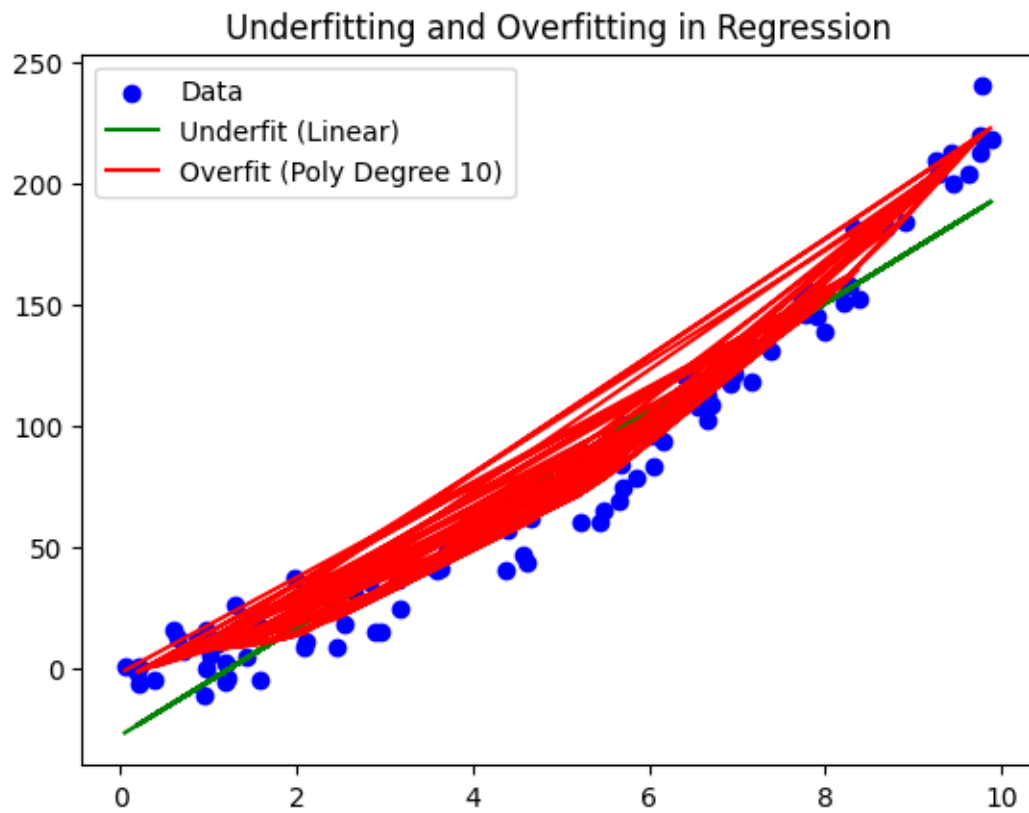
# Create a dataset (quadratic data)
np.random.seed(0)
X = np.random.rand(100, 1) * 10 # Random points from 0 to 10
y = 2 * (X ** 2) + 3 * X + np.random.randn(100, 1) * 10 # Quadratic relation
    ↳ with some noise

# Underfitting: Linear model (degree = 1)
poly_features_underfit = PolynomialFeatures(degree=1)
X_poly_underfit = poly_features_underfit.fit_transform(X)
lin_reg_underfit = LinearRegression()
lin_reg_underfit.fit(X_poly_underfit, y)
y_pred_underfit = lin_reg_underfit.predict(X_poly_underfit)

# Overfitting: High-degree polynomial model (degree = 10)
poly_features_overfit = PolynomialFeatures(degree=10)
X_poly_overfit = poly_features_overfit.fit_transform(X)
lin_reg_overfit = LinearRegression()
lin_reg_overfit.fit(X_poly_overfit, y)
y_pred_overfit = lin_reg_overfit.predict(X_poly_overfit)

# Plot results
plt.scatter(X, y, color='blue', label="Data")
plt.plot(X, y_pred_underfit, color='green', label="Underfit (Linear)")
plt.plot(X, y_pred_overfit, color='red', label="Overfit (Poly Degree 10)")
plt.title("Underfitting and Overfitting in Regression")
plt.legend()
plt.show()

```



1-train_test_split

October 20, 2024

1 Train Test Split

The train-test split is a critical concept in machine learning for evaluating the performance of a model. It involves dividing the dataset into two distinct parts: the training set and the test set. This method ensures that we can assess how well the model generalizes to unseen data, which is the key to making accurate predictions in real-world applications.

Purpose of Train-Test Split: The main goal is to evaluate the model's ability to generalize to new, unseen data by testing it on a separate test set that the model has never encountered during training. If a model performs well on both the training data and the test data, it's likely to perform well on new, real-world data.

Key Terminology:

1. **Training Set:** The subset of the dataset used to train the model. The model learns patterns, relationships, and trends from this data.
2. **Test Set:** The subset used to evaluate the model after training. It provides an unbiased estimate of the model's performance on unseen data.

Why Train-Test Split is Important:

- **Avoid Overfitting:** Without a train-test split, the model might memorize the training data (overfitting) rather than learning generalizable patterns. This would result in poor performance on new data.
- **Unbiased Evaluation:** It provides an unbiased measure of how well the model generalizes beyond the training data.
- **Generalization Performance:** It allows us to assess the model's generalization error, i.e., how it will perform on future data.

Train-Test Split Ratio:

- **80:20 Split:** Commonly used when the dataset is sufficiently large. This balance ensures that there is enough data for the model to learn (80%) while still keeping a good portion (20%) for evaluating the model's performance.
- **70:30 Split:** Used when you want more data for testing, but still have enough data for training. This might be more appropriate when the dataset is relatively smaller, as having more test data can give a clearer picture of generalization.

- 90:10 Split: Used when the dataset is very large. A small test set is enough to provide a good estimate of model performance since the model has a huge amount of data to train on.

```
[23]: import pandas as pd
from sklearn.model_selection import train_test_split
```

```
[24]: data = pd.read_csv('500hits.csv', encoding='latin-1')
data.head()
```

```
[24]:
```

	PLAYER	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	\
0	Ty Cobb	24	3035	11434	2246	4189	724	295	117	726	1249	
1	Stan Musial	22	3026	10972	1949	3630	725	177	475	1951	1599	
2	Tris Speaker	22	2789	10195	1882	3514	792	222	117	724	1381	
3	Derek Jeter	20	2747	11195	1923	3465	544	66	260	1311	1082	
4	Honus Wagner	21	2792	10430	1736	3430	640	252	101	0	963	

	SO	SB	CS	BA	HOF
0	357	892	178	0.366	1
1	696	78	31	0.331	1
2	220	432	129	0.345	1
3	1840	358	97	0.310	1
4	327	722	15	0.329	1

```
[25]: X = data.drop(columns=['PLAYER', 'HOF'])
y = data['HOF']

X.shape
y.shape
```

```
[25]: (465,)
```

```
[26]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=11,
↳ test_size=0.2)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(372, 14)
```

```
(93, 14)
```

```
(372,)
```

```
(93,)
```

```
[27]: print(X_train.describe().round(3))
```

	YRS	G	AB	R	H	2B	3B	\
count	372.000	372.000	372.000	372.000	372.000	372.000	372.000	
mean	17.011	2046.522	7526.078	1154.126	2177.995	382.505	78.094	
std	2.662	351.233	1302.406	291.308	426.615	97.173	48.798	

min	11.000	1331.000	4981.000	651.000	1660.000	177.000	3.000
25%	15.000	1797.500	6507.500	936.000	1838.000	312.000	41.000
50%	17.000	1992.000	7237.000	1099.000	2080.500	367.000	67.000
75%	19.000	2245.500	8198.250	1305.000	2383.750	436.250	108.000
max	26.000	3308.000	12364.000	2295.000	4189.000	792.000	309.000

	HR	RBI	BB	SO	SB	CS	BA
count	372.000	372.000	372.000	372.000	372.000	372.000	372.000
mean	202.642	901.073	780.105	850.323	196.927	58.987	0.289
std	141.726	484.370	327.453	472.918	185.586	49.322	0.021
min	9.000	0.000	239.000	0.000	7.000	0.000	0.246
25%	79.750	645.000	536.500	448.000	64.500	23.000	0.274
50%	185.500	977.500	719.000	844.000	141.000	52.000	0.288
75%	293.250	1218.500	961.250	1234.250	285.500	84.000	0.300
max	755.000	2297.000	2190.000	1936.000	1406.000	335.000	0.366

```
[28]: print(X_test.describe().round(3))
```

	YRS	G	AB	R	H	2B	3B \
count	93.000	93.000	93.000	93.000	93.000	93.000	93.000
mean	17.204	2057.409	7452.968	1135.065	2139.258	374.742	80.398
std	3.154	368.580	1265.371	283.877	415.165	93.929	51.792
min	11.000	1399.000	5472.000	601.000	1660.000	206.000	14.000
25%	15.000	1820.000	6622.000	935.000	1818.000	310.000	45.000
50%	17.000	1997.000	7359.000	1108.000	2054.000	361.000	68.000
75%	19.000	2282.000	8096.000	1283.000	2256.000	432.000	99.000
max	25.000	2850.000	10876.000	1859.000	3430.000	668.000	252.000

	HR	RBI	BB	SO	SB	CS	BA
count	93.000	93.000	93.000	93.000	93.000	93.000	93.000
mean	194.677	867.011	797.387	836.065	191.817	54.473	0.287
std	151.600	495.127	328.755	552.309	166.926	42.508	0.022
min	15.000	0.000	266.000	15.000	8.000	0.000	0.248
25%	78.000	618.000	527.000	359.000	61.000	15.000	0.272
50%	151.000	926.000	750.000	745.000	137.000	50.000	0.285
75%	291.000	1138.000	937.000	1179.000	271.000	83.000	0.299
max	612.000	1922.000	1747.000	2597.000	744.000	173.000	0.340

2-feature_scaling

October 20, 2024

1 Feature Scaling

is a crucial preprocessing step in machine learning where we transform the data to ensure that the features are on a similar scale. Many machine learning algorithms compute distances (e.g., Euclidean distance) or gradients that are highly sensitive to the magnitude of the input data. If one feature has a large range of values compared to others, it can dominate and distort the model's predictions. Feature scaling addresses this issue by bringing the features into comparable ranges.

Why Feature Scaling is Important:

1. **Improves Model Convergence:** Algorithms like Gradient Descent and its variants benefit from feature scaling as it makes the optimization process faster and more stable.
2. **Distance-based Algorithms:** Algorithms like KNN, K-Means clustering, and SVM rely on distance metrics. If the features have different ranges, the model could misinterpret the importance of features based on their magnitude.
3. **Regularization and PCA:** Feature scaling ensures that regularization techniques (like Lasso or Ridge) and dimensionality reduction methods (like PCA) treat all features equally.

Techniques for Feature Scaling

1. **Normalization (Min-Max Scaling)** Normalization scales the data into a fixed range, typically between 0 and 1, or sometimes -1 and 1. It's commonly used when you know that the distribution of data does not follow a Gaussian distribution, or when the algorithm expects the data to be within a particular range.
2. **Standardization (Z-Score Normalization)** Standardization transforms the data to have a mean of 0 and a standard deviation of 1. It's useful when the feature distribution follows a Gaussian distribution (bell-shaped curve) or when the algorithm assumes the data is normally distributed (e.g., logistic regression, linear regression, or SVM).

1.0.1 When to Use Normalization vs. Standardization:

- Normalization is preferred for algorithms like:
 - Neural Networks (often benefit from data scaled between 0 and 1),
 - K-Nearest Neighbors (KNN),
 - Support Vector Machines (SVM) when using an RBF kernel.
- Standardization is preferred for:
 - Linear regression,
 - Logistic regression,
 - SVM (with linear kernel),

- Principal Component Analysis (PCA),
- Regularized models (Lasso, Ridge).

```
[1]: import pandas as pd
```

```
df = pd.read_csv('500hits.csv', encoding='latin-1')
```

```
df = df.drop(columns=['PLAYER', 'CS'])
```

```
df.describe().round(3)
```

```
[1]:
```

	YRS	G	AB	R	H	2B	3B	\
count	465.000	465.000	465.000	465.000	465.000	465.000	465.000	
mean	17.049	2048.699	7511.456	1150.314	2170.247	380.953	78.555	
std	2.765	354.392	1294.066	289.635	424.191	96.483	49.363	
min	11.000	1331.000	4981.000	601.000	1660.000	177.000	3.000	
25%	15.000	1802.000	6523.000	936.000	1838.000	312.000	41.000	
50%	17.000	1993.000	7241.000	1104.000	2076.000	366.000	67.000	
75%	19.000	2247.000	8180.000	1296.000	2375.000	436.000	107.000	
max	26.000	3308.000	12364.000	2295.000	4189.000	792.000	309.000	

	HR	RBI	BB	SO	SB	BA	HOF
count	465.000	465.000	465.000	465.000	465.000	465.000	465.000
mean	201.049	894.260	783.561	847.471	195.905	0.289	0.329
std	143.623	486.193	327.432	489.224	181.846	0.021	0.475
min	9.000	0.000	239.000	0.000	7.000	0.246	0.000
25%	79.000	640.000	535.000	436.000	63.000	0.273	0.000
50%	178.000	968.000	736.000	825.000	137.000	0.287	0.000
75%	292.000	1206.000	955.000	1226.000	285.000	0.300	1.000
max	755.000	2297.000	2190.000	2597.000	1406.000	0.366	2.000

```
[2]: X1 = df.iloc[:, 0:13]
```

```
[3]: X2 = df.iloc[:, 0:13]
```

```
[4]: from sklearn.preprocessing import StandardScaler
```

```
scaleStandard = StandardScaler()
```

```
X1 = scaleStandard.fit_transform(X1)
```

```
X1 = pd.DataFrame(X1,
```

```
    columns=['YRS', 'G', 'AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'BB', 'SO', 'SB', 'BA'])
```

```
X1.head()
```

```
[4]:
```

	YRS	G	AB	R	H	2B	3B	\
0	2.516295	2.786078	3.034442	3.787062	4.764193	3.559333	4.389485	
1	1.792237	2.760655	2.677044	2.760530	3.444971	3.569709	1.996457	
2	1.792237	2.091184	2.075964	2.528955	3.171214	4.264876	2.909053	
3	1.068180	1.972543	2.849554	2.670665	3.055576	1.691719	-0.254611	
4	1.430208	2.099658	2.257758	2.024329	2.972977	2.687780	3.517449	

	HR	RBI	BB	SO	SB	BA
0	-0.585841	-0.346449	1.423013	-1.003628	3.832067	3.648290
1	1.909487	2.175837	2.493089	-0.309948	-0.649080	1.996159
2	-0.585841	-0.350567	1.826585	-1.283965	1.299723	2.657012
3	0.410896	0.858071	0.912434	2.030966	0.892346	1.004881
4	-0.697364	-1.841290	0.548609	-1.065016	2.896201	1.901752

```
[5]: X1.describe().round(3)
```

```
[5]:
```

	YRS	G	AB	R	H	2B	3B	HR \
count	465.000	465.000	465.000	465.000	465.000	465.000	465.000	465.000
mean	-0.000	0.000	0.000	0.000	0.000	0.000	-0.000	-0.000
std	1.001	1.001	1.001	1.001	1.001	1.001	1.001	1.001
min	-2.190	-2.027	-1.958	-1.899	-1.204	-2.116	-1.532	-1.339
25%	-0.742	-0.697	-0.765	-0.741	-0.784	-0.715	-0.762	-0.851
50%	-0.018	-0.157	-0.209	-0.160	-0.222	-0.155	-0.234	-0.161
75%	0.706	0.560	0.517	0.504	0.483	0.571	0.577	0.634
max	3.240	3.557	3.754	3.956	4.764	4.265	4.673	3.861

	RBI	BB	SO	SB	BA
count	465.000	465.000	465.000	465.000	465.000
mean	0.000	0.000	-0.000	0.000	0.000
std	1.001	1.001	1.001	1.001	1.001
min	-1.841	-1.665	-1.734	-1.040	-2.016
25%	-0.524	-0.760	-0.842	-0.732	-0.742
50%	0.152	-0.145	-0.046	-0.324	-0.081
75%	0.642	0.524	0.775	0.490	0.533
max	2.888	4.300	3.580	6.662	3.648

```
[6]: from sklearn.preprocessing import MinMaxScaler

scaleMinMax = MinMaxScaler(feature_range=(0,1))
X2 = scaleMinMax.fit_transform(X2)
X2 = pd.DataFrame(X2,
                  columns=['YRS', 'G', 'AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'BB', 'SO', 'SB', 'BA'])
X2.head()
```

```
[6]:
```

	YRS	G	AB	R	H	2B	3B \
0	0.866667	0.861912	0.874035	0.971074	1.000000	0.889431	0.954248
1	0.733333	0.857360	0.811459	0.795750	0.778964	0.891057	0.568627
2	0.733333	0.737481	0.706217	0.756198	0.733096	1.000000	0.715686
3	0.600000	0.716237	0.841663	0.780401	0.713721	0.596748	0.205882
4	0.666667	0.738998	0.738047	0.670012	0.699881	0.752846	0.813725

	HR	RBI	BB	SO	SB	BA
0	0.144772	0.316064	0.517683	0.137466	0.632595	1.000000

```

1  0.624665  0.849369  0.697078  0.268002  0.050751  0.708333
2  0.144772  0.315194  0.585341  0.084713  0.303788  0.825000
3  0.336461  0.570744  0.432086  0.708510  0.250893  0.533333
4  0.123324  0.000000  0.371092  0.125915  0.511079  0.691667

```

```
[7]: X2.describe().round(3)
```

```

[7]:
      count  YRS      G      AB      R      H      2B      3B      HR  \
count  465.000  465.000  465.000  465.000  465.000  465.000  465.000  465.000
mean    0.403   0.363   0.343   0.324   0.202   0.332   0.247   0.257
std     0.184   0.179   0.175   0.171   0.168   0.157   0.161   0.193
min     0.000   0.000   0.000   0.000   0.000   0.000   0.000   0.000
25%     0.267   0.238   0.209   0.198   0.070   0.220   0.124   0.094
50%     0.400   0.335   0.306   0.297   0.164   0.307   0.209   0.227
75%     0.533   0.463   0.433   0.410   0.283   0.421   0.340   0.379
max     1.000   1.000   1.000   1.000   1.000   1.000   1.000   1.000

      count  RBI      BB      SO      SB      BA
count  465.000  465.000  465.000  465.000  465.000
mean    0.389   0.279   0.326   0.135   0.356
std     0.212   0.168   0.188   0.130   0.177
min     0.000   0.000   0.000   0.000   0.000
25%     0.279   0.152   0.168   0.040   0.225
50%     0.421   0.255   0.318   0.093   0.342
75%     0.525   0.367   0.472   0.199   0.450
max     1.000   1.000   1.000   1.000   1.000

```

3-one_hot_encoder

October 20, 2024

1 One Hot Encoder

One-Hot Encoding is a technique used to represent categorical data as binary vectors. Machine learning models typically expect numerical input, but many datasets contain categorical variables (e.g., “color” can be “red,” “blue,” or “green”). One-hot encoding allows us to transform these categorical features into a format that the model can understand without assuming any ordinal relationship between the categories.

Why One-Hot Encoding is Important:

1. **Handling Categorical Data:** Many machine learning models (e.g., linear regression, decision trees, SVMs) cannot work with raw categorical data directly. One-hot encoding converts the categorical data into binary vectors that the model can process.
2. **Avoiding Misinterpretation of Ordinal Relationships:** By assigning each category its own binary feature, one-hot encoding ensures that the model does not assume any inherent order or relationship between categories, which could mislead the learning process.

How One-Hot Encoding Works: Assume we have a categorical feature “Color” with three possible values: “Red,” “Green,” and “Blue.” One-hot encoding creates a binary vector for each category:

Color	One-Hot Encoding
Red	[1, 0, 0]
Green	[0, 1, 0]
Blue	[0, 0, 1]

When to Use One-Hot Encoding:

- **Non-Ordinal Categorical Data:** One-hot encoding is used for nominal (non-ordinal) categorical data, where there is no intrinsic order among the categories (e.g., “color” or “country”).
- **Features with Small Cardinality:** Works well when the categorical feature has a small number of unique categories. When the number of unique categories (cardinality) is large, one-hot encoding can lead to a high-dimensional, sparse matrix, which can become computationally expensive and increase memory usage.

Summary:

- One-Hot Encoding is used to convert categorical variables into binary vectors.

- It is essential for handling non-ordinal categorical data in machine learning.
- It works by creating a separate binary column for each unique category.
- Care should be taken with features that have high cardinality, as it can lead to large, sparse matrices.
- The dummy variable trap can be avoided by dropping one of the one-hot encoded columns when working with linear models.

```
[2]: import pandas as pd
d = {
    'sales': [100000, 222000, 1000000, 522000, 111111, 222222, 1111111, 20000, 75000, 90000, 1000000, 10000],
    'city': ['Tampa', 'Tampa', 'Orlando', 'Jacksonville', 'Miami', 'Jacksonville', 'Miami', 'Miami', 'Orlando',
    'size': ['Small',
    'Medium', 'Large', 'Large', 'Small', 'Medium', 'Large', 'Small', 'Medium', 'Medium', 'Medium', 'Small']
}

df = pd.DataFrame(data=d)
df.head()
```

```
[2]:      sales      city  size
0   100000     Tampa  Small
1    222000     Tampa  Medium
2  1000000   Orlando   Large
3   522000 Jacksonville   Large
4   111111      Miami   Small
```

```
[3]: df['city'].unique()
```

```
[3]: array(['Tampa', 'Orlando', 'Jacksonville', 'Miami'], dtype=object)
```

```
[7]: from sklearn.preprocessing import OneHotEncoder

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False).
    set_output(transform='pandas')
ohetransform = ohe.fit_transform(df)
ohetransform
```

```
[7]:      sales_10000  sales_20000  sales_75000  sales_90000  sales_100000  \
0             0.0           0.0           0.0           0.0             1.0
1             0.0           0.0           0.0           0.0             0.0
2             0.0           0.0           0.0           0.0             0.0
3             0.0           0.0           0.0           0.0             0.0
4             0.0           0.0           0.0           0.0             0.0
5             0.0           0.0           0.0           0.0             0.0
6             0.0           0.0           0.0           0.0             0.0
```

7	0.0	1.0	0.0	0.0	0.0
8	0.0	0.0	1.0	0.0	0.0
9	0.0	0.0	0.0	1.0	0.0
10	0.0	0.0	0.0	0.0	0.0
11	1.0	0.0	0.0	0.0	0.0

	sales_1111111	sales_222000	sales_222222	sales_522000	sales_1000000 \
0	0.0	0.0	0.0	0.0	0.0
1	0.0	1.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	1.0
3	0.0	0.0	0.0	1.0	0.0
4	1.0	0.0	0.0	0.0	0.0
5	0.0	0.0	1.0	0.0	0.0
6	0.0	0.0	0.0	0.0	0.0
7	0.0	0.0	0.0	0.0	0.0
8	0.0	0.0	0.0	0.0	0.0
9	0.0	0.0	0.0	0.0	0.0
10	0.0	0.0	0.0	0.0	1.0
11	0.0	0.0	0.0	0.0	0.0

	sales_1111111	city_Jacksonville	city_Miami	city_Orlando	city_Tampa \
0	0.0	0.0	0.0	0.0	1.0
1	0.0	0.0	0.0	0.0	1.0
2	0.0	0.0	0.0	1.0	0.0
3	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0
5	0.0	1.0	0.0	0.0	0.0
6	1.0	0.0	1.0	0.0	0.0
7	0.0	0.0	1.0	0.0	0.0
8	0.0	0.0	0.0	1.0	0.0
9	0.0	0.0	0.0	1.0	0.0
10	0.0	0.0	0.0	1.0	0.0
11	0.0	0.0	0.0	1.0	0.0

	size_Large	size_Medium	size_Small
0	0.0	0.0	1.0
1	0.0	1.0	0.0
2	1.0	0.0	0.0
3	1.0	0.0	0.0
4	0.0	0.0	1.0
5	0.0	1.0	0.0
6	1.0	0.0	0.0
7	0.0	0.0	1.0
8	0.0	1.0	0.0
9	0.0	1.0	0.0
10	0.0	1.0	0.0
11	0.0	0.0	1.0

```
[9]: df = pd.concat([df, ohetransform], axis=1)
df
```

```
[9]:
```

	sales	city	size	sales_10000	sales_20000	sales_75000	\
0	100000	Tampa	Small	0.0	0.0	0.0	
1	222000	Tampa	Medium	0.0	0.0	0.0	
2	1000000	Orlando	Large	0.0	0.0	0.0	
3	522000	Jacksonville	Large	0.0	0.0	0.0	
4	111111	Miami	Small	0.0	0.0	0.0	
5	222222	Jacksonville	Medium	0.0	0.0	0.0	
6	1111111	Miami	Large	0.0	0.0	0.0	
7	20000	Miami	Small	0.0	1.0	0.0	
8	75000	Orlando	Medium	0.0	0.0	1.0	
9	90000	Orlando	Medium	0.0	0.0	0.0	
10	1000000	Orlando	Medium	0.0	0.0	0.0	
11	10000	Orlando	Small	1.0	0.0	0.0	

	sales_90000	sales_100000	sales_111111	sales_222000	...	sales_522000	\
0	0.0	1.0	0.0	0.0	...	0.0	
1	0.0	0.0	0.0	1.0	...	0.0	
2	0.0	0.0	0.0	0.0	...	0.0	
3	0.0	0.0	0.0	0.0	...	1.0	
4	0.0	0.0	1.0	0.0	...	0.0	
5	0.0	0.0	0.0	0.0	...	0.0	
6	0.0	0.0	0.0	0.0	...	0.0	
7	0.0	0.0	0.0	0.0	...	0.0	
8	0.0	0.0	0.0	0.0	...	0.0	
9	1.0	0.0	0.0	0.0	...	0.0	
10	0.0	0.0	0.0	0.0	...	0.0	
11	0.0	0.0	0.0	0.0	...	0.0	

	sales_1000000	sales_1111111	city_Jacksonville	city_Miami	city_Orlando	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	1.0	
3	0.0	0.0	1.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	0.0	
5	0.0	0.0	1.0	0.0	0.0	
6	0.0	1.0	0.0	1.0	0.0	
7	0.0	0.0	0.0	1.0	0.0	
8	0.0	0.0	0.0	0.0	1.0	
9	0.0	0.0	0.0	0.0	1.0	
10	1.0	0.0	0.0	0.0	1.0	
11	0.0	0.0	0.0	0.0	1.0	

	city_Tampa	size_Large	size_Medium	size_Small
0	1.0	0.0	0.0	1.0

1	1.0	0.0	1.0	0.0
2	0.0	1.0	0.0	0.0
3	0.0	1.0	0.0	0.0
4	0.0	0.0	0.0	1.0
5	0.0	0.0	1.0	0.0
6	0.0	1.0	0.0	0.0
7	0.0	0.0	0.0	1.0
8	0.0	0.0	1.0	0.0
9	0.0	0.0	1.0	0.0
10	0.0	0.0	1.0	0.0
11	0.0	0.0	0.0	1.0

[12 rows x 39 columns]

4-ordinal_encoder

October 20, 2024

1 Ordinal Encoder

is a technique used to convert categorical features into numerical values, where the categories have a clear order or ranking. Unlike one-hot encoding, which treats all categories as independent, ordinal encoding assigns each unique category an integer value based on its rank or order.

Why Ordinal Encoding is Important:

- **Ordered Categories:** It is specifically used for ordinal data, where the categories have a meaningful order or ranking, but the distances between these categories might not be uniform.
- **Simpler Representation:** It provides a simpler, more compact representation of categorical features compared to one-hot encoding, especially when the feature has an inherent order and a small number of categories.

When to Use Ordinal Encoding:

- **Ordinal Categorical Data:** Use ordinal encoding when the categorical feature has an intrinsic order. For example, “low,” “medium,” and “high” represent ordered categories, but one-hot encoding would not capture the ranking.
- **When Preserving Order is Important:** If the relationship between categories is important for the model to understand, ordinal encoding helps preserve this information. **Smaller Cardinality:** When dealing with ordinal features with a small number of categories, ordinal encoding is efficient and works well with most machine learning algorithms.

How Ordinal Encoding Works: Let’s take an example where we have a categorical feature “Quality” with values: “Low,” “Medium,” and “High.” Ordinal encoding assigns each category a numerical value based on the order:

Quality Low 0 Medium 1 High 2

Color	Ordinal Encoding
Low	0
Medium	1
High	2

In this case, “Low” is mapped to 0, “Medium” to 1, and “High” to 2, which captures the inherent ranking of the categories.

When Not to Use Ordinal Encoding:

- Non-Ordinal Categorical Data: If the categorical feature has no inherent order (e.g., “color” or “country”), ordinal encoding may mislead the model because it implies a relationship or ranking between the categories that doesn’t exist.
- Assumed Numerical Relationships: Some machine learning models, like linear regression, may interpret ordinal values as having linear relationships (i.e., assuming that the difference between “Medium” and “Low” is the same as between “High” and “Medium”), which is not always correct.

Use Cases of Ordinal Encoding:

- Education Levels: Categories such as “Primary,” “Secondary,” and “Tertiary” have an inherent order, making them a perfect candidate for ordinal encoding.
- Customer Satisfaction Levels: Responses like “Very Dissatisfied,” “Dissatisfied,” “Neutral,” “Satisfied,” and “Very Satisfied” can be encoded in their natural order.
- Rating Scales: Ratings such as “Low,” “Medium,” and “High” in product reviews can be ordinal encoded since they imply a ranking.

Summary:

- Ordinal Encoding converts categorical data into numerical data, preserving the order of categories.
- It is used when the categorical variable has an inherent order (e.g., “low,” “medium,” “high”).
- It is compact and efficient but may lead to misleading interpretations if used on non-ordinal categorical data.
- Some algorithms (e.g., decision trees, random forests) work well with ordinal encoding, but others (e.g., linear models) may misinterpret the encoding.

```
[1]: import pandas as pd

d = {'sales': [100000, 222000, 1000000, 522000, 111111, 222222, 111111, 20000, 75000, 90000, 1000000, 10000],
     'city': ['Tampa', 'Tampa', 'Orlando', 'Jacksonville', 'Miami', 'Jacksonville', 'Miami', 'Miami', 'Orlando',
              'Medium', 'Large', 'Large', 'Small', 'Medium', 'Large', 'Small', 'Medium', 'Medium', 'Medium', 'Small'],
     'size': ['Small', 'Medium', 'Large', 'Small', 'Medium', 'Large', 'Small', 'Medium', 'Medium', 'Medium', 'Small']}

df = pd.DataFrame(data=d)
df.head()
```

```
[1]:
```

	sales	city	size
0	100000	Tampa	Small
1	222000	Tampa	Medium
2	1000000	Orlando	Large
3	522000	Jacksonville	Large

```
4    111111      Miami    Small
```

```
[2]: df['size'].unique()
      sizes = ['Small', 'Medium', 'Large']
```

```
[3]: from sklearn.preprocessing import OrdinalEncoder
      enc = OrdinalEncoder(categories = [sizes])
      size_enc = enc.fit_transform(df[['size']])
      size_enc
```

```
[3]: array([[0.],
            [1.],
            [2.],
            [2.],
            [0.],
            [1.],
            [2.],
            [0.],
            [1.],
            [1.],
            [1.],
            [0.]])
```

```
[4]: df['size_enc'] = size_enc
      df
```

```
[4]:
```

	sales	city	size	size_enc
0	100000	Tampa	Small	0.0
1	222000	Tampa	Medium	1.0
2	1000000	Orlando	Large	2.0
3	522000	Jacksonville	Large	2.0
4	111111	Miami	Small	0.0
5	222222	Jacksonville	Medium	1.0
6	1111111	Miami	Large	2.0
7	20000	Miami	Small	0.0
8	75000	Orlando	Medium	1.0
9	90000	Orlando	Medium	1.0
10	1000000	Orlando	Medium	1.0
11	10000	Orlando	Small	0.0

5-handling-missing-data

October 20, 2024

1 Handling Missing Data

Handling missing data is a crucial step in the data preprocessing phase of machine learning projects. Missing values can arise from various reasons, such as incomplete data collection, sensor errors, or data entry issues. If not handled properly, missing data can significantly degrade model performance or cause errors during training and testing.

Why Handling Missing Data is Important:

1. **Maintaining Data Integrity:** Missing values can skew analysis results, leading to incorrect conclusions and poor model performance.
2. **Ensuring Model Robustness:** Most machine learning models (especially algorithms like linear regression, logistic regression, or SVM) cannot handle missing values directly, so preprocessing is necessary to ensure smooth execution.
3. **Preserving Dataset Size:** Rather than discarding records with missing values, handling them appropriately allows you to use as much of your data as possible.

Different Techniques for Handling Missing Data:

1. **Removing Missing Data:**
 - **Dropping Rows:** If the number of rows with missing values is small, you can drop those rows.
 - **Dropping Columns:** If an entire column has too many missing values (e.g., more than 50%), you may choose to drop that feature.
2. **Imputation (Filling Missing Data):**
 - **Mean/Median/Mode Imputation:** Replace missing values with the mean, median, or mode of the non-missing values in that column.
 - **Forward/Backward Fill:** Use previous or next values to fill missing entries (mainly used in time-series data).
 - **Predictive Imputation:** Use a machine learning model to predict the missing values based on other features.
 - **K-Nearest Neighbors (KNN) Imputation:** Use KNN to find similar data points and fill missing values based on the nearest neighbors.
3. **Flagging/Indicator Variables:**

- Create a new binary feature indicating whether the value in the original column was missing or not. This approach allows the model to know which data points had missing values.
4. Using Algorithms that Handle Missing Data:
- Some machine learning algorithms like decision trees and random forests can handle missing values internally without requiring imputation.

Practical Considerations:

1. Nature of Missing Data: Before choosing a method, understand the nature of missing data. Is it Missing Completely at Random (MCAR), Missing at Random (MAR), or Not Missing at Random (NMAR)? Different techniques are better suited for different types of missingness.
2. Preserving Relationships: While imputation preserves data, it may introduce bias or distort relationships between features, especially if the missing data is not random..
3. Multiple Imputation: For more advanced scenarios, techniques like Multiple Imputation (e.g., MICE - Multiple Imputation by Chained Equations) are used to generate multiple datasets with different imputations, combining their results to reduce imputation bias.

Summary:

- Removing Missing Data: Drop rows or columns with missing values if the missing proportion is small.
- Imputation: Replace missing values with mean, median, mode, or more advanced techniques like KNN or predictive imputation.
- Flagging: Create binary flags for missing values to allow the model to capture this information.
- Advanced Algorithms: Some algorithms like decision trees and XGBoost handle missing values naturally without preprocessing.

```
[1]: import pandas as pd
import numpy as np

miles = pd.DataFrame({'farthest_run_mi': [50,62, np.nan,100,26,13,31,50]})
miles
```

```
[1]:   farthest_run_mi
0          50.0
1          62.0
2           NaN
3         100.0
4          26.0
5          13.0
6          31.0
7          50.0
```

```
[2]: miles.isna().sum()
```

```
[2]: farthest_run_mi    1
      dtype: int64
```

```
[3]: from sklearn.impute import SimpleImputer
      imp_mean = SimpleImputer(strategy='mean')
      imp_mean.fit_transform(miles)
```

```
[3]: array([[ 50.      ],
             [ 62.      ],
             [ 47.42857143],
             [100.      ],
             [ 26.      ],
             [ 13.      ],
             [ 31.      ],
             [ 50.      ]])
```

```
[4]: imp_median = SimpleImputer(strategy='median')
      imp_median.fit_transform(miles)
```

```
[4]: array([[ 50.],
             [ 62.],
             [ 50.],
             [100.],
             [ 26.],
             [ 13.],
             [ 31.],
             [ 50.]])
```

```
[5]: imp_mode = SimpleImputer(strategy='most_frequent')
      imp_mode.fit_transform(miles)
```

```
[5]: array([[ 50.],
             [ 62.],
             [ 50.],
             [100.],
             [ 26.],
             [ 13.],
             [ 31.],
             [ 50.]])
```

```
[6]: imp_constant = SimpleImputer(strategy='constant', fill_value=13)
      imp_constant.fit_transform(miles)
```

```
[6]: array([[ 50.],
             [ 62.],
             [ 13.],
             [100.],
             [ 26.]])
```

```

[ 13.],
[ 31.],
[ 50.]]))

```

```

[7]: names = pd.DataFrame({'names': ['Ryan', 'Nolan', 'Honus', 'Wagner', np.nan,
↳ 'Ruth']})
names

```

```

[7]:      names
0      Ryan
1      Nolan
2      Honus
3      Wagner
4         NaN
5       Ruth

```

```

[8]: imp_constant_cat = SimpleImputer(strategy='constant', fill_value='babe')
imp_constant_cat.fit_transform(names)

```

```

[8]: array([[ 'Ryan'],
        [ 'Nolan'],
        [ 'Honus'],
        [ 'Wagner'],
        [ 'babe'],
        [ 'Ruth']], dtype=object)

```

```

[9]: imp_mean_marked = SimpleImputer(strategy='mean', add_indicator= True)
imp_mean_marked.fit_transform(miles)

```

```

[9]: array([[ 50.         ,  0.         ],
        [ 62.         ,  0.         ],
        [ 47.42857143,  1.         ],
        [100.         ,  0.         ],
        [ 26.         ,  0.         ],
        [ 13.         ,  0.         ],
        [ 31.         ,  0.         ],
        [ 50.         ,  0.         ]])

```

```

[13]: new_df = pd.DataFrame({
        'Name': ['Ryan', 'Nolan', 'Walter', 'Honus', 'Christy', np.nan, 'Napoleon',
↳ 'Tris'],
        'farthest_run_mi': [50,62, np.nan,100,26,13,31,50]
    })
new_df

```

```

[13]:      Name  farthest_run_mi
0      Ryan             50.0

```


1	Nolan	62.0
2	Walter	NaN
3	Honus	100.0
4	Christy	26.0
5	NaN	13.0
6	Napoleon	31.0
7	Tris	50.0

```
[14]: from sklearn.compose import make_column_transformer
ct = make_column_transformer(
    (imp_constant_cat, ['Name']),
    (imp_mean, ['farthest_run_mi']),
    remainder='drop'
)
ct.set_output(transform='pandas')
```

```
[14]: ColumnTransformer(transformers=[('simpleimputer-1',
                                       SimpleImputer(fill_value='babe',
                                                       strategy='constant'),
                                       ['Name']),
                                       ('simpleimputer-2', SimpleImputer(),
                                       ['farthest_run_mi'])])
```

```
[15]: df_pandas = ct.fit_transform(new_df)
df_pandas
```

```
[15]: simpleimputer-1__Name  simpleimputer-2__farthest_run_mi
0           Ryan          50.000000
1           Nolan          62.000000
2           Walter          47.428571
3           Honus          100.000000
4           Christy         26.000000
5           babe           13.000000
6           Napoleon         31.000000
7           Tris           50.000000
```

6-data_preprocessing_with_column_transformer

October 20, 2024

1 Data Preprocessing With Column Transformer

The ColumnTransformer in Python's sklearn library is a powerful tool that allows you to apply different preprocessing steps to different columns of your dataset. This is particularly useful when you have a mixture of numerical and categorical data and need to apply distinct transformations to each type of data. Using the ColumnTransformer, you can combine transformations like scaling numerical features, encoding categorical variables, and imputing missing values, all in a single unified process.

Why Use ColumnTransformer?

1. **Multiple Data Types:** It handles datasets with mixed data types (numerical, categorical, text, etc.).
2. **Efficient Pipeline:** You can combine different preprocessing steps into a pipeline, making the code cleaner and easier to maintain.
3. **Selective Preprocessing:** Apply different preprocessing steps to specific columns, instead of preprocessing the entire dataset uniformly.
4. **Integration with Pipelines:** It can be combined with machine learning models in a Pipeline, allowing for a streamlined workflow from preprocessing to model training and evaluation.

```
[2]: import pandas as pd

d = {'sales': [
    100000, 222000, 1000000, 522000, 111111, 222222, 1111111, 20000, 75000, 90000, 1000000, 10000],
     'city': [
    'Tampa', 'Tampa', 'Orlando', 'Jacksonville', 'Miami', 'Jacksonville', 'Miami', 'Miami', 'Orlando',
    'size': ['Small',
    'Medium', 'Large', 'Large', 'Small', 'Medium', 'Large', 'Small', 'Medium', 'Medium', 'Medium', 'Small']

df = pd.DataFrame(data=d)
df
```

```
[2]:      sales      city      size
0   100000     Tampa   Small
1   222000     Tampa  Medium
2  1000000   Orlando   Large
```

3	522000	Jacksonville	Large
4	111111	Miami	Small
5	222222	Jacksonville	Medium
6	111111	Miami	Large
7	20000	Miami	Small
8	75000	Orlando	Medium
9	90000	Orlando	Medium
10	1000000	Orlando	Medium
11	10000	Orlando	Small

```
[7]: from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.compose import make_column_transformer

ohe = OneHotEncoder(sparse_output=False)
ode = OrdinalEncoder()

ct = make_column_transformer(
    (ohe, ['city']),
    (ode, ['size']),
    remainder='passthrough'
)

ct.set_output(transform='pandas')
```

```
[7]: ColumnTransformer(remainder='passthrough',
                        transformers=[('onehotencoder',
                                      OneHotEncoder(sparse_output=False), ['city']),
                                      ('ordinalencoder', OrdinalEncoder(), ['size'])])
```

```
[8]: df_pandas = ct.fit_transform(df)
df_pandas
```

```
[8]:
```

	onehotencoder__city_Jacksonville	onehotencoder__city_Miami	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	1.0	0.0	
4	0.0	1.0	
5	1.0	0.0	
6	0.0	1.0	
7	0.0	1.0	
8	0.0	0.0	
9	0.0	0.0	
10	0.0	0.0	
11	0.0	0.0	

	onehotencoder__city_Orlando	onehotencoder__city_Tampa	\
--	-----------------------------	---------------------------	---

0	0.0	1.0
1	0.0	1.0
2	1.0	0.0
3	0.0	0.0
4	0.0	0.0
5	0.0	0.0
6	0.0	0.0
7	0.0	0.0
8	1.0	0.0
9	1.0	0.0
10	1.0	0.0
11	1.0	0.0

	ordinalencoder__size	remainder__sales
0	2.0	100000
1	1.0	222000
2	0.0	1000000
3	0.0	522000
4	2.0	111111
5	1.0	222222
6	0.0	1111111
7	2.0	20000
8	1.0	75000
9	1.0	90000
10	1.0	1000000
11	2.0	10000

```
[9]: ct2 = make_column_transformer(
      (ohe, [1]),
      (ode, [2]),
      remainder='drop'
    )

    ct2.set_output(transform='pandas')
```

```
[9]: ColumnTransformer(transformers=[('onehotencoder',
                                     OneHotEncoder(sparse_output=False), [1]),
                                     ('ordinalencoder', OrdinalEncoder(), [2])])
```

```
[11]: df_pandas2 = ct2.fit_transform(df)
      df_pandas2
```

```
[11]: onehotencoder__city_Jacksonville  onehotencoder__city_Miami  \
0                                     0.0                        0.0
1                                     0.0                        0.0
2                                     0.0                        0.0
3                                     1.0                        0.0
```

4	0.0	1.0
5	1.0	0.0
6	0.0	1.0
7	0.0	1.0
8	0.0	0.0
9	0.0	0.0
10	0.0	0.0
11	0.0	0.0

	onehotencoder__city_Orlando	onehotencoder__city_Tampa \
0	0.0	1.0
1	0.0	1.0
2	1.0	0.0
3	0.0	0.0
4	0.0	0.0
5	0.0	0.0
6	0.0	0.0
7	0.0	0.0
8	1.0	0.0
9	1.0	0.0
10	1.0	0.0
11	1.0	0.0

	ordinalencoder__size
0	2.0
1	1.0
2	0.0
3	0.0
4	2.0
5	1.0
6	0.0
7	2.0
8	1.0
9	1.0
10	1.0
11	2.0

```
[12]: ct3 = make_column_transformer(
      (ohe, [1]),
      ('passthrough', ['size']),
      remainder='drop'
    )

    ct3.set_output(transform='pandas')
```

```
[12]: ColumnTransformer(transformers=[('onehotencoder',
                                       OneHotEncoder(sparse_output=False), [1]),
```

```
('passthrough', 'passthrough', ['size']))
```

```
[13]: ct3.fit_transform(df)
```

```
[13]:
```

	onehotencoder__city_Jacksonville	onehotencoder__city_Miami	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	1.0	0.0	
4	0.0	1.0	
5	1.0	0.0	
6	0.0	1.0	
7	0.0	1.0	
8	0.0	0.0	
9	0.0	0.0	
10	0.0	0.0	
11	0.0	0.0	

	onehotencoder__city_Orlando	onehotencoder__city_Tampa	passthrough__size
0	0.0	1.0	Small
1	0.0	1.0	Medium
2	1.0	0.0	Large
3	0.0	0.0	Large
4	0.0	0.0	Small
5	0.0	0.0	Medium
6	0.0	0.0	Large
7	0.0	0.0	Small
8	1.0	0.0	Medium
9	1.0	0.0	Medium
10	1.0	0.0	Medium
11	1.0	0.0	Small

7-KNN

October 20, 2024

1 K Nearest Neighbor

K-Nearest Neighbor (KNN) is a supervised machine learning algorithm used for classification and regression tasks. It classifies a new data point based on how its neighbors (the closest data points in the feature space) are classified. In essence, KNN assigns the class of the majority of its K nearest neighbors to the new point. It's a non-parametric and lazy learning algorithm, meaning it makes decisions based on the entire dataset rather than assuming a specific structure (non-parametric), and it doesn't explicitly learn a model during training (lazy learning).

- **Non-parametric:** No assumptions about the data distribution.
- **Lazy learning:** The model simply stores the data and waits until a query (test data) is made before computing anything.

When to Use KNN? KNN is suitable when you have a labeled dataset and you want to classify new, unseen data. However, it is most effective under the following conditions:

- **Small to medium-sized datasets:** KNN can become computationally expensive with large datasets, as it requires scanning through all data points to find the neighbors.
- **Low-dimensional data:** With high-dimensional data, the distance between data points tends to increase, making it harder for the algorithm to distinguish between neighbors (this is known as the curse of dimensionality).
- **Non-linear decision boundaries:** If the decision boundaries between classes are not linear (i.e., they can't be split by a line or hyperplane), KNN can work well since it uses local information. Balanced classes: KNN works best when the number of samples in each class is roughly equal.

How Does KNN Work?

1. **Choose K (the number of neighbors):** This is a user-defined parameter. Typically, an odd number is chosen to avoid ties in classification. For example, K=3 means the algorithm looks at the 3 nearest neighbors of a data point.
2. **Calculate Distance:** For a new data point, the algorithm calculates the distance between this point and all other points in the training dataset. The most common distance metric used is Euclidean distance.

Other distance metrics can be used, like Manhattan distance or Minkowski distance, depending on the dataset and the problem.

3. **Find the K Nearest Neighbors:** After calculating the distance between the new data point and every other point in the dataset, KNN selects the K closest points.
4. **Assign a Class or Predict Value:**
 - **For classification:** The class label that appears most frequently among the K nearest neighbors is assigned to the new data point.
 - **For regression:** The mean (or sometimes median) of the values of the K nearest neighbors is taken as the predicted value for the new data point.
5. **Result:** The algorithm outputs either the predicted class label (for classification) or a predicted value (for regression).

Advantages of KNN:

- **Simple to implement:** KNN requires no assumptions about the underlying data, making it easy to use.
- **Flexible:** Can be used for both classification and regression problems.
- **No training phase:** As a lazy learner, KNN does not require a separate training phase, making it efficient in terms of time before making predictions.

Disadvantages of KNN:

- **Computationally expensive:** As it compares every new point with all points in the training data, it can be slow with large datasets.
- **Memory-intensive:** Since the entire dataset needs to be stored, it requires more memory as the dataset grows.
- **Sensitive to irrelevant features:** If the dataset has many irrelevant features, the distance metric may be distorted, leading to poor performance.
- **Sensitive to the choice of K:** A poor choice of K can lead to underfitting (if K is too large) or overfitting (if K is too small).

Real-World Applications:

1. **Image Recognition:** KNN can classify images based on similarity, such as facial recognition or object classification.
2. **Recommender Systems:** KNN can be used to find similar items to recommend to users based on their preferences.
3. **Medical Diagnosis:** It's used to classify diseases based on the symptoms and historical data of patients.
4. **Finance:** KNN is used to detect outliers or categorize financial transactions (e.g., fraud detection).
5. **Document Classification:** Used in text mining to classify documents into predefined categories.


```
[1]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report
```

```
[2]: df = pd.read_csv('500hits.csv', encoding="latin-1")
df.head()
```

```
[2]:
```

	PLAYER	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	\
0	Ty Cobb	24	3035	11434	2246	4189	724	295	117	726	1249	
1	Stan Musial	22	3026	10972	1949	3630	725	177	475	1951	1599	
2	Tris Speaker	22	2789	10195	1882	3514	792	222	117	724	1381	
3	Derek Jeter	20	2747	11195	1923	3465	544	66	260	1311	1082	
4	Honus Wagner	21	2792	10430	1736	3430	640	252	101	0	963	

	SO	SB	CS	BA	HOF
0	357	892	178	0.366	1
1	696	78	31	0.331	1
2	220	432	129	0.345	1
3	1840	358	97	0.310	1
4	327	722	15	0.329	1

```
[3]: df = df.drop(columns=['PLAYER', 'CS'])
df.head()
```

```
[3]:
```

	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	SO	SB	BA	\
0	24	3035	11434	2246	4189	724	295	117	726	1249	357	892	0.366	
1	22	3026	10972	1949	3630	725	177	475	1951	1599	696	78	0.331	
2	22	2789	10195	1882	3514	792	222	117	724	1381	220	432	0.345	
3	20	2747	11195	1923	3465	544	66	260	1311	1082	1840	358	0.310	
4	21	2792	10430	1736	3430	640	252	101	0	963	327	722	0.329	

	HOF
0	1
1	1
2	1
3	1
4	1

```
[6]: X = df.iloc[:,0:13]
y = df.iloc[:,13]
```

```
[7]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=11,
↳test_size=0.2)

scaler = MinMaxScaler(feature_range=(0,1))
```

```
X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)
```

```
[10]: knn = KNeighborsClassifier(n_neighbors=8)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)
y_pred
```

```
[10]: array([0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
          0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
          0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0,
          1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
          0, 0, 0, 1, 1], dtype=int64)
```

```
[11]: knn.score(X_test, y_test)
```

```
[11]: 0.8279569892473119
```

1.1 Confusion Matrix

```
[12]: cm = confusion_matrix(y_test, y_pred)
cm
```

```
[12]: array([[55, 12],
          [ 4, 22]], dtype=int64)
```

1.2 Classification Report

```
[15]: cr = classification_report(y_test, y_pred)
print(cr)
```

	precision	recall	f1-score	support
0	0.93	0.82	0.87	67
1	0.65	0.85	0.73	26
accuracy			0.83	93
macro avg	0.79	0.83	0.80	93
weighted avg	0.85	0.83	0.83	93

```
[16]: print(knn.n_samples_fit_)
```

```
372
```

8-Decision_Tree

October 20, 2024

1 Decision Tree

A Decision Tree is a supervised machine learning algorithm used for classification and regression tasks. It works by breaking down a dataset into smaller and smaller subsets while incrementally developing a tree structure. Each internal node of the tree represents a test or decision based on an attribute (feature), each branch represents the outcome of the decision, and each leaf node represents a class label (in classification) or a continuous value (in regression).

- **Root node:** The top node, which represents the entire dataset and is split based on the most important feature.
- **Internal nodes:** Represent the attributes or features that are used to split the data.
- **Leaf nodes:** Represent the final class or output value.

The main goal of a Decision Tree is to create a model that predicts the target value by learning simple decision rules inferred from the data features.

When to Use Decision Trees? Decision Trees are useful when you want a model that:

- **Can handle both categorical and numerical data:** Decision Trees are flexible and can work well with different types of data.
- **Is easy to interpret:** Since Decision Trees mimic human decision-making processes, they are intuitive and easy to understand, even for non-experts.
- **Requires minimal data preprocessing:** Unlike many other algorithms, Decision Trees do not require normalization or scaling of data.
- **Can model non-linear relationships:** Decision Trees can capture complex patterns in data, including interactions between features.

However, Decision Trees are especially effective when:

- The dataset is not too large, as deep trees can become computationally expensive.
- There is a need for a simple, interpretable model.
- You are working with a problem that involves a sequence of decision

How Does a Decision Tree Work?

1. **Feature Selection and Splitting:** The tree starts with all the data at the root node. It then selects a feature and splits the data based on this feature to form child nodes. The feature is chosen by evaluating different splitting criteria (more on this later).
2. **Recursive Partitioning:** This process of splitting the data continues recursively at each child node, selecting the best feature at each step until one of the following conditions is met:
 - All samples at a node belong to the same class (for classification).
 - The node reaches a pre-defined depth (maximum depth of the tree).
 - The number of samples at a node is less than the minimum split size.
3. **Prediction:** Once the tree has been constructed, it can be used to classify new samples by passing them down the tree, following the decisions at each node until a leaf node is reached. The class label (or value for regression) at the leaf node is the model's prediction.

Who Should Use Decision Trees? Decision Trees are ideal for:

- **Business analysts and decision-makers:** The model is easy to interpret and can provide insights into the data and important decision points.
- **Data scientists and machine learning engineers** who need to solve classification or regression problems.
- **Researchers** in fields like biology, medicine, or finance who deal with complex datasets involving interactions between multiple features.

Advantages of Decision Trees:

- **Easy to interpret:** Even non-technical people can understand the decision process.
- **Handles both numerical and categorical data:** Can work with different types of features without requiring feature transformation.
- **No need for feature scaling:** Unlike algorithms like SVM or KNN, Decision Trees don't require scaling.
- **Can capture non-linear relationships:** Decision Trees can split data at any point, allowing them to model complex relationships.

Disadvantages of Decision Trees:

- **Prone to overfitting:** Decision Trees can create very complex models that overfit the training data, especially when the tree grows too deep.
- **Unstable:** Small changes in the data can result in significantly different trees.
- **Bias towards dominant classes:** In unbalanced datasets, the tree might be biased toward the dominant class.

To avoid these issues, Decision Trees are often pruned (cutting off branches that do not provide additional information) or ensemble methods like Random Forest or Gradient Boosting are used to create a more robust model.

Real-World Applications:

1. **Loan Default Prediction:** Financial institutions use Decision Trees to classify loan applicants based on their risk profile.
2. **Medical Diagnosis:** Decision Trees help in diagnosing diseases based on patient symptoms and medical records.
3. **Customer Churn Prediction:** Companies use Decision Trees to predict if a customer will leave based on historical data.
4. **Marketing:** Decision Trees are used to segment customers based on purchasing behavior and demographics.
5. **Fraud Detection:** Decision Trees help identify fraudulent transactions by learning decision rules from historical data.

```
[142]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, classification_report

df = pd.read_csv('500hits.csv', encoding="latin-1")
df.head()
```

```
[142]:
```

	PLAYER	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	\
0	Ty Cobb	24	3035	11434	2246	4189	724	295	117	726	1249	
1	Stan Musial	22	3026	10972	1949	3630	725	177	475	1951	1599	
2	Tris Speaker	22	2789	10195	1882	3514	792	222	117	724	1381	
3	Derek Jeter	20	2747	11195	1923	3465	544	66	260	1311	1082	
4	Honus Wagner	21	2792	10430	1736	3430	640	252	101	0	963	

	SO	SB	CS	BA	HOF
0	357	892	178	0.366	1
1	696	78	31	0.331	1
2	220	432	129	0.345	1
3	1840	358	97	0.310	1
4	327	722	15	0.329	1

```
[143]: df = df.drop(columns=['PLAYER', 'CS'])
df.head()
```

```
[143]:
```

	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	SO	SB	BA	\
0	24	3035	11434	2246	4189	724	295	117	726	1249	357	892	0.366	
1	22	3026	10972	1949	3630	725	177	475	1951	1599	696	78	0.331	
2	22	2789	10195	1882	3514	792	222	117	724	1381	220	432	0.345	
3	20	2747	11195	1923	3465	544	66	260	1311	1082	1840	358	0.310	
4	21	2792	10430	1736	3430	640	252	101	0	963	327	722	0.329	

	HOF
0	1
1	1
2	1
3	1
4	1

```

0    1
1    1
2    1
3    1
4    1

```

```

[144]: X = df.iloc[:, 0:13]
y = df.iloc[:, 13]

X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=17,
↳test_size=0.2)

dtc = DecisionTreeClassifier()

dtc.get_params()

```

```

[144]: {'ccp_alpha': 0.0,
'class_weight': None,
'criterion': 'gini',
'max_depth': None,
'max_features': None,
'max_leaf_nodes': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'monotonic_cst': None,
'random_state': None,
'splitter': 'best'}

```

```

[145]: dtc.fit(X_train, y_train)

```

```

[145]: DecisionTreeClassifier()

```

```

[146]: y_prediction = dtc.predict(X_test)

print(confusion_matrix(y_test, y_prediction))

```

```

[[52  9]
 [11 21]]

```

```

[147]: print(classification_report(y_test, y_prediction))

```

	precision	recall	f1-score	support
0	0.83	0.85	0.84	61
1	0.70	0.66	0.68	32
accuracy			0.78	93

macro avg	0.76	0.75	0.76	93
weighted avg	0.78	0.78	0.78	93

```
[148]: dtc.feature_importances_
```

```
[148]: array([0.          , 0.02598978, 0.03173403, 0.03633493, 0.39759474,
          0.06589131, 0.01565832, 0.05810452, 0.04515849, 0.12784069,
          0.04098071, 0.05207056, 0.10264192])
```

```
[151]: features = pd.DataFrame(dtc.feature_importances_, index=X.columns)
features
```

```
[151]:          0
YRS  0.000000
G    0.025990
AB   0.031734
R    0.036335
H    0.397595
2B   0.065891
3B   0.015658
HR   0.058105
RBI  0.045158
BB   0.127841
SO   0.040981
SB   0.052071
BA   0.102642
```

```
[156]: dtc2 = DecisionTreeClassifier(criterion='entropy', ccp_alpha=0.04)

dtc2.fit(X_train, y_train)

y_pred2 = dtc2.predict(X_test)

print(confusion_matrix(y_test, y_pred2))

print(classification_report(y_test, y_pred2))
```

```
[[50 11]
 [ 9 23]]
```

	precision	recall	f1-score	support
0	0.85	0.82	0.83	61
1	0.68	0.72	0.70	32
accuracy			0.78	93
macro avg	0.76	0.77	0.77	93
weighted avg	0.79	0.78	0.79	93

```
[157]: features2 = pd.DataFrame(dtc2.feature_importances_, index=X.columns)
features2
```

```
[157]:
```

	0
YRS	0.000000
G	0.000000
AB	0.000000
R	0.000000
H	0.837977
2B	0.000000
3B	0.000000
HR	0.000000
RBI	0.000000
BB	0.000000
SO	0.000000
SB	0.000000
BA	0.162023

9-Random_Forest

October 20, 2024

1 Random Forest

Random Forest is a supervised learning algorithm that is used for both classification and regression tasks. It builds multiple decision trees (hence the term “forest”) and combines their outputs to make a final prediction. It is an ensemble method, meaning it aggregates the predictions of many base models (in this case, decision trees) to create a more robust and accurate model.

Each tree in the Random Forest is trained on a different random subset of the data, and the splits within each tree are made using random subsets of features. This randomness helps the model avoid overfitting and improves generalization to unseen data.

- **For classification**, Random Forest uses a majority vote to determine the class label.
- **For regression**, it averages the predictions from all trees.

When to Use Random Forest? Random Forest is versatile and can be used in many situations, but it is especially effective when:

- **You need a powerful model that handles both classification and regression problems.**
- **The dataset has many features:** Random Forest can handle high-dimensional datasets and can even provide feature importance, helping identify the most influential variables.
- **The dataset is large:** Unlike a single decision tree, Random Forest can handle large datasets more effectively because it mitigates overfitting by averaging the predictions of many trees.
- **The data has missing values:** Random Forest can handle datasets with some missing values without requiring extensive preprocessing.
- **You need a model that can generalize well to unseen data:** Random Forest is often one of the best choices in terms of accuracy and robustness because it reduces overfitting.

How Does Random Forest Work?

1. **Create Multiple Subsets (Bootstrapping):** Random Forest randomly samples the training data with replacement to create multiple different training subsets (called bootstrap samples). Each decision tree is trained on one of these subsets.
2. **Random Feature Selection (Feature Bagging):** At each split in a tree, Random Forest only considers a random subset of the features rather than all of them. This ensures that individual trees are diverse and reduces the risk that some features dominate the model.

3. **Build Decision Trees:** Each bootstrap sample is used to build a decision tree. Each tree is grown to its maximum depth (typically unpruned), and the algorithm selects the best split from the subset of features at each node.
4. **Make Predictions:**
 - **For classification:** After training all the trees, Random Forest classifies new data points by making each tree “vote” on the predicted class. The final classification is the class with the most votes.
 - **For regression:** The prediction for each data point is the average of the predictions from all trees.
5. **Aggregate Predictions:** Random Forest aggregates the results of all the trees to make a final prediction. By combining the output of multiple decision trees, Random Forest can achieve better accuracy and stability than any single tree.

Who Should Use Random Forest? Random Forest is suitable for:

- **Data scientists and machine learning engineers** who need a strong baseline model.
- **Business analysts and decision-makers** who want to interpret feature importance or detect important decision-making factors.
- **Researchers and practitioners** in fields like finance, healthcare, and bioinformatics, where datasets can be large, complex, and noisy.
- **Beginners in machine learning** who want an easy-to-use and effective algorithm without requiring much fine-tuning.

Advantages of Random Forest:

- **Reduces overfitting:** By averaging the predictions of multiple decision trees, Random Forest mitigates overfitting that often occurs with a single decision tree.
- **Handles large datasets well:** Random Forest can scale to large datasets and many features effectively.
- **Works with both classification and regression problems:** It is versatile and can handle different types of machine learning problems.
- **Feature importance:** Random Forest can provide insights into which features are most important for making predictions.
- **Robust to outliers:** Since individual trees are built on random subsets of data, outliers have less impact on the final model.

Disadvantages of Random Forest:

- **Slower and more computationally expensive:** Since it trains multiple decision trees, Random Forest can be slower and require more memory compared to simpler models like decision trees or linear models.
- **Less interpretable:** While decision trees are easy to interpret, a Random Forest with hundreds of trees is more difficult to explain and understand.

- **Risk of overfitting in noisy data:** While it generally reduces overfitting, Random Forest can still overfit if there is too much noise in the data or if there are too many trees.

Real-World Applications:

- **Credit Scoring and Risk Assessment:** Banks and financial institutions use Random Forest to assess the risk of loan applicants defaulting based on their historical data.
- **Customer Churn Prediction:** Random Forest helps predict which customers are likely to leave a company based on their behavior and interaction history.
- **Fraud Detection:** Random Forest can identify fraudulent transactions by learning patterns from historical financial transaction data.
- **Healthcare and Medical Diagnosis:** Random Forest is used to classify patients based on symptoms, helping doctors make diagnostic decisions.
- **Stock Market Prediction:** Random Forest can be used to predict stock prices based on historical trends and market conditions.

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

df = pd.read_csv('500hits.csv', encoding="latin-1")
df.head()
```

```
[1]:
```

	PLAYER	YRS	G	AB	R	H	2B	3B	HR	RBI	BB	\
0	Ty Cobb	24	3035	11434	2246	4189	724	295	117	726	1249	
1	Stan Musial	22	3026	10972	1949	3630	725	177	475	1951	1599	
2	Tris Speaker	22	2789	10195	1882	3514	792	222	117	724	1381	
3	Derek Jeter	20	2747	11195	1923	3465	544	66	260	1311	1082	
4	Honus Wagner	21	2792	10430	1736	3430	640	252	101	0	963	

	S0	SB	CS	BA	HOF
0	357	892	178	0.366	1
1	696	78	31	0.331	1
2	220	432	129	0.345	1
3	1840	358	97	0.310	1
4	327	722	15	0.329	1

```
[2]: df = df.drop(columns=['PLAYER', 'CS'])

X = df.iloc[:, 0:13]
y = df.iloc[:, 13]
```

```
[3]: X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=17,
↳test_size=0.2)
```

```
rf = RandomForestClassifier()
rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)
y_pred
```

```
[3]: array([0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0,
          0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
          1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
          0, 0, 1, 0, 0], dtype=int64)
```

```
[4]: rf.score(X_test, y_test)
```

```
[4]: 0.8279569892473119
```

```
[5]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.90	0.87	61
1	0.79	0.69	0.73	32
accuracy			0.83	93
macro avg	0.82	0.79	0.80	93
weighted avg	0.83	0.83	0.82	93

```
[6]: features = pd.DataFrame(rf.feature_importances_, index=X.columns)
features
```

```
[6]:      0
YRS  0.028679
G    0.082555
AB   0.081906
R    0.117477
H    0.129573
2B   0.061198
3B   0.047405
HR   0.055006
RBI  0.109791
BB   0.048961
SO   0.041604
SB   0.046483
BA   0.149363
```

```
[7]: rf2 = RandomForestClassifier(
      n_estimators=1000,
```

```

        criterion='entropy',
        min_samples_split=10,
        max_depth=14, random_state=42
    )

    rf2.fit(X_train, y_train)

```

```
[7]: RandomForestClassifier(criterion='entropy', max_depth=14, min_samples_split=10,
                             n_estimators=1000, random_state=42)
```

```
[8]: rf2.score(X_test, y_test)
```

```
[8]: 0.8494623655913979
```

```
[9]: y_pred2 = rf2.predict(X_test)
      print(classification_report(y_test, y_pred2))
```

	precision	recall	f1-score	support
0	0.85	0.93	0.89	61
1	0.85	0.69	0.76	32
accuracy			0.85	93
macro avg	0.85	0.81	0.82	93
weighted avg	0.85	0.85	0.85	93

10-SVM

October 20, 2024

1 Support Vector Machine

Support Vector Machine (SVM) is a supervised machine learning algorithm used for both classification and regression tasks, though it is mostly used for classification. The goal of SVM is to find a hyperplane in an N-dimensional space (where N is the number of features) that distinctly classifies the data points. The best hyperplane is the one that maximizes the margin between the two classes, meaning it is as far as possible from the nearest data points of each class.

- **Support vectors** are the data points that are closest to the hyperplane. They are the most critical elements of the dataset because they define the decision boundary.
- **Margin** is the distance between the hyperplane and the nearest data points from either class. SVM seeks to maximize this margin to improve model generalization.

When to Use Support Vector Machine? SVM is particularly useful in scenarios where:

- **You have small or medium-sized datasets:** SVM works best on smaller datasets as it can be computationally expensive for large datasets.
- **The data is linearly separable:** SVM is most effective when the data can be clearly separated into two distinct classes by a straight line (in 2D) or a hyperplane (in higher dimensions).
- **You need a model that can handle high-dimensional data:** SVM is capable of performing well even when the number of features exceeds the number of samples, making it suitable for problems like text classification.
- **The classes are well separated:** SVM works well when there is a clear margin between the classes.
- **You need to avoid overfitting:** SVM focuses on maximizing the margin between classes, which often helps avoid overfitting in classification tasks.

How Does Support Vector Machine Work?

1. **Hyperplane Calculation:** In a binary classification problem, the goal of SVM is to find a hyperplane that best separates the data into two classes. In two-dimensional space, this hyperplane is a line, but in higher-dimensional space, it becomes a plane or hyperplane.
2. **Maximizing the Margin:** SVM tries to find the hyperplane that maximizes the margin between the two classes. The margin is the distance between the hyperplane and the closest

data points from each class (called support vectors). The larger the margin, the better the generalization ability of the classifier.

3. Linear vs. Non-Linear Classification:

- **Linear SVM:** If the data is linearly separable, SVM will find the optimal hyperplane that separates the two classes.
- **Non-Linear SVM (Kernel SVM):** In many real-world problems, the data is not linearly separable. SVM uses a technique called the kernel trick to map the data into a higher-dimensional space where a linear separation becomes possible. Some commonly used kernels include:
 - **Polynomial kernel:** For capturing polynomial relationships between features.
 - **Radial Basis Function (RBF) kernel (Gaussian kernel):** Useful for complex, non-linear relationships.

4. Soft Margin vs. Hard Margin:

- **Hard margin SVM:** Enforces that no data points from either class are allowed to be within the margin or misclassified. This is rarely used because real-world data often contains noise and overlaps.
- **Soft margin SVM:** Allows some misclassification or violation of the margin (controlled by a parameter C) to account for noisy or overlapping data.

5. Classification and Prediction:

Once the SVM model has found the optimal hyperplane, new data points are classified by determining on which side of the hyperplane they fall.

Who Should Use Support Vector Machine? SVM is a good choice for:

- **Researchers and practitioners** in fields like bioinformatics, image recognition, and text categorization, where datasets might be small but have many features.
- **Data scientists** looking for a robust classifier that works well in high-dimensional spaces.
- **Machine learning engineers** dealing with classification problems where the data is non-linearly separable but a strong decision boundary is needed.
- **Business analysts** who need an accurate and interpretable model when the classes are well-separated.###

Advantages of Support Vector Machine:

- **Effective in high-dimensional spaces:** SVM works well when the number of features is large, even more than the number of samples.
- **Robust to overfitting:** By maximizing the margin, SVM often generalizes well to unseen data, especially with appropriate tuning of the C parameter.
- **Flexibility with non-linear data:** With the use of kernels, SVM can effectively classify data that is not linearly separable.
- **Memory efficient:** SVM uses only a subset of the training points (the support vectors), making it efficient in memory usage.

Disadvantages of Support Vector Machine:

- **Computationally expensive:** SVM can be slow and memory-intensive, especially for large datasets.
- **Hard to interpret:** SVMs are not as easy to interpret as models like decision trees or logistic regression.
- **Requires careful tuning of parameters:** The performance of SVM depends on the choice of kernel, regularization parameter C, and other hyperparameters. Poor tuning can lead to suboptimal results.
- **Not suitable for large datasets:** SVM can struggle with very large datasets because of its computational complexity.

Real-World Applications:

- **Image Recognition:** SVM is widely used in image classification tasks, such as handwriting recognition or face detection.
- **Bioinformatics:** SVMs are used to classify proteins and genes, making them useful in genomics and biological sequence analysis.
- **Text Classification:** SVM is commonly used in text categorization (e.g., spam detection) because of its ability to handle high-dimensional data.
- **Medical Diagnosis:** SVM helps classify diseases based on symptoms or medical images, such as classifying tumors as malignant or benign.
- **Stock Market Prediction:** SVM is used to predict stock market trends by classifying historical data into different categories.

```
[21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

mean1 = 55
std_dev1 = 10
num_samples = 500

column1_numbers = np.random.normal(mean1, std_dev1, num_samples)
column1_numbers = np.clip(column1_numbers, 30, 120)
column1_numbers = np.round(column1_numbers).astype(int)

mean2 = 18
std_dev2 = 3

column2_numbers = np.random.normal(mean2, std_dev2, num_samples)
column2_numbers = np.clip(column2_numbers, 12, 26)
```



```

column2_numbers = np.round(column2_numbers).astype(int)

column3_numbers = np.random.randint(2, size=num_samples)
column3_numbers[column1_numbers > mean1] =1

data = {"Miles_Per_week": column1_numbers,
        "Farthest_run": column2_numbers,
        "Qualified_Boston_Marathon": column3_numbers}

df = pd.DataFrame(data)

df

```

```

[21]:      Miles_Per_week  Farthest_run  Qualified_Boston_Marathon
0              53             16                1
1              60             20                1
2              44             12                1
3              65             18                1
4              70             21                1
..           ...             ...                ...
495             58             17                1
496             55             12                0
497             43             18                1
498             73             22                1
499             47             19                0


```

[500 rows x 3 columns]

```

[22]: plt.figure(figsize=(10,6))
plt.scatter(df['Miles_Per_week'], df['Farthest_run'],
           ↪c=df['Qualified_Boston_Marathon'], cmap='jet_r')
plt.xlabel("Miles Per Week")
plt.ylabel("Farthest Run")
plt.title("Scatter Plot")
plt.colorbar(label="Qualified Boston Marathon")
plt.show()

```



```
[23]: X = df.iloc[:,0:2]
y = df.iloc[:,2]

X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=33,
↳test_size=0.2)
model = SVC()

model.fit(X_train, y_train)
print(model.score(X_test, y_test))
```

0.76

```
[24]: model_reg0 = SVC(C=0.1)
model_reg0.fit(X_train, y_train)
print(model_reg0.score(X_test, y_test))
```

0.76

```
[25]: model_reg1 = SVC(C=1)
model_reg1.fit(X_train, y_train)
print(model_reg1.score(X_test, y_test))
```

0.76

```
[26]: model_reg2 = SVC(C=1000)
      model_reg2.fit(X_train, y_train)
      print(model_reg2.score(X_test, y_test))
```

0.73

1.0.1 Gamma

```
[27]: model_gamma0 = SVC(gamma=0.1)
      model_gamma0.fit(X_train, y_train)
      model_gamma0.score(X_test, y_test)
```

[27]: 0.71

```
[28]: model_gamma1 = SVC(gamma=1)
      model_gamma1.fit(X_train, y_train)
      model_gamma1.score(X_test, y_test)
```

[28]: 0.74

```
[29]: model_gamma2 = SVC(gamma=1000)
      model_gamma2.fit(X_train, y_train)
      model_gamma2.score(X_test, y_test)
```

[29]: 0.74

1.0.2 Kernel

```
[30]: model_kernel0 = SVC(kernel='linear')
      model_kernel0.fit(X_train, y_train)
      model_kernel0.score(X_test, y_test)
```

[30]: 0.74

```
[31]: model_kernel0 = SVC(kernel='rbf')
      model_kernel0.fit(X_train, y_train)
      model_kernel0.score(X_test, y_test)
```

[31]: 0.76

11-Naive_Bayes

October 20, 2024

1 Naive Bayes Classifier

The **Naive Bayes classifier** is a family of probabilistic classifiers that rely on Bayes' Theorem and the assumption of feature independence. The classifier is called naive because it assumes that all the features in a dataset are independent of each other, which is rarely true in real-world scenarios, but it simplifies the computation.

Naive Bayes works by calculating the probabilities for each class and then predicting the class with the highest probability based on the input features. It is a powerful and simple algorithm for classification tasks, especially in domains like text classification (e.g., spam detection, sentiment analysis).

When to Use Naive Bayes Classifier? Naive Bayes is particularly useful in the following scenarios:

- **When working with high-dimensional datasets:** It performs well with many features, such as in text classification where each word can be a feature.
- **When the assumption of conditional independence holds:** Even when this assumption is not completely true, Naive Bayes often performs surprisingly well in practice.
- **For real-time predictions:** Naive Bayes is computationally efficient and can handle large datasets quickly, making it suitable for real-time applications.
- **Text classification, sentiment analysis, and spam filtering:** These tasks involve high-dimensional and sparse data, where Naive Bayes shines.
- **Categorical input features:** The algorithm handles categorical variables well, especially in its Multinomial Naive Bayes variant.

Who Should Use Naive Bayes Classifier? Naive Bayes is an excellent choice for:

- **Data scientists and machine learning practitioners** who need a fast, simple, and interpretable model for classification tasks.
- **Business analysts** looking for text-based solutions, such as customer sentiment analysis, email spam filtering, or document categorization.
- **Researchers and developers** working with large-scale text processing, like in Natural Language Processing (NLP).
- **Beginners in machine learning**, as Naive Bayes is easy to implement and understand, while still being effective in many real-world scenarios.

Advantages of Naive Bayes Classifier:

- **Fast and efficient:** Naive Bayes is computationally efficient and works well with large datasets.
- **Handles high-dimensional data:** Especially useful for text classification problems where the data is sparse and high-dimensional.
- **Simple to implement:** The algorithm is simple to code and interpret, making it a good starting point for beginners.
- **Requires less training data:** It performs well with relatively small datasets because it estimates fewer parameters.
- **Performs well in practice:** Even though the independence assumption rarely holds true, Naive Bayes often works surprisingly well.

1.0.1 Disadvantages of Naive Bayes Classifier:

- **Strong independence assumption:** The assumption that features are conditionally independent given the class can lead to suboptimal results if this condition is not met in the data.
- **Zero probability problem:** If a feature value that wasn't seen in the training data appears in the test set, the model will assign a zero probability to that class. This can be mitigated using techniques like Laplace smoothing.
- **Not suitable for continuous data without modifications:** The basic form of Naive Bayes does not handle continuous data well unless it is adapted (e.g., Gaussian Naive Bayes).
- **Not ideal for highly correlated features:** Naive Bayes struggles when features are correlated since the independence assumption is violated.

Real-World Applications of Naive Bayes Classifier:

- **Email Spam Detection:** Classifies emails as spam or not based on the occurrence of certain words.
- **Text Classification:** Categorizes documents or news articles into different topics based on word occurrences (Multinomial Naive Bayes).
- **Sentiment Analysis:** Determines whether a review or a tweet is positive or negative by analyzing word frequencies (commonly used in NLP).
- **Medical Diagnosis:** Predicts diseases based on symptoms by calculating the probability of each condition.
- **Recommender Systems:** Helps in filtering and recommending content based on a user's previous interactions and preferences.

```
[55]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
```

```

from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV

num_items = 41
data = {
    'City Population': np.random.randint(10000, 1000000, num_items),
    'Continent': np.random.choice(['Asia', 'Europe', 'North America', 'South_
    ↪America'], num_items),
    'Venue Capacity': np.random.randint(500, 20000, num_items),
    'Day Of Week': np.random.choice(['Monday', 'Tuesday', 'Wednesday',
    ↪Thursday', 'Friday', 'Saturday', 'Sunday'], num_items),
    'Multiple Concerts': np.random.randint(0, 2, num_items),
    'Sold Out': np.random.randint(0, 2, num_items)
}

df = pd.DataFrame(data)
#df = pd.read_csv('pear_jam_tour2.csv', encoding="unicode_escape")
df

```

```

[55]:
   City Population  Continent  Venue Capacity  Day Of Week  \
0          491817  North America          16665    Thursday
1          701437         Europe           9246      Monday
2          166930         Europe           8403    Thursday
3          588655  South America           4672    Saturday
4          388841  South America           1405      Friday
5          882813         Asia           18118      Sunday
6          158157         Asia            1666  Wednesday
7          117802         Asia          14382      Sunday
8           44163  North America            2502    Saturday
9          893911  North America          17897      Friday
10         981604         Europe          15859    Tuesday
11         780283  North America          19257  Wednesday
12         265129         Asia           4524    Thursday
13         131209         Asia          14857    Thursday
14         405167         Europe           9712      Monday
15          84701  North America            2482    Thursday
16         822257  North America           5842      Sunday
17         424221  South America          17013    Saturday
18         305672         Asia           2628      Sunday
19         720117  North America           3999    Thursday
20         285205  North America           6290    Saturday
21         121093  North America          10833    Thursday
22         240183         Asia           8535    Saturday
23         368908  North America          19451    Tuesday
24         989683  North America           8082    Saturday
25         743758  South America          10294      Monday

```

26	951668	Asia	15488	Friday
27	356715	Europe	1282	Wednesday
28	428998	South America	17548	Friday
29	490859	Europe	9096	Friday
30	815163	Asia	12759	Tuesday
31	568414	North America	3550	Friday
32	569051	South America	5245	Sunday
33	327147	Asia	4974	Wednesday
34	703689	North America	19689	Monday
35	368834	Europe	15419	Monday
36	259567	North America	16369	Tuesday
37	654837	North America	19461	Friday
38	560703	North America	19339	Thursday
39	466105	South America	12232	Tuesday
40	167668	South America	2890	Thursday

	Multiple Concerts	Sold Out
0	0	0
1	0	0
2	0	0
3	0	1
4	1	1
5	1	1
6	1	0
7	0	0
8	0	0
9	1	1
10	1	1
11	0	1
12	0	1
13	1	1
14	0	0
15	0	0
16	0	1
17	1	0
18	0	1
19	0	0
20	1	0
21	1	1
22	1	1
23	1	1
24	0	0
25	0	0
26	0	1
27	0	0
28	0	0
29	0	0

30	1	1
31	0	1
32	1	1
33	0	0
34	0	0
35	0	0
36	0	0
37	1	1
38	1	0
39	0	1
40	0	1

```
[56]: df2 = pd.get_dummies(df[['Continent', 'Day Of Week']])

df2
```

```
[56]:
```

	Continent_Asia	Continent_Europe	Continent_North America	\
0	False	False	True	
1	False	True	False	
2	False	True	False	
3	False	False	False	
4	False	False	False	
5	True	False	False	
6	True	False	False	
7	True	False	False	
8	False	False	True	
9	False	False	True	
10	False	True	False	
11	False	False	True	
12	True	False	False	
13	True	False	False	
14	False	True	False	
15	False	False	True	
16	False	False	True	
17	False	False	False	
18	True	False	False	
19	False	False	True	
20	False	False	True	
21	False	False	True	
22	True	False	False	
23	False	False	True	
24	False	False	True	
25	False	False	False	
26	True	False	False	
27	False	True	False	
28	False	False	False	
29	False	True	False	

30	True	False	False
31	False	False	True
32	False	False	False
33	True	False	False
34	False	False	True
35	False	True	False
36	False	False	True
37	False	False	True
38	False	False	True
39	False	False	False
40	False	False	False

	Continent_South America	Day Of Week_Friday	Day Of Week_Monday \
0	False	False	False
1	False	False	True
2	False	False	False
3	True	False	False
4	True	True	False
5	False	False	False
6	False	False	False
7	False	False	False
8	False	False	False
9	False	True	False
10	False	False	False
11	False	False	False
12	False	False	False
13	False	False	False
14	False	False	True
15	False	False	False
16	False	False	False
17	True	False	False
18	False	False	False
19	False	False	False
20	False	False	False
21	False	False	False
22	False	False	False
23	False	False	False
24	False	False	False
25	True	False	True
26	False	True	False
27	False	False	False
28	True	True	False
29	False	True	False
30	False	False	False
31	False	True	False
32	True	False	False
33	False	False	False

34	False	False	True
35	False	False	True
36	False	False	False
37	False	True	False
38	False	False	False
39	True	False	False
40	True	False	False

	Day Of Week_Saturday	Day Of Week_Sunday	Day Of Week_Thursday \
0	False	False	True
1	False	False	False
2	False	False	True
3	True	False	False
4	False	False	False
5	False	True	False
6	False	False	False
7	False	True	False
8	True	False	False
9	False	False	False
10	False	False	False
11	False	False	False
12	False	False	True
13	False	False	True
14	False	False	False
15	False	False	True
16	False	True	False
17	True	False	False
18	False	True	False
19	False	False	True
20	True	False	False
21	False	False	True
22	True	False	False
23	False	False	False
24	True	False	False
25	False	False	False
26	False	False	False
27	False	False	False
28	False	False	False
29	False	False	False
30	False	False	False
31	False	False	False
32	False	True	False
33	False	False	False
34	False	False	False
35	False	False	False
36	False	False	False
37	False	False	False

38	False	False	True
39	False	False	False
40	False	False	True

	Day Of Week_Tuesday	Day Of Week_Wednesday
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	True
7	False	False
8	False	False
9	False	False
10	True	False
11	False	True
12	False	False
13	False	False
14	False	False
15	False	False
16	False	False
17	False	False
18	False	False
19	False	False
20	False	False
21	False	False
22	False	False
23	True	False
24	False	False
25	False	False
26	False	False
27	False	True
28	False	False
29	False	False
30	True	False
31	False	False
32	False	False
33	False	True
34	False	False
35	False	False
36	True	False
37	False	False
38	False	False
39	True	False
40	False	False

```
[57]: df3 = pd.concat([df, df2], axis=1)
```

```
df3
```

```
[57]:
```

	City	Population	Continent	Venue	Capacity	Day Of Week	\
0		491817	North America		16665	Thursday	
1		701437	Europe		9246	Monday	
2		166930	Europe		8403	Thursday	
3		588655	South America		4672	Saturday	
4		388841	South America		1405	Friday	
5		882813	Asia		18118	Sunday	
6		158157	Asia		1666	Wednesday	
7		117802	Asia		14382	Sunday	
8		44163	North America		2502	Saturday	
9		893911	North America		17897	Friday	
10		981604	Europe		15859	Tuesday	
11		780283	North America		19257	Wednesday	
12		265129	Asia		4524	Thursday	
13		131209	Asia		14857	Thursday	
14		405167	Europe		9712	Monday	
15		84701	North America		2482	Thursday	
16		822257	North America		5842	Sunday	
17		424221	South America		17013	Saturday	
18		305672	Asia		2628	Sunday	
19		720117	North America		3999	Thursday	
20		285205	North America		6290	Saturday	
21		121093	North America		10833	Thursday	
22		240183	Asia		8535	Saturday	
23		368908	North America		19451	Tuesday	
24		989683	North America		8082	Saturday	
25		743758	South America		10294	Monday	
26		951668	Asia		15488	Friday	
27		356715	Europe		1282	Wednesday	
28		428998	South America		17548	Friday	
29		490859	Europe		9096	Friday	
30		815163	Asia		12759	Tuesday	
31		568414	North America		3550	Friday	
32		569051	South America		5245	Sunday	
33		327147	Asia		4974	Wednesday	
34		703689	North America		19689	Monday	
35		368834	Europe		15419	Monday	
36		259567	North America		16369	Tuesday	
37		654837	North America		19461	Friday	
38		560703	North America		19339	Thursday	
39		466105	South America		12232	Tuesday	
40		167668	South America		2890	Thursday	

	Multiple Concerts	Sold Out	Continent_Asia	Continent_Europe	\
0	0	0	False	False	
1	0	0	False	True	
2	0	0	False	True	
3	0	1	False	False	
4	1	1	False	False	
5	1	1	True	False	
6	1	0	True	False	
7	0	0	True	False	
8	0	0	False	False	
9	1	1	False	False	
10	1	1	False	True	
11	0	1	False	False	
12	0	1	True	False	
13	1	1	True	False	
14	0	0	False	True	
15	0	0	False	False	
16	0	1	False	False	
17	1	0	False	False	
18	0	1	True	False	
19	0	0	False	False	
20	1	0	False	False	
21	1	1	False	False	
22	1	1	True	False	
23	1	1	False	False	
24	0	0	False	False	
25	0	0	False	False	
26	0	1	True	False	
27	0	0	False	True	
28	0	0	False	False	
29	0	0	False	True	
30	1	1	True	False	
31	0	1	False	False	
32	1	1	False	False	
33	0	0	True	False	
34	0	0	False	False	
35	0	0	False	True	
36	0	0	False	False	
37	1	1	False	False	
38	1	0	False	False	
39	0	1	False	False	
40	0	1	False	False	

	Continent_North America	Continent_South America	Day Of Week_Friday	\
0	True	False	False	
1	False	False	False	
2	False	False	False	

3	False	True	False
4	False	True	True
5	False	False	False
6	False	False	False
7	False	False	False
8	True	False	False
9	True	False	True
10	False	False	False
11	True	False	False
12	False	False	False
13	False	False	False
14	False	False	False
15	True	False	False
16	True	False	False
17	False	True	False
18	False	False	False
19	True	False	False
20	True	False	False
21	True	False	False
22	False	False	False
23	True	False	False
24	True	False	False
25	False	True	False
26	False	False	True
27	False	False	False
28	False	True	True
29	False	False	True
30	False	False	False
31	True	False	True
32	False	True	False
33	False	False	False
34	True	False	False
35	False	False	False
36	True	False	False
37	True	False	True
38	True	False	False
39	False	True	False
40	False	True	False

	Day Of Week_Monday	Day Of Week_Saturday	Day Of Week_Sunday \
0	False	False	False
1	True	False	False
2	False	False	False
3	False	True	False
4	False	False	False
5	False	False	True
6	False	False	False

7	False	False	True
8	False	True	False
9	False	False	False
10	False	False	False
11	False	False	False
12	False	False	False
13	False	False	False
14	True	False	False
15	False	False	False
16	False	False	True
17	False	True	False
18	False	False	True
19	False	False	False
20	False	True	False
21	False	False	False
22	False	True	False
23	False	False	False
24	False	True	False
25	True	False	False
26	False	False	False
27	False	False	False
28	False	False	False
29	False	False	False
30	False	False	False
31	False	False	False
32	False	False	True
33	False	False	False
34	True	False	False
35	True	False	False
36	False	False	False
37	False	False	False
38	False	False	False
39	False	False	False
40	False	False	False

	Day Of Week_Thursday	Day Of Week_Tuesday	Day Of Week_Wednesday
0	True	False	False
1	False	False	False
2	True	False	False
3	False	False	False
4	False	False	False
5	False	False	False
6	False	False	True
7	False	False	False
8	False	False	False
9	False	False	False
10	False	True	False

11	False	False	True
12	True	False	False
13	True	False	False
14	False	False	False
15	True	False	False
16	False	False	False
17	False	False	False
18	False	False	False
19	True	False	False
20	False	False	False
21	True	False	False
22	False	False	False
23	False	True	False
24	False	False	False
25	False	False	False
26	False	False	False
27	False	False	True
28	False	False	False
29	False	False	False
30	False	True	False
31	False	False	False
32	False	False	False
33	False	False	True
34	False	False	False
35	False	False	False
36	False	True	False
37	False	False	False
38	True	False	False
39	False	True	False
40	True	False	False

```
[58]: df4 = df3.drop(columns=['Continent', 'Day Of Week'], axis=1)
df4
```

```
[58]:
```

	City	Population	Venue Capacity	Multiple Concerts	Sold Out	\
0		491817	16665	0	0	
1		701437	9246	0	0	
2		166930	8403	0	0	
3		588655	4672	0	1	
4		388841	1405	1	1	
5		882813	18118	1	1	
6		158157	1666	1	0	
7		117802	14382	0	0	
8		44163	2502	0	0	
9		893911	17897	1	1	
10		981604	15859	1	1	
11		780283	19257	0	1	

12	265129	4524	0	1
13	131209	14857	1	1
14	405167	9712	0	0
15	84701	2482	0	0
16	822257	5842	0	1
17	424221	17013	1	0
18	305672	2628	0	1
19	720117	3999	0	0
20	285205	6290	1	0
21	121093	10833	1	1
22	240183	8535	1	1
23	368908	19451	1	1
24	989683	8082	0	0
25	743758	10294	0	0
26	951668	15488	0	1
27	356715	1282	0	0
28	428998	17548	0	0
29	490859	9096	0	0
30	815163	12759	1	1
31	568414	3550	0	1
32	569051	5245	1	1
33	327147	4974	0	0
34	703689	19689	0	0
35	368834	15419	0	0
36	259567	16369	0	0
37	654837	19461	1	1
38	560703	19339	1	0
39	466105	12232	0	1
40	167668	2890	0	1

	Continent_Asia	Continent_Europe	Continent_North America	\
0	False	False	True	
1	False	True	False	
2	False	True	False	
3	False	False	False	
4	False	False	False	
5	True	False	False	
6	True	False	False	
7	True	False	False	
8	False	False	True	
9	False	False	True	
10	False	True	False	
11	False	False	True	
12	True	False	False	
13	True	False	False	
14	False	True	False	
15	False	False	True	

16	False	False	True
17	False	False	False
18	True	False	False
19	False	False	True
20	False	False	True
21	False	False	True
22	True	False	False
23	False	False	True
24	False	False	True
25	False	False	False
26	True	False	False
27	False	True	False
28	False	False	False
29	False	True	False
30	True	False	False
31	False	False	True
32	False	False	False
33	True	False	False
34	False	False	True
35	False	True	False
36	False	False	True
37	False	False	True
38	False	False	True
39	False	False	False
40	False	False	False

	Continent_South America	Day Of Week_Friday	Day Of Week_Monday \
0	False	False	False
1	False	False	True
2	False	False	False
3	True	False	False
4	True	True	False
5	False	False	False
6	False	False	False
7	False	False	False
8	False	False	False
9	False	True	False
10	False	False	False
11	False	False	False
12	False	False	False
13	False	False	False
14	False	False	True
15	False	False	False
16	False	False	False
17	True	False	False
18	False	False	False
19	False	False	False

20	False	False	False
21	False	False	False
22	False	False	False
23	False	False	False
24	False	False	False
25	True	False	True
26	False	True	False
27	False	False	False
28	True	True	False
29	False	True	False
30	False	False	False
31	False	True	False
32	True	False	False
33	False	False	False
34	False	False	True
35	False	False	True
36	False	False	False
37	False	True	False
38	False	False	False
39	True	False	False
40	True	False	False

	Day Of Week_Saturday	Day Of Week_Sunday	Day Of Week_Thursday \
0	False	False	True
1	False	False	False
2	False	False	True
3	True	False	False
4	False	False	False
5	False	True	False
6	False	False	False
7	False	True	False
8	True	False	False
9	False	False	False
10	False	False	False
11	False	False	False
12	False	False	True
13	False	False	True
14	False	False	False
15	False	False	True
16	False	True	False
17	True	False	False
18	False	True	False
19	False	False	True
20	True	False	False
21	False	False	True
22	True	False	False
23	False	False	False

24	True	False	False
25	False	False	False
26	False	False	False
27	False	False	False
28	False	False	False
29	False	False	False
30	False	False	False
31	False	False	False
32	False	True	False
33	False	False	False
34	False	False	False
35	False	False	False
36	False	False	False
37	False	False	False
38	False	False	True
39	False	False	False
40	False	False	True

	Day Of Week_Tuesday	Day Of Week_Wednesday
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	True
7	False	False
8	False	False
9	False	False
10	True	False
11	False	True
12	False	False
13	False	False
14	False	False
15	False	False
16	False	False
17	False	False
18	False	False
19	False	False
20	False	False
21	False	False
22	False	False
23	True	False
24	False	False
25	False	False
26	False	False
27	False	True

28	False	False
29	False	False
30	True	False
31	False	False
32	False	False
33	False	True
34	False	False
35	False	False
36	True	False
37	False	False
38	False	False
39	True	False
40	False	False

```
[59]: X = df4.drop(columns=['Sold Out'], axis=1)
y = df4['Sold Out']

X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=33,
↳test_size=0.2)
gnb = GaussianNB()

gnb.fit(X_train, y_train)
y_pred = gnb.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.80	0.80	5
1	0.75	0.75	0.75	4
accuracy			0.78	9
macro avg	0.78	0.78	0.78	9
weighted avg	0.78	0.78	0.78	9

```
[60]: gnb.score(X_train, y_train)
```

```
[60]: 0.5625
```

```
[61]: gnb.score(X_test, y_test)
```

```
[61]: 0.7777777777777778
```

1.0.2 Add in parameter

```
[66]: param_grid = {  
        'var_smoothing': [0.00000001, 0.000000001, 0.00000001]  
    }  
  
    grid_search = GridSearchCV(gnb,param_grid, cv=5, scoring='accuracy', n_jobs=-1)  
    grid_search.fit(X_train, y_train)
```

```
[66]: GridSearchCV(cv=5, estimator=GaussianNB(), n_jobs=-1,  
        param_grid={'var_smoothing': [1e-08, 1e-09, 1e-08]},  
        scoring='accuracy')
```

```
[67]: grid_search.best_params_
```

```
[67]: {'var_smoothing': 1e-08}
```

```
[68]: grid_search.best_score_
```

```
[68]: 0.4095238095238095
```

12-Logistical_Regression

October 20, 2024

1 Logistic Regression

Logistic Regression is a supervised learning algorithm used for classification tasks, not regression, despite its name. It estimates the probability that a given input point belongs to a certain class by applying a logistic (sigmoid) function to a linear combination of input features. The output is a probability between 0 and 1, making it ideal for binary classification problems (two possible outcomes, like yes/no, spam/not spam).

Logistic regression predicts a binary outcome based on one or more predictor variables (features) and is one of the foundational techniques in machine learning and statistics.

When to Use Logistic Regression? Logistic Regression is commonly used when:

- **The dependent variable is binary:** For example, a classification problem where the output is either “1” or “0” (e.g., true/false, success/failure).
- **You need probabilistic interpretations:** Logistic regression outputs probabilities that a data point belongs to a certain class, which makes it useful in areas where you need both the prediction and the confidence of that prediction.
- **When interpretability is important:** It provides clear and interpretable coefficients, which can explain the impact of each feature on the prediction.
- **The relationship between features and the target is linear in the log-odds:** While logistic regression assumes a linear relationship between the features and the log-odds of the outcome, it can still handle non-linear relationships by transforming the input features (e.g., polynomial or interaction terms).

Typical use cases include:

- **Medical Diagnosis:** Predicting whether a patient has a disease or not.
- **Email Classification:** Determining whether an email is spam or not.
- **Customer Churn:** Predicting whether a customer will leave a service or remain.
- **Credit Scoring:** Estimating the probability of a customer defaulting on a loan.

Who Should Use Logistic Regression? Logistic Regression is ideal for:

- **Data scientists and machine learning practitioners** who need a simple, interpretable model for binary classification tasks.

- **Business analysts and researchers** looking for an explainable model where they can interpret the effect of each feature on the outcome.
- **Medical professionals** making decisions based on classification problems like disease detection.
- **Beginner machine learning students** since logistic regression is one of the simplest and easiest algorithms to learn and apply.

Advantages of Logistic Regression:

- **Simplicity and Interpretability:** It is easy to implement and interpret, making it a popular choice for binary classification problems.
- **Probabilistic Output:** Logistic regression provides probabilities for class membership, which can be useful for understanding the confidence of predictions.
- **Fast and Efficient:** It performs well on relatively small datasets and is computationally efficient.
- **No Feature Scaling Required:** While performance can be improved by scaling features, logistic regression does not require it.
- **Linear Decision Boundary:** The decision boundary is linear, which makes it suitable for linearly separable datasets.

Disadvantages of Logistic Regression:

- **Limited to Linear Boundaries:** Logistic regression assumes a linear relationship between the features and the log-odds, making it unsuitable for complex or highly non-linear problems.
- **Assumes Independence of Features:** Like Naive Bayes, logistic regression assumes that the features are independent of each other, which may not hold true in many cases.
- **Cannot Handle Multiclass Problems Directly:** Logistic regression is a binary classifier by default, but techniques like One-vs-Rest or Softmax Regression can extend it to multiclass problems.
- **Sensitive to Outliers:** Logistic regression can be sensitive to outliers, which can skew the decision boundary.

Real-World Applications of Logistic Regression:

1. **Credit Scoring:** Predicting the probability of a customer defaulting on a loan.
2. **Medical Diagnosis:** Predicting the likelihood of a disease based on symptoms.
3. **Customer Retention:** Estimating the probability that a customer will churn or stay with the company.
4. **Marketing Campaigns:** Predicting whether a customer will respond to an advertisement or not.
5. **Email Classification:** Classifying emails as spam or non-spam.


```
[1]: import pandas as pd

d = {
    'miles_per_week': [
        ↪[37,39,46,51,88,17,18,20,21,22,23,24,25,27,28,29,30,31,32,33,34,38,40,42,57,68,35,36,41,43,
        'completed_50m_ultra': [
        ↪['no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no','no',
        ]
    }

df = pd.DataFrame(data=d)
df
```

```
[1]:      miles_per_week  completed_50m_ultra
0              37              no
1              39              no
2              46              no
3              51              no
4              88              no
..           ...           ...
96              67              yes
97              74              yes
98              79              yes
99              90              yes
100             112              yes
```

[101 rows x 2 columns]

```
[2]: from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import OrdinalEncoder
from sklearn.linear_model import LogisticRegression
from matplotlib import pyplot as plt

import seaborn as sns

finished_race = ['no', 'yes']
enc = OrdinalEncoder(categories=[finished_race])
df['completed_50m_ultra'] = enc.fit_transform(df[['completed_50m_ultra']])

df
```

```
[2]:      miles_per_week  completed_50m_ultra
0              37              0.0
1              39              0.0
2              46              0.0
3              51              0.0
4              88              0.0
```

```


..          ...          ...
96          67          1.0
97          74          1.0
98          79          1.0
99          90          1.0
100         112          1.0

```

```
[101 rows x 2 columns]
```


```
[3]: plt.scatter(df['miles_per_week'], df['completed_50m_ultra'])
```

```
[3]: <matplotlib.collections.PathCollection at 0x1f0ce21e410>
```



```
[4]: sns.countplot(x='completed_50m_ultra', data=df)
```

```
[4]: <Axes: xlabel='completed_50m_ultra', ylabel='count'>
```



```
[5]: X = df.iloc[:,0:1]
y = df.iloc[:,1]

X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.8,
↳random_state=11)

model = LogisticRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
model.score(X_test, y_test)
```

```
[5]: 0.9047619047619048
```

```
[6]: print(confusion_matrix(y_test, y_pred))
```

```
[[ 5  1]
 [ 1 14]]
```

```
[7]: print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support
```

0.0	0.83	0.83	0.83	6
1.0	0.93	0.93	0.93	15
accuracy			0.90	21
macro avg	0.88	0.88	0.88	21
weighted avg	0.90	0.90	0.90	21

13-Machine_Learning_Pipeline

October 20, 2024

1 Machine Learning Pipeline

```
[96]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline, Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer

from sklearn.tree import DecisionTreeClassifier

import joblib
```

```
[97]: d1 = {
    'Social_media_followers' : [1000000, np.nan, 2000000, 1310000, 1700000, np.
    ↪nan, 4100000, 1600000, 2200000, 1000000],
    'Sold_out': [1,0,0,1,0,0,0,1,0,1]
}

d2 = {
    'Genre':['Rock', 'Metal', 'Bluegrass', 'Rock', np.nan, 'Rock', 'Rock',
    ↪np.nan, 'Bluegrass', 'Rock'],
    'Social_media_followers':[1000000, np.nan, 2000000, 1310000, 1700000,
    ↪np.nan, 4100000, 1600000, 2200000, 1000000],
    'Sold_out':[1,0,0,1,0,0,0,1,0,1]
}

df1 = pd.DataFrame(d1)
df1
```

```
[97]:   Social_media_followers  Sold_out
0          1000000.0          1
1              NaN          0
2          2000000.0          0
3          1310000.0          1
```

4	1700000.0	0
5	NaN	0
6	4100000.0	0
7	1600000.0	1
8	2200000.0	0
9	1000000.0	1

```
[98]: X1 = df1[['Social_media_followers']]
      y1 = df1[['Sold_out']]

      X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.3,
      ↪random_state=19)

      imputer = SimpleImputer(strategy='mean')
      lr = LogisticRegression()

      pipe1 = make_pipeline(imputer, lr)

      pipe1.fit(X1_train, y1_train)
```

c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1229: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
[98]: Pipeline(steps=[('simpleimputer', SimpleImputer()),
                       ('logisticregression', LogisticRegression())])
```

```
[99]: pipe1.score(X1_train, y1_train)
```

```
[99]: 1.0
```

```
[100]: pipe1.score(X1_test, y1_test)
```

```
[100]: 0.6666666666666666
```

```
[101]: pipe1.named_steps.simpleimputer.statistics_
```

```
[101]: array([2051666.66666667])
```

```
[102]: pipe1.named_steps.logisticregression.coef_
```

```
[102]: array([[ -9.72872687e-05]])
```

1.0.1 More Advance Pipeline

```
[103]: df = pd.DataFrame(data=d2)
df
```

```
[103]:
```

	Genre	Social_media_followers	Sold_out
0	Rock	1000000.0	1
1	Metal	NaN	0
2	Bluegrass	2000000.0	0
3	Rock	1310000.0	1
4	NaN	1700000.0	0
5	Rock	NaN	0
6	Rock	4100000.0	0
7	NaN	1600000.0	1
8	Bluegrass	2200000.0	0
9	Rock	1000000.0	1

```
[104]: X = df.iloc[:,0:2]
y = df.iloc[:,2]

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3,
↳random_state=17)

num_cols = ["Social_media_followers"]
cat_cols = ['Genre']

num_pipeline = Pipeline(
    steps = [
        ('impute', SimpleImputer(strategy='mean')),
        ('scale', StandardScaler())
    ]
)

cat_pipeline = Pipeline(steps=[
    ('impute', SimpleImputer(strategy='most_frequent')),
    ('one-hot-encoder', OneHotEncoder(handle_unknown='ignore',
↳sparse_output=False))
])
cat_pipeline
```

```
[104]: Pipeline(steps=[('impute', SimpleImputer(strategy='most_frequent')),
    ('one-hot-encoder',
    OneHotEncoder(handle_unknown='ignore', sparse_output=False))])
```

```
[105]: col_transformer = ColumnTransformer(transformers= [
    ('num_pipeline', num_pipeline, num_cols),
    ('cat_pipeline', cat_pipeline, cat_cols),
],
```

```
remainder='drop', n_jobs=-1
)
```

```
[106]: dtc = DecisionTreeClassifier()
pipefinal = make_pipeline(col_transformer, dtc)
pipefinal.fit(X_train, y_train)
```

```
[106]: Pipeline(steps=[('columntransformer',
                        ColumnTransformer(n_jobs=-1,
                                          transformers=[('num_pipeline',
                                                          Pipeline(steps=[('impute',
                                                                              SimpleImputer()),
                                                                              ('scale',
                                                                              StandardScaler()))],
                                                                              ['Social_media_followers']),
                                                          ('cat_pipeline',
                                                          Pipeline(steps=[('impute',
                                                                              SimpleImputer(strategy='most_frequent')),
                                                                              ('one-hot-
                                                                              encoder',
                                                                              OneHotEncoder(handle_unknown='ignore',
                                                                              sparse_output=False))]),
                                                                              ['Genre'])])),
                        ('decisiontreeclassifier', DecisionTreeClassifier())])
```

```
[107]: pipefinal.score(X_test, y_test)
```

```
[107]: 0.6666666666666666
```

1.1 How to save your pipeline

```
[108]: joblib.dump(pipefinal, 'pipe.joblib')
```

```
[108]: ['pipe.joblib']
```

```
[109]: pipefinal2 = joblib.load('pipe.joblib')
pipefinal2
```

```
[109]: Pipeline(steps=[('columntransformer',
                        ColumnTransformer(n_jobs=-1,
                                          transformers=[('num_pipeline',
                                                          Pipeline(steps=[('impute',
                                                                              SimpleImputer()),
                                                                              ('scale',
                                                                              StandardScaler()))],
                                                                              ['Social_media_followers']),
                                                          ('cat_pipeline',
```



```

SimpleImputer(strategy='most_frequent')),
                                                    Pipeline(steps=[('impute',
                                                                    ('one-hot-
encoder',
OneHotEncoder(handle_unknown='ignore',
               sparse_output=False)))]),
                                                    ['Genre']]])),
('decisiontreeclassifier', DecisionTreeClassifier())])

```

14-Cross_Validation

October 20, 2024

1 Cross Validation

Cross-Validation is a technique used in machine learning to assess the performance of a model by testing it on different subsets of the data. Instead of using just a single training and testing split, cross-validation helps evaluate the model's ability to generalize to unseen data by dividing the dataset into multiple parts (or folds) and training/testing the model on each part.

The most commonly used form of cross-validation is k-fold cross-validation, where the dataset is divided into k subsets or "folds." The model is trained on k-1 folds and tested on the remaining fold. This process is repeated k times, each time using a different fold as the test set and the rest as the training set. The final performance is averaged over all folds.

Cross-validation is important because it gives a better estimate of the model's performance by reducing the risk of overfitting or underfitting the data.

When to Use Cross-Validation? Cross-Validation should be used when:

- **You want to evaluate model performance:** It's used to estimate how well a model generalizes to an independent dataset.
- **You're working with limited data:** Cross-validation helps make the most out of the available data by using all of it for both training and testing across different iterations.
- **You're tuning hyperparameters:** It's often employed when trying to find the optimal hyperparameters of a model (e.g., in grid search).
- **To prevent overfitting:** When you want to avoid overfitting on a particular training set by providing a more robust evaluation of model performance.
- **In model selection:** It is used to compare different models and choose the best-performing one based on a more reliable performance metric.

Use cases include:

- **Machine learning competitions:** Such as Kaggle competitions, where high generalization performance is crucial.
- **Scientific research:** Where precise model evaluation is necessary to make reliable predictions.
- **Production-ready systems:** To validate that the model will work well in the real world, where unseen data is used.

1.0.1 Common Types of Cross-Validation:

1. k-Fold Cross-Validation:

- The most commonly used form, as described above. The number of folds, k , is usually 5 or 10.
- Advantages:
 - Provides a better estimation of model performance.
 - Each instance is used for both training and validation.
- Disadvantages:
 - More computationally expensive than simple train-test split, as the model is trained and evaluated k times.

2. Leave-One-Out Cross-Validation (LOOCV):

- A special case of k -fold where k equals the number of data points. That means for each iteration, only one instance is left out for testing, and the rest is used for training.
- Advantages:
 - No data point is wasted.
 - Provides an unbiased estimate of the generalization error.
- Disadvantages:
 - Computationally expensive for large datasets.
 - High variance, as training on almost the entire dataset may lead to overfitting.

3. Stratified k-Fold Cross-Validation:

- A variation of k -fold cross-validation where the data is split such that the proportion of classes in each fold is the same as the original dataset. This is especially useful in cases of imbalanced datasets.
- Advantages:
 - Better representation of the minority and majority classes in each fold.
 - Reduces bias in datasets with imbalanced class distributions.
- Disadvantages:
 - Still computationally intensive, similar to k -fold cross-validation.

4. Time Series Cross-Validation (Rolling or Expanding Window):

- In time series data, data points are not independent, so traditional k -fold cross-validation may not be suitable. Instead, the training and test splits are done based on time order.
- Advantages:
 - Maintains the temporal order of data, which is important for time series problems.

- Disadvantages:
 - More complex to implement and evaluate.

1.0.2 Who Should Use Cross-Validation?

Cross-validation is suitable for:

- **Data scientists and machine learning practitioners** looking for a robust model evaluation technique.
- **Machine learning beginners** who want to prevent overfitting or underfitting and need to compare different models and tune hyperparameters.
- **Researchers and engineers** working with limited data, where splitting the data into a simple train-test set might not be enough.
- **Kaggle competitors** or those involved in machine learning competitions where cross-validation provides a better estimate of performance on unseen data.

Advantages of Cross-Validation:

- **More reliable evaluation:** By using all the data for training and testing in different iterations, cross-validation gives a better estimate of model performance.
- **Reduces overfitting risk:** Prevents the model from overfitting to a particular train-test split.
- **Maximizes data usage:** Every data point is used for both training and validation, ensuring the model is evaluated on the entire dataset.
- **Model comparison:** Helps compare the performance of different models using a fair and unbiased technique.
- **Hyperparameter tuning:** Aids in finding the best hyperparameters for a model by testing performance across multiple folds.

Disadvantages of Cross-Validation:

- **Computationally expensive:** Training and testing the model multiple times (depending on k) can be time-consuming, especially for large datasets or complex models.
- **Bias-variance tradeoff:** While k -fold cross-validation provides a more reliable estimate, it can still have some variance if k is small. LOOCV reduces bias but may increase variance.

Real-World Applications of Cross-Validation:

- **Model Selection:** Comparing several machine learning algorithms (e.g., decision trees, logistic regression, SVM) and selecting the one with the best cross-validated performance.
- **Hyperparameter Tuning:** Tuning parameters like the depth of a decision tree, the learning rate of a neural network, or the C parameter in an SVM using cross-validation as part of a grid search or randomized search.
- **Time Series Prediction:** Using time-based cross-validation (rolling windows) to validate models for stock prices, weather forecasting, and other time-dependent data.

- **Medical Research:** Evaluating predictive models that classify diseases using limited patient data

1.1 Example

```
[26]: from sklearn.datasets import load_iris
from sklearn.model_selection import KFold, cross_val_score, StratifiedKFold
from sklearn.linear_model import LogisticRegression

# Load dataset
iris = load_iris()
X = iris.data
y = iris.target

# Initialize the model
model = LogisticRegression(max_iter=200)

# Define k-fold cross-validation with 5 folds
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Perform cross-validation and get scores
cv_scores = cross_val_score(model, X, y, cv=kf)

# Print cross-validated accuracy scores
print(f"Cross-validated scores: {cv_scores}")
print(f"Mean score: {cv_scores.mean()}")
```

```
Cross-validated scores: [1.          1.          0.93333333 0.96666667 0.96666667]
Mean score: 0.9733333333333334
```

```
[39]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

np.random.seed(42)
```

```
[28]: fastball_speed = np.random.randint(90,106, size=500)
tommy_john = np.where(fastball_speed > 96, np.random.choice([0,1], size=500,
    ↳p=[0.3, 0.7]), 0)

d = {
    'fastball_speed': fastball_speed,
    'tommy_john': tommy_john
}
```

```
df = pd.DataFrame(d)
df
```

```
[28]:      fastball_speed  tommy_john
0              96             0
1              93             0
2             102             1
3             104             1
4             100             0
..           ...           ...
495            104             1
496            92             0
497            101             1
498            90             0
499            93             0
```

[500 rows x 2 columns]

```
[29]: X = df[['fastball_speed']]
      y = df[['tommy_john']]

      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
      ↪random_state=11)
```

```
[30]: lr = LogisticRegression()

      lr.fit(X_train, y_train.values.ravel())
```

```
[30]: LogisticRegression()
```

```
[31]: lr.score(X_test, y_test)
```

```
[31]: 0.71
```

1.2 Test again

```
[32]: X2_train, X2_test, y2_train, y2_test = train_test_split(X,y, test_size=0.2,
      ↪random_state=25)

      lr = LogisticRegression()
      lr.fit(X2_train, y2_train.values.ravel())
```

```
[32]: LogisticRegression()
```

```
[33]: lr.score(X2_test, y2_test)
```

```
[33]: 0.8
```

```
[34]: cvscore = cross_val_score(lr, X, y.values.ravel(), cv=10)
      cvscore
```

```
[34]: array([0.68, 0.78, 0.72, 0.72, 0.78, 0.82, 0.74, 0.74, 0.76, 0.82])
```

```
[35]: print(np.average(cvscore))
      print(np.std(cvscore))
```

```
0.756
```

```
0.04270831300812523
```

```
[36]: kf = KFold(n_splits=15, shuffle=True, random_state=42)

      kfscore = cross_val_score(lr, X, y.values.ravel(), cv=kf, scoring='f1')
      kfscore2 = cross_val_score(lr, X, y.values.ravel(), cv=kf, scoring='accuracy')

      print(kfscore)
      print(np.average(kfscore))
      print(np.std(kfscore))

      print(kfscore2)
      print(np.average(kfscore2))
      print(np.std(kfscore2))
```

```
[0.66666667 0.82758621 0.72727273 0.7027027 0.6875      0.72727273
 0.74074074 0.72      0.66666667 0.60869565 0.66666667 0.63157895
 0.64516129 0.66666667 0.6      ]
0.6856785107611354
0.05620337068328653
[0.76470588 0.85294118 0.82352941 0.67647059 0.70588235 0.81818182
 0.78787879 0.78787879 0.78787879 0.72727273 0.72727273 0.78787879
 0.66666667 0.78787879 0.63636364]
0.7559120617944146
0.060970963505883366
```

```
[37]: kf3 = StratifiedKFold(n_splits=10, shuffle=True, random_state=11)

      kfscore3 = cross_val_score(lr, X, y.values.ravel(), cv=kf3)

      print(kfscore3)
      print(np.average(kfscore3))
      print(np.std(kfscore3))
```

```
[0.66 0.76 0.76 0.8 0.8 0.74 0.76 0.74 0.78 0.76]
0.756
0.03773592452822642
```

1.3 Pipeline Sample

```
[40]: scales = StandardScaler()
      pipe1 = make_pipeline(scales, lr)
      pipe1.fit(X_train, y_train.values.ravel())
```

```
[40]: Pipeline(steps=[('standardscaler', StandardScaler()),
                       ('logisticregression', LogisticRegression())])
```

```
[42]: scorespipe = cross_val_score(pipe1, X, y.values.ravel(), cv=10)

      print(scorespipe)
      print(np.average(scorespipe))
      print(np.std(scorespipe))
```

```
[0.68 0.78 0.72 0.72 0.78 0.82 0.74 0.74 0.76 0.82]
0.756
0.04270831300812523
```


15-Hyperparameter_Tuning

October 20, 2024

1 Hyperparameter Tuning

Hyperparameter tuning is the process of finding the optimal values for the hyperparameters of a machine learning model. Unlike model parameters, which are learned from the training data (e.g., weights in a neural network, coefficients in a regression model), hyperparameters are set before the learning process and control how the model is trained or the structure of the model.

Examples of hyperparameters include:

- Learning rate in gradient-based algorithms.
- Number of trees in a Random Forest.
- Depth of a decision tree.
- C (regularization) and kernel in Support Vector Machines.
- k in k-nearest neighbors (KNN).
- Batch size or number of layers in neural networks.

Tuning these hyperparameters is crucial because they can have a significant impact on the model's performance. The goal is to find the best combination of hyperparameter values that yields the highest-performing model.

When to Use Hyperparameter Tuning? Hyperparameter tuning is used:

- **When building any machine learning model:** Most machine learning algorithms have hyperparameters that need to be fine-tuned.
- **When optimizing model performance:** After selecting a base model, hyperparameter tuning can help improve its accuracy, precision, recall, or other performance metrics.
- **To avoid overfitting/underfitting:** The right hyperparameters can prevent a model from overfitting (too complex) or underfitting (too simple) the data.
- **During model validation:** When performing cross-validation, tuning helps adjust the model for better generalization to unseen data.

Hyperparameter tuning is commonly used when:

- Building models for production systems where performance is critical.
- Participating in machine learning competitions.

- Deploying models with specific performance constraints, such as low latency in real-time systems.

1.0.1 How Does Hyperparameter Tuning Work?

The tuning process involves searching through a defined set of hyperparameter values and evaluating the model's performance for each combination of values. There are several methods to tune hyperparameters:

1. Grid Search:

A brute-force approach that evaluates all possible combinations of hyperparameters from a specified grid. You define a grid (list of values) for each hyperparameter, and the algorithm tries all possible combinations. Example: If you have two hyperparameters with 3 values each, Grid Search will evaluate all $3 \times 3 = 9$ combinations.

Pros:

- Exhaustive: Tests every combination.
- Simple to implement and understand.

Cons:

- Computationally expensive for large datasets or a large number of hyperparameters.
- May test irrelevant combinations of hyperparameters, leading to inefficiency..

2. Random Search:

- Instead of evaluating all combinations, Random Search randomly samples combinations of hyperparameters.
- You define a range of values for each hyperparameter, and the algorithm randomly selects combinations for evaluation.

Pros:

- More efficient than Grid Search as it focuses on randomly chosen combinations.
- Can still find good hyperparameter values without exhaustive search.

Cons:

- May miss the best combination if it isn't sampled.
- Still computationally expensive if many iterations are performed.

3. Bayesian Optimization:

- Bayesian optimization models the hyperparameter search as a probability problem and selects hyperparameter values based on the likelihood of improving the model's performance.
- This method focuses on exploring the hyperparameter space intelligently rather than randomly, making it more efficient than Grid or Random Search.

Pros:

- More efficient, as it focuses on the most promising hyperparameter areas.

- Can achieve high performance with fewer evaluations.

Cons:

- More complex to implement and requires specialized libraries (e.g., Hyperopt, Optuna).

4. Gradient-Based Optimization:

Similar to gradient descent used in model training, gradient-based methods adjust hyperparameters based on how changes impact performance.

Pros:

- Can quickly converge to optimal hyperparameters.

Cons:

- Limited to continuous hyperparameters (can't handle categorical ones like tree depth or number of estimators).

5. Manual Search:

A trial-and-error approach where hyperparameters are manually tuned based on intuition and experience.

Pros:

- Simple and effective when you have experience with the model or the dataset.

Cons:

- May miss optimal settings due to lack of systematic exploration.
- Time-consuming if tried exhaustively.

1.0.2 Who Should Use Hyperparameter Tuning?

Hyperparameter tuning is essential for:

- Machine learning practitioners and data scientists who want to maximize model performance.
- Researchers testing different hypotheses or models and trying to extract the best possible results from their experiments.
- Beginners who are learning machine learning and want to experience the difference in model behavior with tuned hyperparameters.
- Competitors in machine learning contests, such as Kaggle, where the smallest improvement can make a difference in ranking.

Advantages of Hyperparameter Tuning:

- **Improved performance:** Well-tuned hyperparameters can significantly improve model accuracy, precision, recall, and other metrics.
- **Prevention of overfitting or underfitting:** Tuning can help find the sweet spot between a model that's too complex and one that's too simple.

- **Customization:** Allows for deep customization of machine learning models to suit the specific characteristics of your dataset.

Disadvantages of Hyperparameter Tuning:

- **Computational expense:** For large datasets or complex models, tuning (especially Grid Search) can be time-consuming and resource-intensive.
- **Complexity:** Some tuning methods, such as Bayesian optimization, can be hard to implement and require specialized knowledge.
- **Local minima:** Certain tuning strategies, like gradient-based methods, may get stuck in local minima and miss the global optimum.

1.0.3 Real-World Applications of Hyperparameter Tuning:

- **Deep Learning:** Tuning hyperparameters like learning rate, batch size, and the number of layers to improve model performance on image, text, or speech data.
- **Kaggle Competitions:** Competitors frequently use hyperparameter tuning to optimize their models for the best performance on the leaderboard.
- **Predictive Modeling:** In fields like finance and healthcare, tuning hyperparameters to build accurate predictive models for stock prices, patient outcomes, etc.
- **Natural Language Processing (NLP):** Tuning hyperparameters for algorithms like transformers or recurrent neural networks for better text classification, translation, or sentiment analysis.

```
[30]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV,
↳ RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
```

```
[31]: mean1 = 55
std_dev1 = 10
num_samples = 500

column1_numbers = np.random.normal(mean1, std_dev1, num_samples)
column1_numbers = np.clip(column1_numbers, 30, 120)
column1_numbers = np.round(column1_numbers).astype(int)

mean2 = 18
std_dev2 = 3

column2_numbers = np.random.normal(mean2, std_dev2, num_samples)
column2_numbers = np.clip(column2_numbers, 30, 120)
column2_numbers = np.round(column2_numbers).astype(int)
```

```

column3_numbers = np.random.randint(2, size=num_samples)
column3_numbers[column1_numbers > mean1] = 1

data = {'Miles_Per_Week': column1_numbers, 'Farthest_run': column2_numbers,
        'Qualified_Boston_Marathon': column3_numbers}
df = pd.DataFrame(data)
df

```

```

[31]:      Miles_Per_Week  Farthest_run  Qualified_Boston_Marathon
0                55           30                0
1                59           30                1
2                72           30                1
3                72           30                1
4                30           30                1
..              ...           ...                ...
495              51           30                0
496              51           30                1
497              76           30                1
498              47           30                1
499              61           30                1

```

[500 rows x 3 columns]

```

[32]: X = df.iloc[:, 0:2]
      y = df.iloc[:, 2]

      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,
        random_state=26)

      rf = RandomForestClassifier()
      param_grid = [{
          'n_estimators': [500, 1000, 1500],
          'criterion': ['entropy', 'gini'],
          'min_samples_split': [5,10,15],
          'min_samples_leaf': [1,2,4],
          'max_depth': [10,20,30]
      }]

      grid_search = GridSearchCV(rf, param_grid, cv=2, scoring='accuracy', n_jobs=-1)
      grid_search.fit(X_train, y_train)

```

```

[32]: GridSearchCV(cv=2, estimator=RandomForestClassifier(), n_jobs=-1,
      param_grid=[{'criterion': ['entropy', 'gini'],
                    'max_depth': [10, 20, 30],
                    'min_samples_leaf': [1, 2, 4],
                    'min_samples_split': [5, 10, 15],

```

```
        'n_estimators': [500, 1000, 1500]]],  
        scoring='accuracy')
```

```
[33]: grid_search.best_score_
```

```
[33]: 0.7885714285714286
```

```
[34]: grid_search.best_params_
```

```
[34]: {'criterion': 'gini',  
      'max_depth': 20,  
      'min_samples_leaf': 4,  
      'min_samples_split': 15,  
      'n_estimators': 1500}
```

```
[35]: random_param_grid = [{  
      'n_estimators': [500, 1000, 1500],  
      'criterion': ['entropy', 'gini'],  
      'min_samples_split': [5,10,15],  
      'min_samples_leaf': [1,2,4],  
      'max_depth': [10,20,30]  
    }]  
  
    random_grid_search = RandomizedSearchCV(rf, random_param_grid, cv=5,  
      ↪scoring='accuracy', n_jobs=-1, random_state=26)  
    random_grid_search.fit(X_train, y_train)
```

```
[35]: RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,  
      param_distributions=[{'criterion': ['entropy', 'gini'],  
      'max_depth': [10, 20, 30],  
      'min_samples_leaf': [1, 2, 4],  
      'min_samples_split': [5, 10, 15],  
      'n_estimators': [500, 1000, 1500]}],  
      random_state=26, scoring='accuracy')
```

```
[36]: random_grid_search.best_score_
```

```
[36]: 0.7457142857142858
```

```
[37]: random_grid_search.best_params_
```

```
[37]: {'n_estimators': 1000,  
      'min_samples_split': 10,  
      'min_samples_leaf': 1,  
      'max_depth': 20,  
      'criterion': 'entropy'}
```

16-PCA_Analysis

October 20, 2024

1 PCA Analysis

Principal Component Analysis (PCA) is a dimensionality reduction technique that is used to reduce the number of variables (features) in a dataset while retaining as much of the important information (variance) as possible. PCA does this by identifying directions, called principal components, along which the variance in the data is maximized.

PCA can be thought of as a way to transform a large set of possibly correlated variables into a smaller set of uncorrelated variables (the principal components). The first principal component captures the most variance in the data, the second captures the second most, and so on.

In simpler terms:

- **Dimensionality reduction:** PCA reduces the number of features by creating new ones (principal components) that explain most of the variability.
- **Data simplification:** By transforming data into fewer dimensions, PCA simplifies the data while keeping the most critical information.

1.0.1 When to Use PCA?

PCA is typically used in scenarios where:

- **The data has many features:** When dealing with high-dimensional data, PCA helps reduce complexity by creating a smaller set of meaningful features.
- **There is multicollinearity:** PCA is useful when features in the dataset are highly correlated. The principal components are uncorrelated, which helps in creating a better predictive model.
- **You need visualization of high-dimensional data:** PCA is commonly used for visualizing data in 2D or 3D, even if the original data has more dimensions.
- **To avoid overfitting:** By reducing the number of dimensions, PCA helps prevent overfitting in machine learning models.
- **Preprocessing:** Often used as a data preprocessing step before applying machine learning models like clustering (e.g., K-means), classification (e.g., logistic regression), or regression models.

PCA is used in applications like:

- **Image compression:** Reducing the number of pixels (features) while preserving image quality.

- **Genetics:** Analyzing large datasets of gene expression data.
- **Financial markets:** Reducing the number of stock price indicators while retaining information about market movement.
- **Natural language processing:** Reducing dimensionality of word embeddings or text data before applying machine learning.

1.0.2 How Does PCA Work?

PCA works through a series of mathematical steps that transform the data into a new set of coordinates, called principal components. These components are ordered such that the first one accounts for the most variance, the second for the next largest variance, and so on.

Here's how PCA works step by step:

1. Standardize the Data:

- PCA works best when data is standardized. This is because PCA is affected by the scale of the variables. Standardization transforms the features so that they have a mean of 0 and a standard deviation of 1.

2. Compute the Covariance Matrix:

- The covariance matrix shows how much the features vary from the mean with respect to each other. This helps identify patterns of correlation in the data.
- The covariance matrix is an $m \times m$ matrix (where m is the number of features), with entries representing the covariance between each pair of features.

3. Compute the Eigenvalues and Eigenvectors:

- Eigenvalues represent the amount of variance explained by each principal component, while eigenvectors represent the direction of these components.
- Eigenvalues are used to rank the principal components in order of importance (from the largest to the smallest eigenvalue).

4. Form the Principal Components:

- The principal components are linear combinations of the original features. Each component is a vector in the new feature space. The first component explains the largest amount of variance, the second explains the next largest, and so on.

5. Select the Number of Principal Components (k):

- Usually, you select the number of principal components k that explain a certain threshold of variance (e.g., 90%, 95%). This helps to retain the most important information while reducing the number of dimensions.

6. Project the Data:

- The original data is projected onto the new k -dimensional space (spanned by the selected principal components), creating a reduced-dimensional representation.

Advantages of PCA:

- **Dimensionality reduction:** Significantly reduces the number of features while retaining most of the variance, making the dataset simpler and more manageable.
- **Uncorrelated features:** The principal components are linearly uncorrelated, solving multicollinearity issues in the original features.
- **Data visualization:** PCA helps in visualizing high-dimensional data in two or three dimensions.
- **Speeds up model training:** Reducing the number of features can make machine learning models faster and less prone to overfitting.

Disadvantages of PCA:

- **Loss of interpretability:** The principal components are linear combinations of the original features, which makes them hard to interpret.
- **Sensitive to scaling:** PCA requires features to be on the same scale, so data standardization is often necessary.
- **Linear transformations only:** PCA only captures linear relationships, so it may not work well for datasets with complex, nonlinear structures.
- **Variance-focused:** PCA maximizes variance, but it does not directly optimize for predictive power, so it may not always yield the best features for prediction.

1.0.3 Real-World Applications of PCA:

- **Image Compression:** PCA can reduce the dimensionality of image data by retaining the most important pixel values, enabling image compression without significant loss of quality.
- **Genomics:** PCA is widely used in genomics to reduce the number of gene expression measurements while maintaining the most important patterns of variation across samples.
- **Finance:** In stock market analysis, PCA is used to reduce the number of correlated variables (e.g., stock prices) into principal components that explain overall market movements.
- **Natural Language Processing (NLP):** PCA is often used to reduce the dimensionality of word vectors (embeddings) in text analysis.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

df = pd.read_csv('2022mlbteams.csv')
df
```

```
[1]:
```

	Tm	#Bat	BatAge	R/G	G	PA	AB	R	H	\
0	Arizona Diamondbacks	57	26.5	4.33	162	6027	5351	702	1232	

1	Atlanta Braves	53	27.5	4.87	162	6082	5509	789	1394
2	Baltimore Orioles	58	27.0	4.16	162	6049	5429	674	1281
3	Boston Red Sox	54	28.8	4.54	162	6144	5539	735	1427
4	Chicago Cubs	64	27.9	4.06	162	6072	5425	657	1293
5	Chicago White Sox	44	29.3	4.23	162	6123	5611	686	1435
6	Cincinnati Reds	66	29.4	4.00	162	5978	5380	648	1264
7	Cleveland Guardians	50	25.9	4.31	162	6163	5558	698	1410
8	Colorado Rockies	43	29.1	4.31	162	6105	5540	698	1408
9	Detroit Tigers	53	27.9	3.44	162	5870	5378	557	1240
10	Houston Astros	45	29.3	4.55	162	6054	5409	737	1341
11	Kansas City Royals	55	27.1	3.95	162	6010	5437	640	1327
12	Los Angeles Angels	66	27.9	3.85	162	5977	5423	623	1265
13	Los Angeles Dodgers	51	29.6	5.23	162	6247	5526	847	1418
14	Miami Marlins	56	28.9	3.62	162	5949	5395	586	1241
15	Milwaukee Brewers	51	29.1	4.48	162	6122	5417	725	1271
16	Minnesota Twins	61	26.9	4.30	162	6113	5476	696	1356
17	New York Mets	61	29.7	4.77	162	6176	5489	772	1422
18	New York Yankees	54	30.2	4.98	162	6172	5422	807	1308
19	Oakland Athletics	64	28.3	3.51	162	5863	5314	568	1147
20	Philadelphia Phillies	56	28.1	4.61	162	6077	5496	747	1392
21	Pittsburgh Pirates	68	26.3	3.65	162	5912	5331	591	1186
22	San Diego Padres	55	28.2	4.35	162	6175	5468	705	1317
23	Seattle Mariners	59	27.5	4.26	162	6117	5375	690	1236
24	San Francisco Giants	66	30.0	4.42	162	6117	5392	716	1261
25	St. Louis Cardinals	51	28.8	4.77	162	6165	5496	772	1386
26	Tampa Bay Rays	61	27.0	4.11	162	6008	5412	666	1294
27	Texas Rangers	55	28.0	4.36	162	6029	5478	707	1308
28	Toronto Blue Jays	51	27.1	4.78	162	6158	5555	775	1464
29	Washington Nationals	55	28.7	3.72	162	5998	5434	603	1351

	2B	...	OPS	OPS+	TB	GDP	HBP	SH	SF	IBB	LOB	Playoffs
0	262	...	0.689	95	2061	97	60	31	50	14	1039	0
1	298	...	0.761	109	2443	103	66	1	36	13	1030	1
2	275	...	0.695	99	2119	95	83	12	43	10	1095	0
3	352	...	0.731	102	2268	131	63	12	50	23	1133	0
4	265	...	0.698	94	2097	130	84	19	36	16	1100	0
5	272	...	0.698	97	2172	127	73	16	35	9	1117	1
6	235	...	0.676	85	2003	127	92	12	33	6	1020	0
7	273	...	0.699	102	2126	119	81	22	52	36	1156	1
8	280	...	0.713	91	2203	139	61	10	40	10	1113	1
9	235	...	0.632	82	1859	108	58	10	44	8	1015	0
10	284	...	0.743	111	2293	118	60	9	42	18	1068	1
11	247	...	0.686	93	2064	101	48	20	44	7	1091	0
12	219	...	0.687	93	2116	95	54	25	25	28	1050	0
13	325	...	0.775	115	2441	85	56	3	53	22	1159	1
14	248	...	0.658	85	1961	120	70	4	36	6	1045	0
15	251	...	0.724	103	2213	117	80	11	37	25	1102	1

16	269	...	0.718	105	2195	133	62	10	46	11	1126	0
17	272	...	0.744	113	2261	122	112	20	44	25	1158	1
18	225	...	0.751	112	2311	121	70	14	41	36	1093	1
19	249	...	0.626	84	1837	109	59	22	33	7	969	0
20	255	...	0.739	108	2320	116	52	6	44	15	1075	1
21	221	...	0.655	85	1939	96	54	19	32	14	1016	0
22	275	...	0.700	102	2087	95	65	17	46	24	1174	1
23	229	...	0.704	106	2094	120	89	9	45	17	1129	1
24	255	...	0.705	100	2101	109	95	6	53	14	1115	0
25	290	...	0.745	112	2309	112	80	5	45	11	1132	1
26	296	...	0.686	99	2041	93	57	7	31	13	1074	1
27	224	...	0.696	96	2166	82	47	10	38	12	1007	0
28	307	...	0.760	117	2395	136	55	8	33	13	1111	1
29	252	...	0.688	98	2051	141	60	20	37	12	1099	0

[30 rows x 30 columns]

```
[2]: df.drop(columns=['Tm', '#Bat'], axis=1, inplace=True)
```

```
[3]: X = df.iloc[:, 0:27]
y = df.iloc[:, 27]

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
↳random_state=21)
scaleStandard = StandardScaler()
X_train = scaleStandard.fit_transform(X_train)

df.columns
```

```
[3]: Index(['BatAge', 'R/G', 'G', 'PA', 'AB', 'R', 'H', '2B', '3B', 'HR', 'RBI',
'SB', 'CS', 'BB', 'SO', 'BA', 'OBP', 'SLG', 'OPS', 'OPS+', 'TB', 'GDP',
'HBP', 'SH', 'SF', 'IBB', 'LOB', 'Playoffs'],
dtype='object')
```

```
[4]: X_train = pd.DataFrame(X_train, columns=['BatAge', 'R/G', 'G', 'PA', 'AB', 'R',
↳'H', '2B', '3B', 'HR', 'RBI',
'SB', 'CS', 'BB', 'SO', 'BA', 'OBP', 'SLG', 'OPS', 'OPS+', 'TB', 'GDP',
'HBP', 'SH', 'SF', 'IBB', 'LOB'])

X_train
```

```
[4]:
```

	BatAge	R/G	G	PA	AB	R	H	2B	\
0	-0.032987	-1.708056	0.0	-2.105939	-2.081256	-1.713743	-2.374828	-0.496051	
1	-2.144152	0.099809	0.0	1.070839	1.875833	0.099373	1.267922	0.293245	
2	-0.296883	0.212801	0.0	-0.348122	0.578427	0.224896	-0.144856	-1.318234	
3	-1.088570	-0.713730	0.0	-0.549318	-0.086494	-0.709556	0.118309	-0.561825	

4	1.638352	1.613896	0.0	1.166142	-0.329757	1.619600	-0.144856	-1.285346
5	-1.176535	-0.352157	0.0	-0.570496	-0.491933	-0.346933	-0.338766	1.049654
6	-1.616361	0.145006	0.0	-0.369300	-1.481206	0.155161	-1.197514	-0.068515
7	-0.208917	0.777758	0.0	0.160163	0.870343	0.782778	1.018608	-0.298727
8	0.934630	-0.600738	0.0	-0.888174	-1.010896	-0.597979	-0.754289	-0.956473
9	1.110561	2.178854	0.0	1.960337	1.356871	2.177482	1.378728	2.003386
10	0.318874	-1.233491	0.0	-0.676389	-0.135146	-1.225596	0.450727	-0.397389
11	-0.384848	-0.939713	0.0	-0.898763	-0.313540	-0.946656	-0.740438	-1.482670
12	0.494804	-1.459474	0.0	-1.195263	-0.767632	-1.462696	-1.072857	-0.528938
13	0.670735	0.099809	0.0	0.456662	1.583917	0.099373	1.240220	0.523457
14	0.846665	0.642169	0.0	-0.083390	-0.540586	0.643307	0.312219	0.655006
15	1.198526	1.139331	0.0	1.208499	0.756820	1.131454	1.434131	0.260358
16	-1.176535	-0.239165	0.0	-0.136337	-0.216234	-0.235356	-0.518826	0.359020
17	0.406839	0.619570	0.0	0.869643	1.567699	0.615413	1.503384	2.891344
18	0.406839	1.139331	0.0	1.092017	0.870343	1.131454	0.935503	0.852330
19	-1.264500	0.077211	0.0	0.541376	0.545992	0.071479	0.519981	0.161696
20	-0.384848	-0.465149	0.0	0.107216	-0.281105	-0.472456	-0.352617	0.030147
21	1.462422	0.348391	0.0	0.583733	-0.816285	0.350419	-0.795842	-0.298727
22	0.670735	0.483980	0.0	0.636679	-0.410845	0.475943	-0.657334	-0.430276
23	-0.384848	-1.866244	0.0	-2.031814	-1.043331	-1.867160	-1.086707	-0.956473

	3B	HR	...	SLG	OPS	OPS+	TB	GDP \
0	-1.023632	-1.011987	...	-1.991887	-2.198327	-1.578618	-2.106985	-0.269665
1	1.101096	-1.305671	...	-0.440312	-0.149673	0.310476	-0.112710	0.352639
2	-0.359654	0.779487	...	0.062902	-0.233865	-0.319222	0.163314	-1.949889
3	2.030664	-0.982618	...	-0.566115	-0.514502	-0.634071	-0.540547	-0.767509
4	-1.953200	2.424118	...	1.362870	1.309641	1.359973	1.163902	0.477100
5	-0.758041	-0.953250	...	-0.691919	-0.514502	-0.004373	-0.699261	-1.265353
6	0.171527	0.045276	...	-0.356443	-0.430311	-0.424172	-0.561249	-1.016431
7	0.835505	0.985066	...	1.195132	0.972877	0.940174	1.226008	0.165948
8	-0.625245	-0.453987	...	-0.901591	-0.795139	-1.473669	-0.961484	0.850483
9	1.101096	1.190645	...	2.033822	1.983171	1.674822	2.060981	-1.763197
10	-0.359654	-1.041355	...	-0.691919	-0.458375	-0.109323	-0.630255	1.721710
11	1.101096	0.544539	...	-0.146771	-0.486438	-0.634071	-0.181716	-1.140892
12	-0.359654	-0.806408	...	-1.279001	-1.300287	-1.473669	-1.251310	0.414870
13	1.499482	-0.659566	...	0.188705	0.243219	-0.843970	0.418637	1.597249
14	-1.289223	1.249381	...	1.279001	1.085132	1.255023	1.039691	0.290409
15	0.569914	-0.013461	...	0.775788	1.113195	1.464923	0.818872	0.539331
16	0.304323	-0.013461	...	-0.146771	-0.261928	-0.004373	-0.161014	-1.140892
17	-1.422018	-0.483355	...	0.649984	0.748367	0.310476	0.867176	1.099405
18	-0.226859	0.750118	...	1.111263	1.141259	1.359973	1.150101	-0.082974
19	-0.625245	0.192118	...	0.314508	0.383538	0.625325	0.363432	1.223866
20	1.101096	-0.365882	...	-0.272574	-0.177737	-0.529121	-0.312827	1.037175
21	-0.625245	0.338960	...	-0.146771	0.018709	0.100577	-0.285225	-0.269665
22	-0.758041	1.396224	...	0.649984	0.551920	0.415426	0.487643	0.228178
23	0.569914	-1.804934	...	-1.991887	-2.029944	-1.788518	-1.955172	-0.331896

	HBP	SH	SF	IBB	LOB
0	-0.590017	1.201209	-1.197071	-1.046090	-2.339528
1	0.776338	1.201209	1.465682	2.309111	1.440098
2	-1.335301	-0.462003	-0.496347	-0.467607	-1.571476
3	-1.273194	0.924007	0.344523	-1.046090	0.126324
4	0.093161	0.092401	-0.075912	2.309111	0.166748
5	-0.714231	-0.877807	-1.477361	-0.351910	-0.217278
6	-0.527910	2.448618	1.185393	-0.236214	-0.924695
7	-1.024766	-1.016408	0.344523	-0.120517	-0.197066
8	1.459515	-0.184801	-1.197071	-1.161786	-1.308721
9	-0.776338	-1.432211	1.605827	0.689359	1.500734
10	-0.527910	0.924007	-0.636492	-0.467607	0.288020
11	-0.900552	1.617012	-2.318231	1.383538	-0.702364
12	0.093161	-1.293610	-0.776637	-1.161786	-0.803423
13	-0.465803	-0.462003	-0.216057	-0.699000	0.570986
14	-0.527910	-0.600604	0.064233	0.226572	-0.338549
15	2.701655	0.924007	0.344523	1.036449	1.480522
16	0.900552	-0.184801	0.204378	-0.699000	0.207172
17	-0.341589	-0.184801	1.185393	0.805055	0.975224
18	0.714231	-1.155009	0.484668	-0.583304	0.955012
19	-0.403696	-0.462003	0.624813	-0.583304	0.833741
20	0.962659	0.785406	-0.776637	-0.004821	0.308232
21	1.645836	-1.016408	1.605827	-0.236214	0.611410
22	0.714231	-0.323402	-0.636492	1.036449	0.348655
23	-0.652124	-0.462003	0.344523	-0.930393	-1.409780

[24 rows x 27 columns]

```
[5]: X_train.describe().round(3)
```

```
[5]:
```

	BatAge	R/G	G	PA	AB	R	H	2B	3B	\
count	24.000	24.000	24.0	24.000	24.000	24.000	24.000	24.000	24.000	
mean	0.000	0.000	0.0	-0.000	0.000	0.000	0.000	0.000	0.000	
std	1.022	1.022	0.0	1.022	1.022	1.022	1.022	1.022	1.022	
min	-2.144	-1.866	0.0	-2.106	-2.081	-1.867	-2.375	-1.483	-1.953	
25%	-0.561	-0.629	0.0	-0.597	-0.597	-0.626	-0.744	-0.537	-0.658	
50%	0.143	0.100	0.0	0.012	-0.249	0.099	-0.145	-0.184	-0.293	
75%	0.715	0.625	0.0	0.695	0.785	0.622	0.956	0.400	0.902	
max	1.638	2.179	0.0	1.960	1.876	2.177	1.503	2.891	2.031	

	HR	...	SLG	OPS	OPS+	TB	GDP	HBP	SH	\
count	24.000	...	24.000	24.000	24.000	24.000	24.000	24.000	24.000	
mean	-0.000	...	0.000	-0.000	-0.000	-0.000	0.000	0.000	-0.000	
std	1.022	...	1.022	1.022	1.022	1.022	1.022	1.022	1.022	
min	-1.805	...	-1.992	-2.198	-1.789	-2.107	-1.950	-1.335	-1.432	
25%	-0.843	...	-0.598	-0.493	-0.634	-0.579	-0.830	-0.668	-0.670	
50%	-0.013	...	-0.147	-0.164	-0.004	-0.137	0.197	-0.435	-0.254	

75%	0.757	...	0.681	0.804	0.704	0.831	0.617	0.730	0.924
max	2.424	...	2.034	1.983	1.675	2.061	1.722	2.702	2.449


	SF	IBB	LOB
count	24.000	24.000	24.000
mean	0.000	-0.000	0.000
std	1.022	1.022	1.022
min	-2.318	-1.162	-2.340
25%	-0.672	-0.699	-0.728
50%	0.134	-0.294	0.187
75%	0.520	0.718	0.667
max	1.606	2.309	1.501

[8 rows x 27 columns]


```
[6]: pca1 = PCA()
X_pca1 = pca1.fit_transform(X_train)
pca1.explained_variance_ratio_
```

```
[6]: array([5.03143586e-01, 1.29565631e-01, 8.99453783e-02, 5.97627178e-02,
4.75204409e-02, 3.60767505e-02, 3.53674555e-02, 2.44602746e-02,
1.77297976e-02, 1.52080776e-02, 1.26222548e-02, 1.16581886e-02,
7.54086691e-03, 3.93586244e-03, 3.24827865e-03, 1.37155962e-03,
6.27774032e-04, 1.54100200e-04, 4.62620159e-05, 9.33531667e-06,
4.48882814e-06, 6.43748708e-07, 2.74352526e-07, 3.89906609e-34])
```

```
[7]: plt.bar(range(1, len(pca1.explained_variance_) + 1), pca1.explained_variance_)
plt.ylabel("Explained variance")
plt.xlabel('Components')
plt.plot(range(1, len(pca1.explained_variance_) + 1), np.cumsum(pca1.
    explained_variance_), c='red', label="Cumulative Explained Variance")
plt.legend(loc="upper left")
plt.show()
```



```
[8]: plt.plot(pca1.explained_variance_ratio_)
plt.xlabel('Number of Components')
plt.ylabel('Cumulative explained variance')
plt.show()
```



```
[9]: pca2 = PCA(0.95)
X_pca2 = pca2.fit_transform(X_train)
X_pca2.shape
```


```
[9]: (24, 10)
```

```
[10]: pca2.explained_variance_ratio_
```

```
[10]: array([0.50314359, 0.12956563, 0.08994538, 0.05976272, 0.04752044,
          0.03607675, 0.03536746, 0.02446027, 0.0177298 , 0.01520808])
```

```
[11]: pca2c = PCA(n_components=2)
X_pca2c = pca2c.fit_transform(X_train)

colormap = plt.get_cmap('coolwarm')
plt.figure()
scatter = plt.scatter(X_pca2c[:, 0], X_pca2c[:, 1], c=y_train, cmap=colormap)
plt.xlabel('PCA1')
plt.ylabel('PCA2')
plt.colorbar(scatter, label="Playoffs")
plt.show()
```

17-AdaBoost

October 20, 2024

1 Adaboost

AdaBoost (Adaptive Boosting) is an ensemble learning technique that combines multiple weak classifiers to form a strong classifier. It works by training weak classifiers (often decision stumps or shallow trees) sequentially, where each classifier focuses more on the mistakes made by the previous ones.

The key idea behind AdaBoost is to assign more weight to the misclassified instances from the previous classifier, so that the next classifier can focus on these harder-to-classify examples. By combining the predictions of these weak learners, AdaBoost can create a highly accurate model.

AdaBoost = Adaptive + Boosting

- **Boosting:** Combining weak classifiers to create a strong one.
- **Adaptive:** Adjusting the importance (weights) of misclassified examples.

When to Use AdaBoost? AdaBoost is particularly useful when:

- You want to improve the performance of a weak learner (such as a decision stump).
- You are dealing with binary classification problems, although it can be extended to multi-class classification as well.
- You need an efficient model that can provide high accuracy without too much computational cost.
- Your dataset has moderate noise. However, AdaBoost can be sensitive to outliers, so noisy datasets should be preprocessed carefully.

How Does AdaBoost Work? AdaBoost trains classifiers sequentially, each one focusing more on the mistakes of the previous one. Here's the step-by-step breakdown:

1. Initialize Weights:

- All instances in the training set are given equal weights initially. For N data points, the weight for each data point is $1/N$.

2. Train the First Weak Classifier:

A weak classifier (usually a decision stump) is trained on the weighted data. It predicts the class for each example, and the weighted error is calculated.

3. Update Weights:

- Increase the weights of the misclassified examples so that the next classifier pays more attention to these points.
- Decrease the weights of correctly classified examples.

4. Repeat:

- Train a new weak classifier using the updated weights.
- Adjust the weights based on the errors of the new classifier.
- This process continues for a specified number of iterations or until a stopping criterion is met.

5. Combine Weak Learners:

- After multiple iterations, the weak classifiers are combined to make the final prediction.
- Each classifier is given a weight based on its accuracy, with more accurate classifiers contributing more to the final decision.

1.0.1 Who Should Use AdaBoost?

- **Data scientists and machine learning engineers:** Who want to improve the performance of simple models in tasks like binary classification.
- **Beginners and intermediate learners:** AdaBoost is relatively simple to understand, making it suitable for learners who are exploring boosting techniques.
- **Industries that rely on classification tasks:** AdaBoost can be used in fraud detection, text classification, image recognition, and customer churn prediction.

Advantages of AdaBoost:

- **Improves weak learners:** Transforms weak classifiers into a strong one.
- **No parameter tuning:** Unlike many machine learning algorithms, AdaBoost often works well with minimal tuning.
- **Efficient and scalable:** It can handle large datasets efficiently.
- **Feature selection:** AdaBoost implicitly performs feature selection by giving more importance to features that are useful for classification.

Disadvantages of AdaBoost:

- **Sensitive to noisy data:** Outliers can disrupt the performance of AdaBoost because they get higher weights, causing overfitting.
- **Computation cost:** For very large datasets, the sequential training of weak learners can become computationally expensive.
- **Works best with weak learners:** Using a strong learner as the base model can lead to overfitting.

```
[43]: from sklearn.model_selection import train_test_split, GridSearchCV
      from sklearn.datasets import make_classification
      from sklearn.ensemble import AdaBoostClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix,
    ↪classification_report
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
```

```
[44]: X, y = make_classification(n_samples=2000, n_features=10, n_informative=8,
    ↪n_redundant=2, random_state=11)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
    ↪2, random_state=11)
abc = AdaBoostClassifier()
abc.fit(X_train, y_train)
```

c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```
[44]: AdaBoostClassifier()
```

```
[45]: y_pred = abc.predict(X_test)
y_pred
```

```
[45]: array([0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1,
    1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
    0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
    1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0,
    1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1,
    0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0,
    0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1,
    0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
    0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1,
    0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1,
    0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0,
    1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1,
    1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0,
    1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
    1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1,
    1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0,
    1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0,
    1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1,
    0, 1, 0, 0])
```

```
[46]: print(accuracy_score(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```

0.815
[[171  34]
 [ 40 155]]

```

	precision	recall	f1-score	support
0	0.81	0.83	0.82	205
1	0.82	0.79	0.81	195
accuracy			0.81	400
macro avg	0.82	0.81	0.81	400
weighted avg	0.82	0.81	0.81	400

1.1 Logistic Regression

```
[47]: abclog = AdaBoostClassifier(estimator=LogisticRegression())
      abclog.fit(X_train, y_train)
```

c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```
[47]: AdaBoostClassifier(estimator=LogisticRegression())
```

```
[48]: y_pred2 = abclog.predict(X_test)
      y_pred2
```

```
[48]: array([0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1,
          0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
          0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1,
          1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
          1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1,
          0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
          0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1,
          1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
          1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1,
          0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1,
          0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1,
          1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0,
          0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1,
          1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0,
          1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
          1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
          1, 0, 0, 0, 0])
```

```
[49]: print(accuracy_score(y_test, y_pred2))
      print(confusion_matrix(y_test, y_pred2))
      print(classification_report(y_test, y_pred2))
```

0.785

```
[[167  38]
```

```
 [ 48 147]]
```

	precision	recall	f1-score	support
0	0.78	0.81	0.80	205
1	0.79	0.75	0.77	195
accuracy			0.79	400
macro avg	0.79	0.78	0.78	400
weighted avg	0.79	0.79	0.78	400

1.2 Support Vector Machine

```
[50]: svc = SVC(kernel='linear', probability=True)

abcsvm = AdaBoostClassifier(estimator=svc, n_estimators=25, learning_rate=0.1)
abcsvm.fit(X_train, y_train)
```

c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```
[50]: AdaBoostClassifier(estimator=SVC(kernel='linear', probability=True),
                        learning_rate=0.1, n_estimators=25)
```

```
[51]: y_pred3 = abcsvm.predict(X_test)
      y_pred3
```

```
[51]: array([0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1,
            0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
            0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1,
            1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
            1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1,
            0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
            0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1,
            1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
            1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1,
            0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1,
            0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
            0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
```

```

1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1,
1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1,
0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1,
1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1,
0, 0, 0, 0])

```

```

[52]: print(accuracy_score(y_test, y_pred3))
      print(confusion_matrix(y_test, y_pred3))
      print(classification_report(y_test, y_pred3))

```

0.7925

[[168 37]

[46 149]]

	precision	recall	f1-score	support
0	0.79	0.82	0.80	205
1	0.80	0.76	0.78	195
accuracy			0.79	400
macro avg	0.79	0.79	0.79	400
weighted avg	0.79	0.79	0.79	400

1.3 Hyperparameter Tuning

```

[54]: param_grid = {
      'n_estimators': [1,5,10,25,50,100,500],
      'learning_rate': [0.00001, 0.0001, 0.001, 0.1, 0.5, 1.0]
      }

      abc_grid = GridSearchCV(abc, param_grid, cv=3, n_jobs=-1)
      abc_grid.fit(X_train, y_train)

```

c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```

[54]: GridSearchCV(cv=3, estimator=AdaBoostClassifier(), n_jobs=-1,
                  param_grid={'learning_rate': [1e-05, 0.0001, 0.001, 0.1, 0.5, 1.0],
                              'n_estimators': [1, 5, 10, 25, 50, 100, 500]})

```

```

[55]: abc_grid.best_params_

```

```

[55]: {'learning_rate': 0.1, 'n_estimators': 500}

```

```
[56]: abc2 = AdaBoostClassifier(learning_rate=0.2, n_estimators=500)
      abc2.fit(X_train, y_train)
```

```
c:\Users\ikiga\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\ensemble\_weight_boosting.py:519: FutureWarning: The SAMME.R
algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME
algorithm to circumvent this warning.
      warnings.warn(
```

```
[56]: AdaBoostClassifier(learning_rate=0.2, n_estimators=500)
```

```
[58]: y_pred4 = abc2.predict(X_test)
      y_pred4
```

```
[58]: array([0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1,
          1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
          0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1,
          1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
          1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1,
          0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
          0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1,
          0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0,
          0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1,
          0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1,
          0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
          1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
          1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1,
          1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1,
          0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1,
          1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0,
          1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
          1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0,
          0, 1, 0, 0])
```

```
[59]: print(accuracy_score(y_test, y_pred4))
      print(confusion_matrix(y_test, y_pred4))
      print(classification_report(y_test, y_pred4))
```

```
0.8275
```

```
[[173  32]
```

```
 [ 37 158]]
```

	precision	recall	f1-score	support
0	0.82	0.84	0.83	205
1	0.83	0.81	0.82	195
accuracy			0.83	400
macro avg	0.83	0.83	0.83	400
weighted avg	0.83	0.83	0.83	400

18-Gradient_Boosring

October 20, 2024

1 Gradient Boosting

Gradient Boosting is an ensemble learning technique that builds models in a stage-wise fashion. It creates a strong predictive model by combining the predictions of several weaker models (often decision trees). Unlike AdaBoost, which focuses on misclassified instances, Gradient Boosting minimizes a loss function using the gradient descent optimization algorithm, hence the name “Gradient Boosting.”

1.0.1 When to Use Gradient Boosting?

Gradient Boosting is particularly useful when:

- You have structured or tabular data.
- You want to achieve high accuracy and predictive performance.
- Your data contains complex patterns that simpler models struggle to capture.
- You are interested in both regression and classification tasks.

Gradient Boosting can handle a variety of loss functions and is often employed in Kaggle competitions due to its robustness.

1.0.2 How Does Gradient Boosting Work?

Gradient Boosting works by sequentially adding weak learners (typically shallow decision trees) to the ensemble, each one correcting the errors made by the previous learners. Here’s a step-by-step breakdown:

1. Initialize the Model:

- Start with a base prediction (often the mean of the target values for regression tasks).

2. Calculate Residuals:

- Calculate the residuals (the difference between actual values and predicted values) of the current model.

3. Fit a Weak Learner:

- Train a new weak learner (e.g., decision tree) on the residuals to predict these errors. The idea is to learn the error of the previous model.

4. Update the Model:

- Update the model by adding the predictions of the new weak learner to the current model predictions, scaled by a learning rate (denoted as η).

5. Repeat:

- Repeat steps 2 to 4 for a specified number of iterations or until a stopping criterion is met.

6. Final Prediction:

- The final prediction is a weighted sum of the weak learners' predictions.

1.0.3 Who Should Use Gradient Boosting?

- **Data scientists and machine learning practitioners:** Who want to build predictive models that leverage the power of ensemble learning.
- **Kaggle competitors:** Gradient Boosting is popular in data science competitions for its high predictive accuracy and flexibility.
- **Industries that require predictive analytics:** Applicable in finance, marketing, health-care, and other fields that deal with regression or classification problems.

Advantages of Gradient Boosting:

- **High predictive performance:** Often achieves state-of-the-art results on various tasks.
- **Flexibility:** Can optimize various loss functions and handle different types of data (e.g., continuous, categorical). **Robustness to overfitting:** When used with techniques like early stopping or regularization (e.g., limiting tree depth).
- **Feature importance:** Provides insight into the importance of features in making predictions.

Disadvantages of Gradient Boosting:

- **Sensitive to noise and outliers:** Can overfit if not properly regularized.
- **Longer training times:** Sequential training can be slower compared to parallelized models like Random Forest.
- **Complexity:** Requires careful tuning of hyperparameters (e.g., number of trees, learning rate, tree depth) for optimal performance.

1.0.4 Real-World Applications of Gradient Boosting:

- **Finance:** Credit scoring and risk assessment, where high accuracy is crucial.
- **Customer segmentation:** In marketing analytics, to understand customer behavior.
- **Healthcare:** Predictive models for patient outcomes and disease risk.
- **Recommendation systems:** To personalize user experiences based on historical data.

```
[19]: import pandas as pd

from sklearn import datasets
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
wine = datasets.load_wine(as_frame=True)
wine
```

```
[19]: {'data':      alcohol  malic_acid  ash  alcalinity_of_ash  magnesium
total_phenols  \
0      14.23      1.71  2.43      15.6      127.0      2.80
1      13.20      1.78  2.14      11.2      100.0      2.65
2      13.16      2.36  2.67      18.6      101.0      2.80
3      14.37      1.95  2.50      16.8      113.0      3.85
4      13.24      2.59  2.87      21.0      118.0      2.80
..      ...      ...      ...      ...      ...      ...
173    13.71      5.65  2.45      20.5      95.0      1.68
174    13.40      3.91  2.48      23.0      102.0      1.80
175    13.27      4.28  2.26      20.0      120.0      1.59
176    13.17      2.59  2.37      20.0      120.0      1.65
177    14.13      4.10  2.74      24.5      96.0      2.05

      flavanoids  nonflavanoid_phenols  proanthocyanins  color_intensity  hue
\
0      3.06      0.28      2.29      5.64  1.04
1      2.76      0.26      1.28      4.38  1.05
2      3.24      0.30      2.81      5.68  1.03
3      3.49      0.24      2.18      7.80  0.86
4      2.69      0.39      1.82      4.32  1.04
..      ...      ...      ...      ...      ...
173    0.61      0.52      1.06      7.70  0.64
174    0.75      0.43      1.41      7.30  0.70
175    0.69      0.43      1.35     10.20  0.59
176    0.68      0.53      1.46      9.30  0.60
177    0.76      0.56      1.35      9.20  0.61

      od280/od315_of_diluted_wines  proline
0      3.92     1065.0
1      3.40     1050.0
2      3.17     1185.0
3      3.45     1480.0
4      2.93      735.0
..      ...      ...
173    1.74      740.0
174    1.56      750.0
175    1.56      835.0
176    1.62      840.0
177    1.60      560.0
```

```

[178 rows x 13 columns],
'target': 0      0
1      0
2      0
3      0
4      0
..
173    2
174    2
175    2
176    2
177    2
Name: target, Length: 178, dtype: int32,
'frame':      alcohol  malic_acid  ash  alkalinity_of_ash  magnesium
total_phenols \
0      14.23      1.71  2.43      15.6      127.0      2.80
1      13.20      1.78  2.14      11.2      100.0      2.65
2      13.16      2.36  2.67      18.6      101.0      2.80
3      14.37      1.95  2.50      16.8      113.0      3.85
4      13.24      2.59  2.87      21.0      118.0      2.80
..      ...      ...  ...      ...      ...      ...
173    13.71      5.65  2.45      20.5      95.0      1.68
174    13.40      3.91  2.48      23.0      102.0      1.80
175    13.27      4.28  2.26      20.0      120.0      1.59
176    13.17      2.59  2.37      20.0      120.0      1.65
177    14.13      4.10  2.74      24.5      96.0      2.05

      flavanoids  nonflavanoid_phenols  proanthocyanins  color_intensity  hue
\
0      3.06      0.28      2.29      5.64  1.04
1      2.76      0.26      1.28      4.38  1.05
2      3.24      0.30      2.81      5.68  1.03
3      3.49      0.24      2.18      7.80  0.86
4      2.69      0.39      1.82      4.32  1.04
..      ...      ...      ...      ...  ...
173    0.61      0.52      1.06      7.70  0.64
174    0.75      0.43      1.41      7.30  0.70
175    0.69      0.43      1.35     10.20  0.59
176    0.68      0.53      1.46      9.30  0.60
177    0.76      0.56      1.35      9.20  0.61

      od280/od315_of_diluted_wines  proline  target
0      3.92    1065.0      0
1      3.40    1050.0      0
2      3.17    1185.0      0
3      3.45    1480.0      0
4      2.93     735.0      0

```

..
173	1.74	740.0	2
174	1.56	750.0	2
175	1.56	835.0	2
176	1.62	840.0	2
177	1.60	560.0	2

```
[178 rows x 14 columns],
'target_names': array(['class_0', 'class_1', 'class_2'], dtype='<U7'),
'DESCR': '.. _wine_dataset:\n\nWine recognition
dataset\n-----\n\n**Data Set Characteristics:**\n\nNumber of
Instances: 178\nNumber of Attributes: 13 numeric, predictive attributes and the
class\nAttribute Information:\n    - Alcohol\n    - Malic acid\n    - Ash\n
- Alcalinity of ash\n    - Magnesium\n    - Total phenols\n    - Flavanoids\n
- Nonflavanoid phenols\n    - Proanthocyanins\n    - Color intensity\n    -
Hue\n    - OD280/OD315 of diluted wines\n    - Proline\n    - class:\n    -
class_0\n    - class_1\n    - class_2\n\nSummary
Statistics:\n\n===== \n
Min  Max  Mean  SD\n===== \n
===== \nAlcohol:          11.0  14.8   13.0   0.8\nMalic Acid:
0.74  5.80   2.34  1.12\nAsh:          1.36  3.23   2.36
0.27\nAlcalinity of Ash:          10.6  30.0   19.5   3.3\nMagnesium:
70.0 162.0   99.7  14.3\nTotal Phenols:          0.98  3.88   2.29
0.63\nFlavanoids:          0.34  5.08   2.03  1.00\nNonflavanoid
Phenols:          0.13  0.66   0.36  0.12\nProanthocyanins:          0.41
3.58   1.59  0.57\nColour Intensity:          1.3  13.0   5.1   2.3\nHue:
0.48  1.71   0.96  0.23\nOD280/OD315 of diluted wines: 1.27  4.00   2.61
0.71\nProline:          278  1680   746
315\n===== \n\nMissing
Attribute Values: None\nClass Distribution: class_0 (59), class_1 (71), class_2
(48)\nCreator: R.A. Fisher\nDonor: Michael Marshall
(MARSHALL%PLU@io.arc.nasa.gov)\nDate: July, 1988\n\nThis is a copy of UCI ML
Wine recognition datasets.\nhttps://archive.ics.uci.edu/ml/machine-learning-
databases/wine/wine.data\n\nThe data is the results of a chemical analysis of
wines grown in the same\nregion in Italy by three different cultivators. There
are thirteen different\nmeasurements taken for different constituents found in
the three types of\nwine.\n\nOriginal Owners:\nForina, M. et al, PARVUS -\nAn
Extendible Package for Data Exploration, Classification and
Correlation.\nInstitute of Pharmaceutical and Food Analysis and
Technologies,\nVia Brigata Salerno, 16147 Genoa, Italy.\n\nCitation:\n\nLichman,
M. (2013). UCI Machine Learning Repository\n[https://archive.ics.uci.edu/ml].
Irvine, CA: University of California,\nSchool of Information and Computer
Science.\n\n|details-start|\n**References**\n|details-split|\n\n(1) S.
Aeberhard, D. Coomans and O. de Vel,\nComparison of Classifiers in High
Dimensional Settings,\nTech. Rep. no. 92-02, (1992), Dept. of Computer Science
and Dept. of\nMathematics and Statistics, James Cook University of North
Queensland.\n(Also submitted to Technometrics).\n\nThe data was used with many
```

others for comparing various classifiers. The classes are separable, though only RDA has achieved 100% correct classification. (RDA : 100%, QDA 99.4%, LDA 98.9%, 1NN 96.1% (z-transformed data)) (All results using the leave-one-out technique)

(2) S. Aeberhard, D. Coomans and O. de Vel, "THE CLASSIFICATION PERFORMANCE OF RDA" Tech. Rep. no. 92-01, (1992), Dept. of Computer Science and Dept. of Mathematics and Statistics, James Cook University of North Queensland. (Also submitted to Journal of Chemometrics).

```
{
  'feature_names': ['alcohol',
    'malic_acid',
    'ash',
    'alcalinity_of_ash',
    'magnesium',
    'total_phenols',
    'flavanoids',
    'nonflavanoid_phenols',
    'proanthocyanins',
    'color_intensity',
    'hue',
    'od280/od315_of_diluted_wines',
    'proline']}]}
```

```
[20]: X = wine['data']
      y = wine['target']
```

```
[21]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=17)

gbr = GradientBoostingClassifier()
gbr.fit(X_train, y_train)
```

```
[21]: GradientBoostingClassifier()
```

```
[22]: cross_val_score(gbr, X_train, y_train, cv=3, n_jobs=-1).mean()
```

```
[22]: 0.9221335697399526
```

```
[23]: param_grid = {
      'n_estimators': [10,50,100, 500],
      'learning_rate': [0.0001, 0.001, 0.01, 0.1, 1],
      'max_depth': [3,5,7,9]
    }

gbr2 = GridSearchCV(gbr, param_grid, cv=3, n_jobs=-1)
gbr2.fit(X_train, y_train)
```

```
[23]: GridSearchCV(cv=3, estimator=GradientBoostingClassifier(), n_jobs=-1,
      param_grid={'learning_rate': [0.0001, 0.001, 0.01, 0.1, 1],
```

```
'max_depth': [3, 5, 7, 9],  
'n_estimators': [10, 50, 100, 500]})
```

```
[24]: gbr2.best_params_
```

```
[24]: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 500}
```

```
[25]: gbr2.best_score_
```

```
[25]: 0.9434101654846335
```


19-Extra_Trees_Classifier

October 20, 2024

1 Extra Tress Classifier

Extra Trees, short for Extremely Randomized Trees, is an ensemble learning technique that builds a collection of decision trees in a random manner. Like Random Forest, it constructs multiple decision trees, but it introduces additional randomness into the tree-building process. This results in an ensemble method that typically provides better performance and is less prone to overfitting than traditional decision trees.

Extra Trees can be used for both classification and regression tasks. The key idea is to leverage the randomness in both the selection of features and the thresholds for splitting nodes to create a diverse set of trees.

Extra Trees Classifier is a powerful and efficient ensemble method that excels in classification tasks. Its ability to introduce randomness while leveraging the strengths of decision trees makes it a valuable tool in the machine learning toolkit.

1.0.1 When to Use Extra Trees Classifier?

Extra Trees Classifier is particularly useful when:

- You want to improve the performance of a single decision tree or even Random Forest by introducing more randomness.
- Your dataset is large and complex, and you are looking for a robust model that can capture intricate patterns.
- You are interested in reducing overfitting compared to other tree-based models.
- You need an efficient model for classification or regression tasks that can handle high-dimensional data.

1.0.2 How Does Extra Trees Classifier Work?

The Extra Trees Classifier builds decision trees using the following steps:

1. Bootstrapping:

- Unlike Random Forest, which samples the training data with replacement, Extra Trees uses the entire dataset for each tree, ensuring all data points are considered.

2. Random Feature Selection:

- For each node in a tree, a random subset of features is selected. Instead of considering all features, the model randomly selects a subset of features to determine the best split.

3. Threshold Selection:

- Instead of using the best split point based on the optimization criterion (like Gini impurity or information gain), Extra Trees selects the split threshold randomly from the possible values for the selected features. This means that the trees are built using more randomness, hence the name “Extremely Randomized.”

4. Tree Construction:

- The process of splitting nodes continues recursively until a stopping criterion is reached (e.g., maximum depth, minimum samples per leaf).

5. Prediction:

- For classification tasks, predictions are made based on the majority vote of all the trees in the ensemble. For regression tasks, the average of the predictions from all trees is used.

1.0.3 Who Should Use Extra Trees Classifier?

- **Data scientists and machine learning engineers:** Who want to build robust classification models that leverage tree ensembles.
- **Kaggle competitors:** Extra Trees is often used in competitions for its high accuracy and efficiency.
- **Industries requiring fast and accurate models:** Suitable for applications in finance, healthcare, and customer analytics, where quick and reliable predictions are needed.

Advantages of Extra Trees Classifier:

- **High accuracy:** Typically provides better accuracy than single decision trees and sometimes outperforms Random Forest due to increased randomness.
- **Less prone to overfitting:** The additional randomness in feature selection and threshold determination helps reduce overfitting.
- **Fast training time:** Since all data points are used without bootstrapping and splits are chosen randomly, Extra Trees can train faster than other tree-based methods.
- **Feature importance:** Like other tree-based models, Extra Trees can provide insights into feature importance.

Disadvantages of Extra Trees Classifier:

- **Interpretability:** While feature importance is available, the ensemble nature makes it harder to interpret compared to a single decision tree.
- **Sensitivity to noise:** Although less than traditional decision trees, Extra Trees can still be sensitive to noise and outliers in the data.
- **Limited extrapolation:** The predictions can be less reliable for unseen data outside the range of the training data.

1.0.4 Real-World Applications of Extra Trees Classifier:

- **Credit scoring:** To assess risk based on various financial indicators.

- **Image classification:** Used in computer vision tasks to classify images based on features extracted from pixel data.
- **Customer segmentation:** For marketing strategies and customer behavior analysis.
- **Medical diagnosis:** In predicting patient outcomes based on clinical data.

```
[7]: from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.ensemble import ExtraTreesClassifier

X, y = make_classification(n_features=11, random_state=21)

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,
                                                    random_state=16)

etc = ExtraTreesClassifier(random_state=21)
etc.fit(X_train, y_train)
```

```
[7]: ExtraTreesClassifier(random_state=21)
```

```
[5]: cross_val_score(etc, X_train, y_train, scoring='accuracy', cv=5, n_jobs=-1).
      mean()
```

```
[5]: 0.9375
```

```
[8]: param_grid = {
      'criterion': ['gini', 'entropy'],
      'n_estimators': [100, 250, 500],
      'min_samples_leaf': [5,15,25],
      'max_features': [3,5,7,9,11]
    }

etc2 = GridSearchCV(etc,param_grid, cv=3, n_jobs=-1)

etc2.fit(X_train, y_train)
```

```
[8]: GridSearchCV(cv=3, estimator=ExtraTreesClassifier(random_state=21), n_jobs=-1,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_features': [3, 5, 7, 9, 11],
                              'min_samples_leaf': [5, 15, 25],
                              'n_estimators': [100, 250, 500]})
```

```
[10]: etc2.best_params_
```

```
[10]: {'criterion': 'gini',
      'max_features': 5,
      'min_samples_leaf': 5,
```

```
'n_estimators': 500}
```

```
[12]: etc2.best_score_
```

```
[12]: 0.9377967711301044
```

20-Voting_Classifier

October 20, 2024

1 Voting Classifier

A **Voting Classifier** is an ensemble learning technique that combines multiple machine learning models (often referred to as “base learners” or “weak learners”) to make a single, more robust prediction. The idea is to leverage the strengths of different algorithms and reduce the risk of overfitting while improving the overall accuracy of predictions.

The Voting Classifier is a powerful ensemble method that improves prediction accuracy and robustness by aggregating multiple machine learning models. Its flexibility and effectiveness make it a valuable tool in many predictive modeling tasks.

There are two main types of voting:

- **Hard Voting:** The predicted class is the one that receives the majority of votes from the base models.
- **Soft Voting:** The predicted class is based on the average predicted probabilities of each class, and the class with the highest average probability is chosen.

1.0.1 When to Use Voting Classifier?

Voting Classifier is particularly useful when:

- You want to improve prediction accuracy by combining the strengths of different models.
- You have multiple algorithms that perform well individually, and you want to harness their collective performance.
- You want to reduce variance and increase robustness, especially in scenarios where different models may have different error patterns.

1.0.2 How Does Voting Classifier Work?

The Voting Classifier operates as follows:

1. Model Selection:

- Choose a diverse set of models (e.g., Logistic Regression, Decision Trees, Support Vector Machines, etc.) that will serve as base learners.

2. Training:

- Each model is trained independently on the same training dataset.

3. Prediction:

- For a new data point:
 - **Hard Voting:** Each model predicts a class label, and the class with the most votes is selected as the final prediction.
 - **Soft Voting:** Each model outputs the predicted probabilities for each class. The probabilities are averaged across all models, and the class with the highest average probability is selected.

4. Final Output:

- The Voting Classifier provides the final class label for the data point based on the voting mechanism.

1.0.3 Who Should Use Voting Classifier?

- **Data scientists and machine learning practitioners:** Who aim to build robust models by combining the strengths of multiple algorithms.
- **Kaggle competitors:** The Voting Classifier is a common strategy in competitions to boost accuracy by leveraging multiple models.
- **Industries requiring high accuracy:** Useful in fields like finance, healthcare, and marketing, where precise predictions are critical.

Advantages of Voting Classifier:

- **Improved accuracy:** Combining models often leads to better performance than any individual model.
- **Robustness:** Reduces the risk of overfitting and is less sensitive to noise in the data.
- **Flexibility:** Can incorporate different types of models, allowing the use of various algorithms tailored to specific aspects of the data.
- **Easy implementation:** Simple to implement using libraries like scikit-learn.

Disadvantages of Voting Classifier:

- **Complexity:** More complex than using a single model, as it requires managing multiple algorithms.
- **Computationally expensive:** Training multiple models can be time-consuming and resource-intensive, especially with large datasets.
- **Interpretability:** The ensemble nature makes it harder to interpret the model compared to a single algorithm.

1.0.4 Real-World Applications of Voting Classifier:

1. **Sentiment analysis:** Combining different models to classify sentiments from text data.
2. **Fraud detection:** In finance, where multiple algorithms can identify fraudulent patterns.
3. **Medical diagnosis:** Using various models to improve the accuracy of predicting patient outcomes.

4. **Customer churn prediction:** In marketing, to predict which customers are likely to stop using a service.

```
[34]: import pandas as pd

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, VotingClassifier

X, y = make_classification(n_samples=2500, n_features=15, n_informative=8, \
    n_redundant=2, random_state=11)
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.2, \
    random_state=11)
```

```
[35]: gnb = GaussianNB()
gnb.fit(X_train, Y_train)
```

```
[35]: GaussianNB()
```

```
[36]: cross_val_score(gnb, X_train, Y_train, cv=3).mean()
```

```
[36]: 0.7884931408169789
```

```
[37]: lr = LogisticRegression()
lr.fit(X_train, Y_train)
```

```
[37]: LogisticRegression()
```

```
[38]: cross_val_score(lr, X_train, Y_train, cv=3).mean()
```

```
[38]: 0.7455041248144697
```

```
[39]: rfc = RandomForestClassifier()
rfc.fit(X_train, Y_train)
```

```
[39]: RandomForestClassifier()
```

```
[40]: cross_val_score(rfc, X_train, Y_train, cv=3).mean()
```

```
[40]: 0.9044929487208347
```

```
[41]: vc = VotingClassifier([
    ('NaiveBayes', gnb),
    ('LogisticRegression', lr),
    ('RandomForestClassifier', rfc)
```

```
)  
cross_val_score(vc,X_train, Y_train, cv=3).mean()
```

[41]: 0.8315016665841254

```
[42]: param_grid = {  
        'voting': ['hard', 'soft'],  
        'weights': [(1,1,1), (2,1,1), (1,2,1), (1,1,2)]  
    }  
  
    vc2 = GridSearchCV(vc, param_grid, cv=5, n_jobs=-1)  
    vc2.fit(X_train, Y_train)
```

```
[42]: GridSearchCV(cv=5,  
                  estimator=VotingClassifier(estimators=[('NaiveBayes',  
                                                         GaussianNB()),  
                                                         ('LogisticRegression',  
                                                         LogisticRegression()),  
                                                         ('RandomForestClassifier',  
                                                         RandomForestClassifier())]),  
                  n_jobs=-1,  
                  param_grid={'voting': ['hard', 'soft'],  
                              'weights': [(1, 1, 1), (2, 1, 1), (1, 2, 1),  
                                           (1, 1, 2)]})
```

```
[43]: vc2.best_params_
```

[43]: {'voting': 'hard', 'weights': (1, 1, 2)}

```
[44]: vc2.best_score_
```

[44]: 0.875

21-Linear_Regression

October 20, 2024

1 Linear Regression

Linear Regression is a statistical method used to model the relationship between a dependent variable (target) and one or more independent variables (features). It assumes that this relationship can be described by a straight line (in the case of one independent variable) or a hyperplane (in the case of multiple independent variables). The goal of linear regression is to predict the dependent variable by fitting the best line or plane to the given data points.

Linear Regression is one of the simplest and most widely used algorithms in machine learning. Its strengths lie in its interpretability and efficiency, making it a go-to choice for many regression tasks. However, it's important to ensure that the relationship between variables is linear, as the model can fail in more complex scenarios.

1.0.1 When to Use Linear Regression?

Linear Regression is commonly used when:

- You want to model and understand the linear relationship between one or more independent variables and a dependent variable.
- You are looking for a simple and interpretable model for regression tasks.
- The relationship between the variables is approximately linear (i.e., the change in the target is proportional to the change in the feature).
- You need a baseline model for regression tasks to compare with more complex algorithms.

1.0.2 How Does Linear Regression Work?

Linear Regression works by fitting a line (or hyperplane in the case of multiple variables) to minimize the difference between the predicted values and the actual target values. This difference is captured by the residual sum of squares (RSS):

1.0.3 Who Should Use Linear Regression?

- **Data scientists, statisticians, and analysts:** Who want to model and quantify relationships between variables.
- **Business analysts:** To predict trends like sales forecasting, market demand, etc.
- **Researchers:** In social sciences or economics, where interpretability is important and linear models are commonly used.

1.0.4 Advantages of Linear Regression:

- **Simplicity:** Easy to understand, implement, and interpret.
- **Efficiency:** Computationally fast, making it suitable for large datasets.
- **Interpretability:** The coefficients provide insights into how much each independent variable affects the dependent variable.
- **Baselines for comparison:** Often used as a baseline to compare with more complex models.

1.0.5 Disadvantages of Linear Regression:

- **Linearity assumption:** It assumes a linear relationship between the independent and dependent variables, which may not always hold true.
- **Sensitivity to outliers:** Linear regression can be easily influenced by outliers in the data.
- **Multicollinearity:** When independent variables are highly correlated, it can cause issues with coefficient estimation and model interpretation.
- **Underfitting:** Linear regression might not capture complex patterns in the data, leading to underfitting.

1.0.6 Real-World Applications of Linear Regression:

- **Predicting house prices:** Based on factors like square footage, number of rooms, and location.
- **Sales forecasting:** Using historical sales data and other relevant features to predict future sales.
- **Stock market prediction:** To estimate stock prices based on various financial indicators.
- **Healthcare:** Estimating patient outcomes like the length of hospital stay based on medical data.

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

np.random.seed(42)
```

```
[2]: num_samples = 500
years_of_experience = np.random.randint(2,21, size=num_samples)
slope = (200_000 - 60_000) / 18
```

```

intercept = 60_000

salaries = slope * years_of_experience + intercept + np.random.normal(0,
↪10_000, size=num_samples)

data = {'Years_of_Experience': years_of_experience, 'Salary': salaries}
df = pd.DataFrame(data)

df

```

```

[2]:      Years_of_Experience      Salary
0              8  115037.780010
1             16  182309.972927
2             12  156442.408989
3              9  144753.562169
4              8  130798.818454
..          ...          ...
495             18  203824.097462
496              8  123886.744304
497             14  173813.401529
498              5  101780.575328
499              5  123441.890288

```

[500 rows x 2 columns]

```

[3]: df.describe()

```

```

[3]:      Years_of_Experience      Salary
count      500.000000      500.000000
mean         10.616000  142570.011096
std           5.662922   44935.263058
min           2.000000   54881.134555
25%           5.750000  104426.300731
50%          10.000000  139865.032545
75%          16.000000  182341.125962
max          20.000000  241879.376204

```

```

[4]: plt.figure(figsize=(10,6))
sns.scatterplot(x="Years_of_Experience", y='Salary', data=df, color="blue")
sns.regplot(x="Years_of_Experience", y='Salary', data=df, scatter=False,
↪color='red')

plt.xlabel("Years_of_Experience")
plt.ylabel('Salary')
plt.title('Linear Regression Salary')
plt.show()

```



```
[5]: X = df[['Years_of_Experience']]
     y = df[['Salary']]

     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪random_state=42)

     lr = LinearRegression()
     lr.fit(X_train, y_train)
```

```
[5]: LinearRegression()
```

```
[6]: lr.score(X_train, y_train)
```

```
[6]: 0.9488372177019488
```

```
[7]: lr.score(X_test, y_test)
```

```
[7]: 0.9573704756930143
```

```
[8]: y_pred = lr.predict(X_test)
     y_pred
```

```
[8]: array([[107141.01841859],
           [122534.2442647 ],
           [107141.01841859],
```

[161017.30887996],
[184107.14764911],
[107141.01841859],
[145624.08303385],
[153320.69595691],
[168713.92180301],
[161017.30887996],
[130230.85718775],
[76354.56672638],
[84051.17964944],
[76354.56672638],
[76354.56672638],
[207196.98641827],
[199500.37349522],
[122534.2442647],
[199500.37349522],
[207196.98641827],
[214893.59934132],
[130230.85718775],
[91747.79257249],
[161017.30887996],
[161017.30887996],
[191803.76057217],
[168713.92180301],
[161017.30887996],
[76354.56672638],
[191803.76057217],
[145624.08303385],
[207196.98641827],
[199500.37349522],
[99444.40549554],
[153320.69595691],
[130230.85718775],
[91747.79257249],
[184107.14764911],
[76354.56672638],
[130230.85718775],
[191803.76057217],
[84051.17964944],
[207196.98641827],
[114837.63134164],
[176410.53472606],
[130230.85718775],
[122534.2442647],
[84051.17964944],
[137927.4701108],
[184107.14764911],

[76354.56672638],
[191803.76057217],
[168713.92180301],
[199500.37349522],
[137927.4701108],
[84051.17964944],
[214893.59934132],
[199500.37349522],
[122534.2442647],
[153320.69595691],
[76354.56672638],
[107141.01841859],
[130230.85718775],
[145624.08303385],
[214893.59934132],
[161017.30887996],
[137927.4701108],
[207196.98641827],
[99444.40549554],
[76354.56672638],
[191803.76057217],
[107141.01841859],
[91747.79257249],
[84051.17964944],
[145624.08303385],
[168713.92180301],
[199500.37349522],
[137927.4701108],
[184107.14764911],
[122534.2442647],
[199500.37349522],
[91747.79257249],
[199500.37349522],
[114837.63134164],
[107141.01841859],
[137927.4701108],
[153320.69595691],
[214893.59934132],
[184107.14764911],
[76354.56672638],
[168713.92180301],
[184107.14764911],
[184107.14764911],
[161017.30887996],
[207196.98641827],
[84051.17964944],
[145624.08303385],

```
[ 76354.56672638],  
[ 91747.79257249],  
[ 76354.56672638]])
```

```
[9]: mean_absolute_error(y_test, y_pred)
```

```
[9]: 7905.114728709234
```

```
[10]: root_mean_squared_error(y_test, y_pred)
```

```
[10]: 9691.918147584149
```

```
[11]: r2_score(y_test, y_pred)
```

```
[11]: 0.9573704756930143
```

```
[12]: lr.coef_
```

```
[12]: array([[7696.61292305]])
```

```
[13]: lr.intercept_
```

```
[13]: array([60961.34088028])
```

```
[14]: coefficients = lr.coef_  
      intercept = lr.intercept_  
      x = np.linspace(0, 20, 100)  
      y = coefficients*x + intercept  
  
      plt.scatter(x,y,label=f'y - {coefficients[0]} x + {intercept}')  
  
      plt.xlabel('Years of Experience')  
      plt.ylabel("Salary")  
      plt.title("Linear Regression Salary")  
  
      plt.show()
```



22-Multiple_Linear_Regression

October 20, 2024

1 Multiple Linear Regression

```
[1]: import random
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

data = []

for _ in range(500):
    team_name = f"Team {chr(random.randint(65,90))}"
    season = random.randint(2010, 2023)
    wins = random.randint(50, 110)
    losses = 162 - wins
    hits = random.randint(1200,1600)
    doubles = random.randint(200,350)
    triples = random.randint(10,40)
    home_runs = random.randint(100,250)
    strikeouts = random.randint(1000,1500)

    hits_adjusted = hits + (wins-80) * 5
    doubles_adjusted = doubles + (wins-80) * 2
    triples_adjusted = triples + (wins-80) * 3
    home_runs_adjusted = home_runs + (wins-80) * 3
    strikeouts_adjusted = strikeouts - (wins - 80) * 10

    data.append([team_name, season, wins, losses, hits_adjusted, r2_score,
r2_score, doubles_adjusted, triples_adjusted, home_runs_adjusted, strikeouts_adjusted])
```

```
[2]: columns = ["Team", 'Season', 'Wins',"Losses", 'Hits', 'Doubles', "Tripples",
↳ 'HomeRuns', 'Strikesouts']
```

```
df = pd.DataFrame(data, columns=columns)
df
```

```
[2]:
```

	Team	Season	Wins	Losses	Hits	Doubles	Tripples	HomeRuns	\
0	Team R	2018	68	94	1516	197	-17	172	
1	Team Y	2023	65	97	1344	275	-6	153	
2	Team U	2017	62	100	1457	197	-26	152	
3	Team U	2011	62	100	1182	233	-28	68	
4	Team N	2016	75	87	1248	233	0	194	
..	
495	Team D	2010	56	106	1106	152	-40	173	
496	Team O	2022	104	58	1587	310	93	307	
497	Team F	2019	106	56	1612	402	104	285	
498	Team Y	2010	64	98	1320	178	-30	126	
499	Team F	2019	55	107	1090	299	-41	87	


```

Strikesouts
0      1582
1      1495
2      1454
3      1271
4      1058
..      ...
495     1341
496     1130
497     1224
498     1204
499     1573
```

```
[500 rows x 9 columns]
```


```
[3]: sns.lmplot(x="Hits", y="Wins", data=df)
```

```
[3]: <seaborn.axisgrid.FacetGrid at 0x143f5ca4bd0>
```



```
[4]: sns.lmplot(x="Strikesouts", y="Wins", data=df)
```

```
[4]: <seaborn.axisgrid.FacetGrid at 0x143c0b5b450>
```



```
[5]: df2 = df.drop(columns=['Team', 'Season', 'Losses'], axis=1)
df2
```

```
[5]:
```

	Wins	Hits	Doubles	Tripples	HomeRuns	Strikesouts
0	68	1516	197	-17	172	1582
1	65	1344	275	-6	153	1495
2	62	1457	197	-26	152	1454
3	62	1182	233	-28	68	1271
4	75	1248	233	0	194	1058
..
495	56	1106	152	-40	173	1341
496	104	1587	310	93	307	1130
497	106	1612	402	104	285	1224
498	64	1320	178	-30	126	1204
499	55	1090	299	-41	87	1573

```
[500 rows x 6 columns]
```

```
[6]: df2.columns
```

```
[6]: Index(['Wins', 'Hits', 'Doubles', 'Tripples', 'HomeRuns', 'Strikesouts'],  
dtype='object')
```

```
[7]: X = df[['Hits', 'Doubles', 'Tripples', 'HomeRuns', 'Strikesouts']]  
y = df['Wins']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
↳random_state=24)  
  
lr = LinearRegression()  
lr.fit(X_train, y_train)
```

```
[7]: LinearRegression()
```

```
[8]: lr.score(X_test, y_test)
```

```
[8]: 0.9766876951321011
```

```
[9]: lr.score(X_train, y_train)
```

```
[9]: 0.9779524643162276
```

```
[10]: y_pred = lr.predict(X_test)  
y_pred
```

```
[10]: array([ 79.43521812,  84.98752505,  50.06532199,  86.3822643 ,  
        70.2017775 , 105.95777075, 112.95922339,  63.40753436,  
        69.69225294,  61.45407411, 109.65817756,  91.24687093,  
        53.51602938,  93.85145009,  52.19256777, 103.42618223,  
        93.19739631,  49.75559028,  69.68435674,  66.0810152 ,  
        55.59211481,  55.51175199,  83.75514524,  87.95757516,  
        80.68253472,  89.01173975,  66.41065628,  51.2704802 ,  
       114.21228759, 100.30212302,  88.33733575,  78.18776603,  
        84.57104209,  55.42473415,  68.36036516,  77.72199593,  
        63.81868048,  78.23857521,  95.88171008,  46.68880483,  
       106.20061061,  99.65684269, 102.86836133,  76.59226093,  
        97.43890582,  76.61837152,  67.35962044,  96.54023581,  
       98.99463936,  56.07128436,  73.84535217,  66.34627102,  
       100.5049504 ,  74.42990629,  66.88490465, 103.15534167,  
        97.23129157,  64.7906051 ,  69.02154043,  67.47991565,  
        62.22324859,  57.34789107, 103.81785868,  97.64835468,  
       77.99833509,  98.76947005, 106.69044356,  56.70169242,  
       111.37678841,  69.6311106 ,  99.53710032,  61.06668516,  
       104.96695125,  90.96773068,  96.64482364,  49.51410132,  
        67.4316584 ,  56.55871945, 101.21202627, 103.06747698,  
       88.72028133, 100.99416768,  90.87960032,  95.11780101,
```

```
105.01142639, 77.77341573, 57.13396359, 75.36364628,  
62.00493217, 66.91415242, 72.84128309, 97.89586341,  
109.90864271, 112.59300656, 76.58139883, 89.71476103,  
97.45697597, 80.89370933, 110.05559584, 56.32579865])
```

```
[11]: mean_absolute_error(y_test, y_pred)
```

```
[11]: 2.394583294287179
```

```
[12]: root_mean_squared_error(y_test, y_pred)
```

```
[12]: 2.7715279457881903
```

```
[13]: r2_score(y_test, y_pred)
```

```
[13]: 0.9766876951321011
```

```
[14]: lr.coef_
```

```
[14]: array([ 0.00362517,  0.00364042,  0.29551377,  0.01535465, -0.00312123])
```

```
[15]: lr.intercept_
```

```
[15]: 67.51826690267674
```

23-Lasso_Regression

October 20, 2024

1 Lasso Regression

Lasso (Least Absolute Shrinkage and Selection Operator) Regression is a type of linear regression that introduces a penalty equal to the absolute value of the magnitude of the coefficients. This penalty forces some of the coefficients to become exactly zero, effectively performing feature selection. It's a type of regularization technique used to prevent overfitting and simplify models by reducing the number of predictors.

1.0.1 Why is Lasso Regression Important?

- **Feature Selection:** One of the unique advantages of Lasso is that it can shrink some coefficients to zero, essentially performing automatic feature selection. This is especially useful when dealing with datasets that have many irrelevant or redundant features.
- **Prevents Overfitting:** In models with too many features, Lasso reduces overfitting by shrinking coefficients, improving the model's ability to generalize on unseen data.
- **Sparse Models:** By reducing coefficients to zero, Lasso simplifies the model, creating a sparse solution. Sparse models are easier to interpret and often have better prediction accuracy.
- **Works well with high-dimensional data:** In cases where the number of predictors exceeds the number of observations, Lasso is particularly useful because it selects only the most important predictors.

1.0.2 How does Lasso Regression work?

The Lasso regression process follows these steps:

1. **Initialization:** Start by defining a linear regression model with a penalty on the magnitude of the coefficients. This penalty is controlled by a hyperparameter λ .
2. **Regularization:** Introduce the L1 penalty, which is the sum of the absolute values of the regression coefficients. Unlike traditional linear regression, which only minimizes the sum of squared errors, Lasso also minimizes the total size of the coefficients.
3. **Optimization:** Use optimization algorithms like coordinate descent to solve the Lasso objective function. The solver adjusts the coefficients to find the balance between fitting the data and shrinking some coefficients to zero.
4. **Coefficient Shrinkage:** Depending on the value of λ , some of the coefficients may be shrunk to exactly zero. The larger the value of λ , the stronger the regularization effect, meaning more features will be set to zero.

5.Prediction: After training, the resulting model can be used to predict outcomes based on the reduced set of features.

1.0.3 When should you use Lasso Regression?

- **High-dimensional datasets:** When you have more features than observations, Lasso is a good choice because it selects a subset of the most important features, reducing the complexity of the model.
- **Feature Selection:** If you suspect that many of the features in your dataset are irrelevant, Lasso can help automatically select the relevant ones.
- **Overfitting prevention:** When you notice that your model is performing well on training data but poorly on validation or test data, Lasso can be a good solution as it reduces the chance of overfitting by regularizing the coefficients.
- **Sparse solutions:** If you need an interpretable model with fewer predictors, Lasso is suitable because it produces sparse models where some coefficients are zero.

1.0.4 Who uses Lasso Regression?

- **Data Scientists and Machine Learning Engineers:** Especially when working on predictive modeling tasks that involve many features, Lasso is a popular tool for improving model generalizability and interpretability.
- **Researchers:** In fields like genomics, where the number of predictors can be vast, Lasso is frequently used to select important variables while disregarding the rest.
- **Statisticians:** Those dealing with linear regression models that have multicollinearity (high correlation between features) use Lasso to shrink the coefficients and combat the effects of collinearity.

1.0.5 Key Points to Remember:

- The strength of the regularization is controlled by the hyperparameter (or alpha in Python's scikit-learn). The larger the value, the stronger the regularization.
- It's a biased estimator, meaning it sacrifices some bias to reduce variance, leading to a better performance on test data (especially when overfitting is a concern).
- Lasso selects features and shrinks coefficients, making it a good choice for interpretable models.

```
[1]: import pandas as pd
import numpy as np

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score
```



```

ca_housing = fetch_california_housing()

X = ca_housing.data
y = ca_housing.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=19)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)

```

```

[2]: lasso = Lasso()
lasso.fit(X_train, y_train)

```

```

[2]: Lasso()

```

```

[3]: y_pred = lasso.predict(X_test)
y_pred

```

```

[3]: array([2.06708162, 2.06708162, 2.06708162, ..., 2.06708162, 2.06708162,
2.06708162])

```

```

[4]: mean_absolute_error(y_test, y_pred)

```

```

[4]: 0.9119430573559844

```

```

[5]: root_mean_squared_error(y_test, y_pred)

```

```

[5]: 1.1517022831067947

```

```

[6]: r2_score(y_test, y_pred)

```

```

[6]: -4.109353628090062e-05

```

```

[7]: param_grid = {
    'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]
}

lasso_cv = GridSearchCV(lasso, param_grid, cv=3, n_jobs=-1)
lasso_cv.fit(X_train, y_train)

```

```

[7]: GridSearchCV(cv=3, estimator=Lasso(), n_jobs=-1,
    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]})

```

```
[8]: y_pred2 = lasso_cv.predict(X_test)
      y_pred2
```

```
[8]: array([2.72485002, 0.66011565, 2.12095357, ..., 1.92427048, 1.88073015,
          2.3150062 ])
```

```
[9]: mean_absolute_error(y_test, y_pred2)
```

```
[9]: 0.5353500594399226
```

```
[10]: root_mean_squared_error(y_test, y_pred2)
```

```
[10]: 0.7218861196454859
```

```
[11]: r2_score(y_test, y_pred2)
```

```
[11]: 0.60710656378633
```

```
[12]: lasso_cv.best_estimator_
```

```
[12]: Lasso(alpha=0.001)
```

```
[13]: lasso3 = Lasso(alpha=0.001)
      lasso3.fit(X_train, y_train)
```

```
[13]: Lasso(alpha=0.001)
```

```
[14]: lasso3.intercept_
```

```
[14]: 2.0670816194279977
```

```
[15]: lasso3.coef_
```

```
[15]: array([ 0.83673788,  0.12126534, -0.26089701,  0.30370697, -0.00173652,
          -0.02849403, -0.8865986 , -0.86020295])
```

```
[16]: feature_names =
      ↪ ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']

      df = pd.DataFrame({'Feature_Names': feature_names, 'Coef': lasso3.coef_})
      df
```

```
[16]:   Feature_Names      Coef
0      MedInc    0.836738
1    HouseAge    0.121265
2    AveRooms  -0.260897
3    AveBedrms    0.303707
4  Population  -0.001737
```

5	AveOccup	-0.028494
6	Latitude	-0.886599
7	Longitude	-0.860203

24-Ridge_Regression

October 20, 2024

1 Ridge Regression

Ridge Regression is a type of linear regression that adds a regularization term to the model to penalize large coefficients. Unlike ordinary linear regression, which only minimizes the sum of squared errors, Ridge Regression minimizes both the squared errors and a penalty term proportional to the square of the magnitude of the coefficients. This penalty helps reduce model complexity and prevent overfitting.

In Ridge Regression, the penalty term is based on the L2 norm (the square of the coefficients), which encourages small but non-zero coefficients. Unlike Lasso, Ridge does not shrink coefficients to zero; rather, it shrinks them towards zero to reduce variance.

1.0.1 Why is Ridge Regression Important?

- **Multicollinearity Solution:** Ridge Regression is particularly useful when there is multicollinearity (high correlation between independent variables). In such cases, ordinary least squares (OLS) can give unreliable estimates. Ridge stabilizes the regression by adding a penalty that discourages large coefficient values.
- **Overfitting Prevention:** By introducing a regularization term, Ridge Regression controls overfitting, ensuring the model generalizes better to unseen data.
- **Handling High-Dimensional Data:** Ridge works well when the number of features (predictors) is large, especially when the number of features exceeds the number of observations.
- **Interpretability:** While Ridge does not perform feature selection like Lasso, it can still simplify models by shrinking coefficients to smaller values, making the model more interpretable.

1.0.2 How does Ridge Regression work?

- **Define the objective:** Ridge minimizes the sum of squared errors while adding an L2 penalty (the sum of squared coefficients).
- **Choose regularization parameter (λ):** The λ controls the amount of regularization. If $\lambda = 0$, Ridge behaves like standard linear regression. As λ increases, the regularization effect becomes stronger, shrinking the coefficients.
- **Fit the model:** Solving the Ridge Regression objective involves minimizing the cost function, which can be efficiently done using optimization algorithms like gradient descent or closed-form solutions. For Ridge, there is an exact solution:

- **Shrinkage of coefficients:** As λ increases, the model becomes more conservative, shrinking the coefficients to smaller values, but never exactly zero as in Lasso.
- **Prediction:** Once the coefficients are trained, the model can predict outcomes for new data points using the regularized coefficients.

1.0.3 When should you use Ridge Regression?

- **Multicollinearity:** Ridge is ideal when independent variables are highly correlated. It helps by shrinking coefficients, which reduces the variance that can result from multicollinearity.
- **High-dimensional data:** When the dataset has many predictors (especially when the number of predictors exceeds the number of observations), Ridge is more stable than ordinary least squares regression.
- **When all features are important:** If you believe all the features contribute to the target variable (even to a small degree) and you don't want to discard any, Ridge is preferred over Lasso, which might zero out some coefficients.
- **Balancing bias and variance:** Ridge strikes a balance between variance (how much your model changes with different training data) and bias (error introduced by approximating a complex model). It controls variance by shrinking coefficients without introducing significant bias.

1.0.4 Who uses Ridge Regression?

- **Data Scientists and Machine Learning Engineers:** Ridge is a go-to algorithm when building predictive models that need to handle multicollinearity or high-dimensional datasets.
- **Economists and Statisticians:** In fields where datasets often have correlated variables (like economic indicators), Ridge is commonly used to produce more stable and interpretable models.
- **Genomics and Bioinformatics:** Ridge is useful in genomics, where datasets often have a large number of highly correlated predictors (genes), and overfitting is a concern.

1.0.5 Key Differences Between Ridge and Lasso:

- Ridge shrinks coefficients, but it does not set them to exactly zero, meaning all features remain in the model (though their contributions may be small).
- Lasso, on the other hand, can shrink some coefficients to zero, effectively removing features from the model.

1.0.6 Key Points to Remember:

- **Regularization:** The key feature of Ridge is that it adds an L2 penalty (squared coefficients) to the loss function. This penalty reduces the model's complexity by shrinking the coefficients.
- **Bias-Variance Tradeoff:** Ridge Regression helps reduce variance (thus improving generalization) at the cost of introducing a small bias. **No Feature Selection:** Unlike Lasso, Ridge does not perform feature selection. All features are kept in the model but are shrunk towards zero.

```
[1]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import Ridge
from sklearn.datasets import make_regression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

X, y = make_regression(n_samples=100, n_features=4, noise=1, random_state=42 ,
    effective_rank=2)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=19)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)
```

```
[2]: ridge = Ridge()
ridge.fit(X_train, y_train)
```

```
[2]: Ridge()
```

```
[3]: y_pred = ridge.predict(X_test)
y_pred
```

```
[3]: array([ 7.13539375, -8.8002492 ,  8.78912513,  6.09927205,
 5.99994735, -12.30572261, -10.03631065, -1.04649374,
23.52913328,  2.55203747, 11.48889191, -7.61522827,
 0.46017582, -8.62564635, -29.50846326, -7.76254477,
 8.06258548, -7.44014532, -3.84252421, 25.8910913 ])
```

```
[4]: mean_absolute_error(y_test, y_pred)
```

```
[4]: 3.0816824025813956
```

```
[5]: root_mean_squared_error(y_test, y_pred)
```

```
[5]: 11.330220034384492
```

```
[6]: r2_score(y_test, y_pred)
```

```
[6]: 0.9276234847540757
```

```
[7]: param_grid = {  
      'alpha': [ 0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0]  
    }  
  
    ridge_cv = GridSearchCV(ridge, param_grid, cv=3, n_jobs=-1)  
    ridge_cv.fit(X_train, y_train)
```

```
[7]: GridSearchCV(cv=3, estimator=Ridge(), n_jobs=-1,  
      param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0]})
```

```
[8]: y_pred2 = ridge_cv.predict(X_test)  
    y_pred2
```

```
[8]: array([ 7.17119326, -8.89540639,  8.83622445,  6.17650173,  
          6.08043617, -12.40702852, -10.10162896, -1.07770549,  
          23.78518149,  2.57122983, 11.59362725, -7.71459096,  
          0.4722029 , -8.69228577, -29.79550622, -7.85527431,  
          8.12657042, -7.53080381, -3.87581506, 26.15720314])
```

```
[9]: mean_absolute_error(y_test, y_pred2)
```

```
[9]: 3.0816824025813956
```

```
[10]: root_mean_squared_error(y_test, y_pred2)
```

```
[10]: 11.349061639991897
```

```
[11]: r2_score(y_test, y_pred2)
```

```
[11]: 0.9275031261245554
```

```
[12]: ridge_cv.best_estimator_
```

```
[12]: Ridge(alpha=0.0001)
```

```
[13]: ridge_cv.best_estimator_.intercept_
```

```
[13]: 0.15121625800403157
```

```
[14]: ridge_cv.best_estimator_.coef_
```

```
[14]: array([4.3844567 , 5.0806461 , 7.6664195 , 2.89485579])
```

25-Elastic_Net_Regression

October 20, 2024

1 Elastic Net Regression

Elastic Net Regression is a regularization technique that combines both Lasso Regression and Ridge Regression. It introduces penalties to the model, balancing the L1 norm (from Lasso) and the L2 norm (from Ridge) to create a more flexible regularization method. This helps improve predictive accuracy, deal with multicollinearity, and perform automatic feature selection.

Elastic Net allows for the benefits of both Lasso and Ridge:

- **L1 penalty (Lasso)** encourages sparse solutions by shrinking some coefficients to exactly zero, thus selecting features.
- **L2 penalty (Ridge)** shrinks coefficients but keeps them non-zero, thus stabilizing the model when predictors are highly correlated.

1.0.1 Why is Elastic Net Regression Important?

- **Combining Lasso and Ridge:** In situations where Lasso struggles (e.g., highly correlated features), Ridge can compensate. Elastic Net creates a balance between the two, offering a robust model that benefits from both feature selection and coefficient shrinkage.
- **Handles Multicollinearity:** Elastic Net is particularly useful when there are correlated predictors. It keeps Ridge's ability to handle multicollinearity while also keeping Lasso's feature selection capability.
- **Feature Selection:** Like Lasso, Elastic Net can shrink coefficients to zero, effectively eliminating irrelevant features from the model, which makes it interpretable and prevents overfitting.
- **Flexible Regularization:** It introduces two penalties (L1 and L2), giving you the flexibility to fine-tune the degree of regularization needed for the data, based on the dataset's complexity and feature correlation.

1.0.2 How does Elastic Net Regression work?

1. Initialization: Start with the linear regression objective but add two penalty terms: one for Lasso (L1 norm) and one for Ridge (L2 norm).

2. Blending penalties: The model combines the advantages of both regularization techniques. The L1 penalty encourages sparsity by setting some coefficients to zero, while the L2 penalty helps when multicollinearity exists by shrinking the coefficients but keeping them non-zero.

3. Hyperparameter tuning: Elastic Net has two hyperparameters, (Lasso penalty strength) and (Ridge penalty strength). These are often tuned using cross-validation to find the optimal values that prevent overfitting while improving model performance.

4. Optimization: Elastic Net uses algorithms like coordinate descent to solve the objective function, iteratively adjusting the coefficients until an optimal solution is reached.

5. Prediction: After training, the model can predict outcomes based on the selected features and shrunk coefficients.

1.0.3 When should you use Elastic Net Regression?

- **When features are highly correlated:** Lasso struggles with groups of highly correlated features by selecting only one and ignoring the others. Elastic Net, by blending Ridge, can distribute the effect across correlated features.
- **High-dimensional data:** In datasets where the number of features is much larger than the number of observations (e.g., in genomics or text classification), Elastic Net is helpful because it performs both feature selection (like Lasso) and shrinks the coefficients (like Ridge).
- **Sparse models:** If you expect only a subset of your predictors to be relevant, Elastic Net can shrink irrelevant predictors to zero, like Lasso.
- **Handling multicollinearity:** When predictors are highly correlated, Elastic Net is better suited than Lasso alone, since Ridge helps stabilize coefficient estimates in the presence of collinearity.

1.0.4 Who uses Elastic Net Regression?

- **Data Scientists and Machine Learning Practitioners:** Elastic Net is a go-to choice when modeling datasets with many features or when there is multicollinearity, especially in high-stakes predictive tasks.
- **Bioinformaticians and Geneticists:** In fields like genomics, where there are many correlated features (genes), Elastic Net helps in selecting key predictors while accounting for correlations.
- **Economists and Statisticians:** Elastic Net is useful in domains where multicollinearity and high-dimensional datasets are common, as it balances the trade-off between bias and variance.

1.0.5 Key Points to Remember:

- **L1 and L2 regularization:** Elastic Net combines L1 (Lasso) and L2 (Ridge) regularization, balancing feature selection and coefficient shrinkage.
- **Tuning and l1_ratio:** The mix between Lasso and Ridge penalties is controlled by the `l1_ratio` parameter. An `alpha` value helps control the overall strength of the regularization.
- **No perfect separation:** While Lasso is more aggressive in shrinking coefficients to zero, Elastic Net allows coefficients to stay small but non-zero when dealing with correlated features.

1.0.6 Differences Between Elastic Net, Lasso, and Ridge:

- **Lasso (only L1 penalty):** Performs automatic feature selection by shrinking some coefficients to zero, but struggles with correlated features.
- **Ridge (only L2 penalty):** Shrinks coefficients uniformly without zeroing them out, handling multicollinearity but not performing feature selection.
- **Elastic Net:** Balances the two by performing feature selection like Lasso but also shrinking correlated features together like Ridge.

```
[1]: import pandas as pd
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import ElasticNet
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, \
    r2_score

tips = sns.load_dataset('tips')
tips
```

```
[1]:      total_bill  tip  sex smoker  day  time  size
0         16.99  1.01 Female    No  Sun  Dinner     2
1         10.34  1.66  Male    No  Sun  Dinner     3
2         21.01  3.50  Male    No  Sun  Dinner     3
3         23.68  3.31  Male    No  Sun  Dinner     2
4         24.59  3.61 Female    No  Sun  Dinner     4
..          ...   ...   ...   ...   ...   ...   ...
239        29.03  5.92  Male    No  Sat  Dinner     3
240        27.18  2.00 Female   Yes  Sat  Dinner     2
241        22.67  2.00  Male   Yes  Sat  Dinner     2
242        17.82  1.75  Male    No  Sat  Dinner     2
243        18.78  3.00 Female    No  Thur Dinner     2
```

[244 rows x 7 columns]

```
[2]: tips = pd.get_dummies(tips)
tips
```

```
[2]:      total_bill  tip  size  sex_Male  sex_Female  smoker_Yes  smoker_No  \
0         16.99  1.01     2     False        True        False         True
1         10.34  1.66     3        True        False        False         True
2         21.01  3.50     3        True        False        False         True
3         23.68  3.31     2        True        False        False         True
4         24.59  3.61     4     False        True        False         True
..          ...   ...   ...   ...   ...   ...   ...
239        29.03  5.92     3        True        False        False         True
```

240	27.18	2.00	2	False	True	True	False
241	22.67	2.00	2	True	False	True	False
242	17.82	1.75	2	True	False	False	True
243	18.78	3.00	2	False	True	False	True

	day_Thur	day_Fri	day_Sat	day_Sun	time_Lunch	time_Dinner
0	False	False	False	True	False	True
1	False	False	False	True	False	True
2	False	False	False	True	False	True
3	False	False	False	True	False	True
4	False	False	False	True	False	True
..
239	False	False	True	False	False	True
240	False	False	True	False	False	True
241	False	False	True	False	False	True
242	False	False	True	False	False	True
243	True	False	False	False	False	True

[244 rows x 13 columns]

```
[3]: X = tips.drop('tip', axis=1)
      y = tips['tip']

      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
      random_state=19)

      scaler = StandardScaler()

      X_train = scaler.fit_transform(X_train)
      X_test = scaler.fit_transform(X_test)
```

```
[4]: elastic_net = ElasticNet()
      elastic_net.fit(X_train, y_train)
```

```
[4]: ElasticNet()
```

```
[5]: y_pred = elastic_net.predict(X_test)
      y_pred
```

```
[5]: array([3.22787603, 2.75553242, 2.98583549, 2.93657376, 2.81653161,
           2.85814622, 3.02869499, 3.33226823, 2.85601214, 2.97872189,
           2.80728392, 3.14873713, 3.52149022, 2.91701133, 2.92928231,
           2.84907637, 2.89904947, 2.80177087, 2.74681825, 3.23410043,
           3.47738585, 2.85654566, 2.79696919, 2.72014221, 2.85138829,
           2.94262032, 2.94564361, 2.86348143, 2.97498724, 2.942976 ,
           3.1030322 , 2.94422088, 2.80159303, 2.79732487, 2.91896757,
           2.95062313, 3.06248463, 2.93408399, 2.79696919, 2.93070503,
```

```
2.89691539, 2.80159303, 3.20173352, 2.76922611, 2.90313979,  
3.01784673, 2.94937825, 2.86810527, 3.47454041])
```

```
[6]: mean_absolute_error(y_test, y_pred)
```

```
[6]: 1.214329254741493
```

```
[7]: root_mean_squared_error(y_test, y_pred)
```

```
[7]: 2.899097972303773
```

```
[8]: r2_score(y_test, y_pred)
```

```
[8]: 0.14011647876429623
```

```
[9]: param_grid = {  
      'alpha': [ 0.1,0.3,0.5,0.7,0.9,1.0 ],  
      'l1_ratio': [ 0.1,0.3,0.5,0.7,0.9,1.0 ]  
    }  
  
    elastic_net_cv = GridSearchCV(elastic_net, param_grid, cv=3,  
    ↪scoring='neg_root_mean_squared_error', n_jobs=-1)  
    elastic_net_cv.fit(X_train, y_train)
```

```
[9]: GridSearchCV(cv=3, estimator=ElasticNet(), n_jobs=-1,  
      param_grid={'alpha': [0.1, 0.3, 0.5, 0.7, 0.9, 1.0],  
                  'l1_ratio': [0.1, 0.3, 0.5, 0.7, 0.9, 1.0]} ,  
      scoring='neg_root_mean_squared_error')
```

```
[10]: y_pred2 = elastic_net_cv.predict(X_test)  
      y_pred2
```

```
[10]: array([4.36045835, 2.01577541, 3.49350917, 2.6632068 , 2.43904349,  
            2.39559597, 3.21389761, 4.37446046, 2.65737098, 2.7356982 ,  
            2.40921402, 4.13318475, 4.6636574 , 2.67645063, 2.83025002,  
            2.48055813, 2.47870584, 2.54539045, 2.12747365, 4.0536548 ,  
            4.76205469, 2.48556614, 2.21289653, 2.15979944, 2.52823827,  
            2.84575684, 3.18666694, 2.56724603, 2.9643126 , 2.73268614,  
            3.28063625, 3.04227102, 2.65903479, 2.40876079, 2.89783283,  
            2.87108944, 3.6777688 , 2.65517579, 2.19874392, 2.82147533,  
            2.61162889, 2.39085742, 3.45504734, 2.28645428, 2.73256056,  
            3.06932308, 3.04475407, 2.58216077, 5.2470807 ])
```

```
[11]: mean_absolute_error(y_test, y_pred2)
```

```
[11]: 0.9467399418855579
```

```
[12]: root_mean_squared_error(y_test, y_pred2)
```

[12]: 1.9034103248341647

[13]: `r2_score(y_test, y_pred2)`

[13]: 0.43544123444224914

[14]: `elastic_net_cv.best_params_`

[14]: {'alpha': 0.1, 'l1_ratio': 0.1}

[15]: `elastic_net_cv.best_estimator_`

[15]: ElasticNet(alpha=0.1, l1_ratio=0.1)

[16]: `elastic_net_cv.best_score_`

[16]: -0.9574652456524534

26-Stacking_Regressor

October 20, 2024

1 Stacking Regression

A **Stacking Regressor** is an ensemble machine learning technique that combines the predictions of multiple regression models (base models) by training a meta-model (stacker) to make final predictions based on the outputs of those base models. Stacking, also known as stacked generalization, aims to leverage the strengths of several models to improve predictive performance compared to using any single model.

The key idea of stacking is that different models may excel in capturing different aspects of the data. By training a second-level model (meta-model) on the outputs of base models, stacking allows the meta-model to learn how to best combine the predictions of the base models.

1.0.1 Why is Stacking Regressor Important?

- **Boosts model performance:** Since stacking combines the predictive power of several models, it often yields better performance than individual models. It uses the strengths of each model while compensating for their weaknesses.
- **Improves generalization:** Stacking helps in reducing overfitting by leveraging a blend of multiple models. The meta-model is trained to understand which base models perform best under different circumstances.
- **Flexibility:** You can use a wide range of base learners (e.g., decision trees, linear regression, support vector machines, etc.) and combine their predictions in a meaningful way, allowing for diverse model architectures.

1.0.2 How does Stacking Regressor work?

Stacking involves two layers:

1. **Base Models (Level 0):** These are the individual regression models that make predictions on the dataset. Each base model is trained on the original features of the data.
2. **Meta-model (Level 1):** The meta-model (also called the blender or stacker) is trained on the predictions of the base models. It learns to combine these predictions to generate the final output.

Here's a step-by-step breakdown of how stacking works:

1. **Train base models:** Each base model is trained on the training data. These models could be decision trees, random forests, gradient boosting, linear regression, etc.

2. **Generate base model predictions:** Once trained, the base models are used to make predictions on both the training data (for cross-validation) and the test data. These predictions are then used as inputs to the meta-model.
3. **Train the meta-model:** A second model (meta-model), often a simple linear regression or another powerful algorithm, is trained using the predictions from the base models as features. This model learns to combine the outputs of the base models to improve predictive performance.
4. **Final prediction:** During the prediction phase, the base models make predictions on new data, and these predictions are fed into the meta-model, which then produces the final prediction.

1.0.3 When should you use Stacking Regressor?

- **When base models have different strengths:** If your base models have different inductive biases and capture different patterns in the data (e.g., decision trees might capture non-linear relationships, while linear models capture linear trends), stacking can help by learning how to best combine them.
- **High complexity datasets:** Stacking can be useful when the dataset is complex and no single model is able to fully capture the data's underlying patterns. By combining multiple models, stacking often performs better in these cases.
- **To reduce overfitting:** Stacking helps in preventing overfitting by blending multiple models. If individual models tend to overfit, the meta-model can act as a stabilizing factor by learning to trust the better-performing models.
- **When you want to combine models with different architectures:** Stacking gives the flexibility to combine a variety of different models (e.g., tree-based models, linear models, neural networks, etc.) in one powerful ensemble.

1.0.4 Who uses Stacking Regressor?

- **Data Scientists and Machine Learning Engineers:** Stacking is a common technique used in competitions like Kaggle and real-world applications where predictive performance is critical. It's widely adopted when optimizing models for predictive accuracy.
- **Financial Analysts and Economists:** Stacking is used in scenarios such as stock market predictions or economic forecasting, where combining models that capture different signals (e.g., time-series patterns, economic indicators, etc.) can enhance the overall predictive power.
- **Researchers and AI Developers:** Stacking is often used in research when building complex models that need to integrate multiple types of predictions (e.g., biological data analysis, medical imaging, etc.).

1.0.5 Key Points to Remember:

- **Base Models and Meta-model:** The base models can be diverse (linear models, tree-based models, SVM, etc.), and the meta-model is trained on their predictions to combine them effectively.

- **Cross-Validation:** Stacking usually involves cross-validation to avoid overfitting. The meta-model should be trained on out-of-sample predictions to ensure it generalizes well.
- **Customizability:** You can choose any combination of models for the base learners and meta-model depending on the problem and data structure.

1.0.6 How does Stacking differ from Bagging and Boosting?

- **Bagging:** Bagging (e.g., Random Forest) trains multiple models in parallel on different subsets of the data and averages their predictions to reduce variance.
- **Boosting:** Boosting (e.g., Gradient Boosting) trains models sequentially, where each subsequent model attempts to correct the mistakes of the previous one.
- **Stacking:** Stacking trains multiple models in parallel (like Bagging) but uses another model to learn how to best combine the outputs of the base models.
- Lasso Regression captures important features by shrinking irrelevant ones.
- Ridge Regression handles multicollinearity and smooths out feature effects.
- Support Vector Regression captures complex, non-linear relationships.

```
[1]: import pandas as pd
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV, \
    RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import StackingRegressor, RandomForestRegressor, \
    GradientBoostingRegressor, VotingRegressor
from sklearn.linear_model import Ridge, Lasso, LinearRegression
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, \
    r2_score
from sklearn.svm import SVR

mpg = sns.load_dataset('mpg')
mpg
```

```
[1]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    18.0          8         307.0         130.0   3504          12.0
1    15.0          8         350.0         165.0   3693          11.5
2    18.0          8         318.0         150.0   3436          11.0
3    16.0          8         304.0         150.0   3433          12.0
4    17.0          8         302.0         140.0   3449          10.5
..    ...          ...          ...          ...    ...          ...
393  27.0          4         140.0          86.0   2790          15.6
394  44.0          4          97.0          52.0   2130          24.6
395  32.0          4         135.0          84.0   2295          11.6
396  28.0          4         120.0          79.0   2625          18.6
397  31.0          4         119.0          82.0   2720          19.4
```


	model_year	origin	name
0	70	usa	chevrolet chevelle malibu
1	70	usa	buick skylark 320
2	70	usa	plymouth satellite
3	70	usa	amc rebel sst
4	70	usa	ford torino
..
393	82	usa	ford mustang gl
394	82	europe	vw pickup
395	82	usa	dodge rampage
396	82	usa	ford ranger
397	82	usa	chevy s-10

[398 rows x 9 columns]

```
[2]: mpg = mpg.drop('name', axis=1)
      mpg = pd.get_dummies(mpg)
      mpg
```

```
[2]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    18.0           8         307.0         130.0    3504           12.0
1    15.0           8         350.0         165.0    3693           11.5
2    18.0           8         318.0         150.0    3436           11.0
3    16.0           8         304.0         150.0    3433           12.0
4    17.0           8         302.0         140.0    3449           10.5
..    ...           ...           ...           ...           ...           ...
393  27.0           4         140.0          86.0    2790           15.6
394  44.0           4          97.0          52.0    2130           24.6
395  32.0           4         135.0          84.0    2295           11.6
396  28.0           4         120.0          79.0    2625           18.6
397  31.0           4         119.0          82.0    2720           19.4
```

	model_year	origin_europe	origin_japan	origin_usa
0	70	False	False	True
1	70	False	False	True
2	70	False	False	True
3	70	False	False	True
4	70	False	False	True
..
393	82	False	False	True
394	82	True	False	False
395	82	False	False	True
396	82	False	False	True
397	82	False	False	True

[398 rows x 10 columns]

```
[3]: pd.DataFrame(mpg.isnull().sum().sort_values(ascending=False))
```

```
[3]:
horsepower    0
mpg            0
cylinders      0
displacement  0
weight         0
acceleration  0
model_year    0
origin_europe 0
origin_japan  0
origin_usa    0
```

```
[4]: mpg['horsepower'].fillna(mpg['horsepower'].mean(), inplace=True)
```

C:\Users\ikiga\AppData\Local\Temp\ipykernel_26612\370553803.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
mpg['horsepower'].fillna(mpg['horsepower'].mean(), inplace=True)
```

```
[5]: mpg
```

```
[5]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	
...	
393	27.0	4	140.0	86.0	2790	15.6	
394	44.0	4	97.0	52.0	2130	24.6	
395	32.0	4	135.0	84.0	2295	11.6	
396	28.0	4	120.0	79.0	2625	18.6	
397	31.0	4	119.0	82.0	2720	19.4	
	model_year	origin_europe	origin_japan	origin_usa			
0	70	False	False	True			
1	70	False	False	True			

2	70	False	False	True
3	70	False	False	True
4	70	False	False	True
..
393	82	False	False	True
394	82	True	False	False
395	82	False	False	True
396	82	False	False	True
397	82	False	False	True

[398 rows x 10 columns]

```
[6]: X = mpg.drop('mpg', axis=1)
      y = mpg['mpg']

      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,
      ↪random_state=19)
```

1.1 Linear Regression

```
[7]: lr = LinearRegression()
      lr.fit(X_train, y_train)
```

```
[7]: LinearRegression()
```

```
[8]: y_pred = lr.predict(X_test)
      y_pred
```

```
[8]: array([24.30115472, 24.73892314, 24.96215822, 17.38881632, 15.86675499,
          27.76426172, 17.05954093, 10.49344079, 20.347386 , 25.11254933,
          19.36068674, 28.56363082, 26.2720292 , 17.73036858, 20.66471276,
          28.7669855 , 28.98129852, 29.25304272, 16.8393709 ,  8.76273604,
          16.24941413, 13.29303025, 27.03902874, 23.9878045 ,  8.31184254,
          31.52160011, 13.42789179, 35.09654981, 11.22026997, 23.39755039,
          33.81737584, 29.59696005, 15.57771451, 23.05695715, 31.56588099,
          36.11360599, 27.33151931, 12.33155235, 28.54718257, 15.3977963 ,
          31.21498729, 23.50190767, 14.44084076, 27.66645708, 26.78275774,
          25.05523353, 24.20075891, 16.30434811, 15.63453382, 12.61013012,
          17.36880625, 20.27285977, 30.10153781, 16.80607781, 33.29417373,
          11.53912186, 32.05759396, 21.35355061, 18.08124065, 30.24876403,
          20.89162341, 20.33755243, 10.89210702, 35.17824516, 32.51214434,
          15.01604073, 25.98942786, 15.12905345, 32.04280232, 24.17805781,
          29.82643316, 22.55608296, 14.09889069, 14.59105924, 26.25541638,
          23.19242234, 12.07920811, 33.66219498, 29.56124407, 20.90794551])
```

```
[9]: mean_absolute_error(y_test, y_pred)
```

```
[9]: 2.3241334520650594
```

```
[10]: root_mean_squared_error(y_test, y_pred)
```

```
[10]: 8.30859663913005
```

```
[11]: r2_score(y_test, y_pred)
```

```
[11]: 0.8325409560733986
```

1.2 Random Forest Regressor

```
[12]: rfr = RandomForestRegressor(random_state=13)  
rfr.fit(X_train, y_train)
```

```
[12]: RandomForestRegressor(random_state=13)
```

```
[13]: y_pred_rfr = rfr.predict(X_test)  
y_pred_rfr
```

```
[13]: array([24.91 , 26.604, 24.747, 15.894, 15.851, 25.264, 18.296, 13.44 ,  
        18.43 , 25.158, 17.164, 32.449, 26.397, 16.445, 19.475, 26.181,  
        28.641, 33.474, 15.54 , 12.64 , 15.985, 14.75 , 24.137, 22.409,  
        12.15 , 30.689, 13.895, 37.527, 13.52 , 21.971, 35.155, 27.763,  
        15.234, 26.289, 33.153, 36.126, 23.992, 13.88 , 35.015, 14.225,  
        31.179, 25.104, 15.482, 31.032, 25.78 , 26.102, 21.299, 16.106,  
        14.604, 14.18 , 16.227, 18.722, 35.506, 16.804, 34.797, 13.82 ,  
        32.936, 18.86 , 18.847, 36.382, 15.164, 21.129, 13.38 , 36.05 ,  
        33.419, 16.248, 24.139, 16.082, 32.642, 23.887, 30.994, 20.967,  
        14.33 , 15.03 , 24.663, 21.33 , 14.09 , 38.027, 31.285, 19.26 ])
```

```
[14]: mean_absolute_error(y_test, y_pred_rfr)
```

```
[14]: 1.6326000000000005
```

```
[15]: root_mean_squared_error(y_test, y_pred_rfr)
```

```
[15]: 5.647237075000005
```

```
[16]: r2_score(y_test, y_pred_rfr)
```

```
[16]: 0.8861804270347423
```

1.3 Ridge Regression

```
[17]: ridge = Ridge()
      param_grid = {
          'alpha': [0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 50, 75]
      }

      ridge_cv = GridSearchCV(ridge, param_grid, cv=5, n_jobs=-1)
      ridge_cv.fit(X_train, y_train)
```

```
[17]: GridSearchCV(cv=5, estimator=Ridge(), n_jobs=-1,
                  param_grid={'alpha': [0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 50,
                                         75]}))
```

```
[18]: ridge_cv.best_estimator_
```

```
[18]: Ridge(alpha=10)
```

```
[19]: y_pred_ridge = ridge_cv.predict(X_test)
      y_pred_ridge
```

```
[19]: array([24.50784329, 25.1453431 , 24.53292975, 17.28315713, 15.76458563,
          27.56386969, 17.08360929, 10.4757291 , 20.27551626, 25.32880598,
          19.37647572, 28.86733796, 26.62154492, 17.73752567, 20.67788725,
          28.61052096, 29.25178742, 29.5422219 , 16.87281192,  8.53318336,
          16.27511627, 13.2104424 , 26.88340032, 24.009771 ,  8.54411519,
          31.26528627, 13.51952031, 34.90097167, 10.98124741, 23.16980535,
          33.58843945, 29.44372653, 15.49893771, 23.46053976, 31.4803678 ,
          35.94195329, 27.1668034 , 12.13433213, 28.84816147, 15.45947453,
          30.97278809, 23.81928018, 14.51031251, 27.2552616 , 26.63822582,
          25.25480835, 24.33315835, 16.31946596, 15.52974001, 12.67976987,
          17.40627283, 20.26742196, 30.05946769, 16.89475666, 33.16874796,
          11.53698617, 31.9061312 , 21.40546978, 18.10722736, 30.55085517,
          20.62056802, 20.53397226, 10.65600453, 35.01686624, 32.35642928,
          15.05582844, 25.86631482, 15.21926104, 31.87832831, 24.08914757,
          29.75506395, 22.43481848, 14.10345592, 14.38519282, 26.3607398 ,
          23.43599278, 12.1534485 , 33.50094882, 29.46266978, 20.95613101])
```

```
[20]: mean_absolute_error(y_test, y_pred_ridge)
```

```
[20]: 2.298757602037677
```

```
[21]: root_mean_squared_error(y_test, y_pred_ridge)
```

```
[21]: 8.045449532074198
```

```
[22]: r2_score(y_test, y_pred_ridge)
```

```
[22]: 0.8378446631702253
```

1.4 Gradient Boosting Regressor

```
[23]: gbr = GradientBoostingRegressor()  
gbr.fit(X_train, y_train)
```

```
[23]: GradientBoostingRegressor()
```

```
[24]: y_pred_gbr = gbr.predict(X_test)  
y_pred_gbr
```

```
[24]: array([24.81614709, 26.92370305, 24.05991901, 17.30771176, 16.19740785,  
        25.28778472, 18.47327473, 13.4293537 , 17.87010961, 25.32283075,  
        17.3784293 , 32.80854588, 26.43231189, 16.71346622, 19.89542764,  
        26.93072095, 29.09128776, 34.5437197 , 16.2287741 , 12.9976502 ,  
        16.15260843, 14.0218489 , 23.430844 , 23.554742 , 12.15018212,  
        29.3384603 , 13.78280808, 37.46136379, 13.83317576, 21.92720442,  
        35.35308293, 27.59229457, 15.38298242, 25.77062188, 32.41174551,  
        36.29002888, 22.94838594, 13.1043108 , 35.75852911, 14.27426203,  
        31.12414823, 24.07273405, 15.53471207, 33.44820815, 26.03408155,  
        25.8403547 , 20.59553436, 16.25817053, 15.70183123, 13.21764247,  
        16.88855731, 18.84642478, 31.8104251 , 16.10117025, 36.1970594 ,  
        13.4293537 , 33.0030821 , 18.74768476, 18.59036535, 35.9761289 ,  
        19.47705059, 20.92402884, 14.42364406, 36.40972479, 32.92852343,  
        16.49541674, 24.5359267 , 14.67871467, 31.81852565, 23.62370365,  
        29.16526417, 20.5589327 , 14.29479101, 15.56229952, 23.83009838,  
        21.90942175, 13.86561836, 34.69275483, 29.33933 , 19.03518476])
```

```
[25]: mean_absolute_error(y_test, y_pred_gbr)
```

```
[25]: 1.7513204620962235
```

```
[26]: root_mean_squared_error(y_test, y_pred_gbr)
```

```
[26]: 5.225603103128156
```

```
[27]: r2_score(y_test, y_pred_gbr)
```

```
[27]: 0.8946784231324356
```

1.5 Stacking Regression

```
[28]: estimators = [  
        ('rfr', rfr),  
        ('ridge', ridge_cv.best_estimator_),  
        ('lr', lr)  
    ]
```

```

sr = StackingRegressor(
    estimators=estimators,
    final_estimator=gbr
)

sr.fit(X_train, y_train)

```

```

[28]: StackingRegressor(estimators=[('rfr', RandomForestRegressor(random_state=13)),
                                   ('ridge', Ridge(alpha=10)),
                                   ('lr', LinearRegression())],
                        final_estimator=GradientBoostingRegressor())

```

```

[29]: y_pred_sr = sr.predict(X_test)
      y_pred_sr

```

```

[29]: array([23.53378874, 25.71455216, 24.09668026, 15.40149474, 15.40615392,
            27.36438654, 17.98058559, 13.43407508, 18.52080362, 24.24343196,
            16.93869117, 30.90969533, 26.90396953, 16.73109235, 18.98478987,
            26.63577215, 25.55108187, 29.9583349 , 15.33757724, 12.2164736 ,
            15.40615392, 15.17557021, 24.92925791, 22.1997565 , 12.13294436,
            33.13916835, 13.79983485, 34.80103745, 13.56035943, 20.45624966,
            36.40332172, 27.01318555, 15.40615392, 23.05829459, 31.64665183,
            36.72001283, 27.15296778, 13.79983485, 32.15462099, 14.16801969,
            34.11696062, 23.53378874, 15.87656151, 32.40313431, 25.59016802,
            25.45964114, 20.68403814, 15.81340739, 15.40615392, 14.3074484 ,
            15.9730475 , 18.43689213, 31.85386936, 16.50287556, 39.40124297,
            13.56035943, 33.96965856, 17.56922769, 18.4320498 , 36.70733073,
            16.3237033 , 20.25150774, 13.56035943, 37.78515789, 34.28042863,
            15.87161528, 25.40213556, 15.87161528, 32.20626472, 23.6224426 ,
            30.78302652, 20.5972072 , 15.00843971, 15.46436181, 25.59016802,
            20.68403814, 14.3074484 , 37.76093249, 30.96411227, 18.98478987])

```

```

[30]: mean_absolute_error(y_test, y_pred_sr)

```

```

[30]: 1.9308715384123967

```

```

[31]: root_mean_squared_error(y_test, y_pred_sr)

```

```

[31]: 6.580558650525768

```

```

[32]: r2_score(y_test, y_pred_sr)

```

```

[32]: 0.8673694117090569

```

1.6 Voting Regressor

```
[33]: vr = VotingRegressor([
        ('rfr', rfr),
        ('gbr', gbr),
        ('lr', lr)
    ], weights=[2,3,1])

vr.fit(X_train, y_train)
```

```
[33]: VotingRegressor(estimators=[('rfr', RandomForestRegressor(random_state=13)),
                                   ('gbr', GradientBoostingRegressor()),
                                   ('lr', LinearRegression())],
                      weights=[2, 3, 1])
```

```
[34]: y_pred_vc = vr.predict(X_test)
y_pred_vc
```

```
[34]: array([24.76159933, 26.45300538, 24.43931921, 16.84999193, 16.02682976,
          25.69260264, 18.17856085, 12.94358365, 18.46961914, 25.23284026,
          17.63732911, 31.98121141, 26.39382748, 16.79346121, 19.88349928,
          26.98685806, 28.9228603 , 33.30536697, 16.10094887, 12.17261444,
          16.11287324, 14.14309616, 24.26759346, 23.24500508, 11.51039815,
          30.15249683, 13.76105267, 37.08910687, 13.29329954, 22.18686061,
          35.0311041 , 27.98330729, 15.36577696, 25.49113713, 32.49028417,
          36.20594877, 24.02677952, 13.23408079, 34.30879498, 14.44509707,
          31.157572 , 24.32135164, 15.3348295 , 31.67918025, 26.07416706,
          25.79671627, 21.43089366, 16.21514328, 15.32467125, 13.43717625,
          16.7480797 , 19.04268902, 32.84100406, 16.45293143, 35.24655866,
          13.24453049, 32.82314004, 19.21943415, 18.59105612, 35.15685845,
          18.27006274, 20.89460649, 13.48210674, 36.08456992, 33.0226191 ,
          16.16638182, 24.64586799, 15.22153291, 32.13039654, 23.51922136,
          29.75720817, 21.02781351, 14.27387728, 15.22299296, 24.51195192,
          21.9301146 , 13.6426772 , 35.63240991, 30.02487235, 19.42224996])
```

```
[35]: root_mean_squared_error(y_test, y_pred_vc)
```

```
[35]: 5.0571779581339475
```

```
[36]: r2_score(y_test, y_pred_vc)
```

```
[36]: 0.8980730173074735
```


1.7 Final Estimator

```
[37]: estimators2 = [  
        ('rfr', rfr),  
        ('ridge', ridge_cv.best_estimator_),  
        ('gbr', gbr)  
    ]  
  
    sr2 = StackingRegressor(  
        estimators=estimators2,  
        final_estimator=vr  
    )  
    sr2.fit(X_train, y_train)
```

```
[37]: StackingRegressor(estimators=[('rfr', RandomForestRegressor(random_state=13)),  
                                     ('ridge', Ridge(alpha=10)),  
                                     ('gbr', GradientBoostingRegressor())],  
                        final_estimator=VotingRegressor(estimators=[('rfr',  
RandomForestRegressor(random_state=13)),  
                                                                    ('gbr',  
GradientBoostingRegressor()),  
                                                                    ('lr',  
LinearRegression())],  
                                                         weights=[2, 3, 1]))
```

```
[38]: y_pred_sr2 = sr2.predict(X_test)  
y_pred_sr2
```

```
[38]: array([24.3491363 , 27.11568059, 23.39061009, 17.05006586, 15.40976879,  
            25.52060588, 18.04750935, 13.39328651, 18.64617906, 24.98338014,  
            17.1792968 , 30.97858195, 28.10750468, 15.68440833, 19.64604119,  
            27.63643889, 26.05606814, 30.67613107, 15.67300316, 12.33635145,  
            15.40526841, 14.087848 , 25.02953828, 22.15264521, 10.96638522,  
            33.17943801, 13.99024305, 35.98606902, 13.3395992 , 21.31222876,  
            36.64232503, 28.26352216, 15.0655988 , 24.74528493, 32.06494113,  
            37.00372476, 25.66198544, 13.83742911, 32.91868345, 13.83196459,  
            33.0385216 , 23.3559947 , 15.62901099, 30.6175818 , 27.05611273,  
            25.57012826, 21.10867599, 15.50163291, 15.26869649, 13.67442293,  
            15.7415394 , 18.49101327, 34.04061254, 16.02565044, 38.78773653,  
            13.52816628, 33.1378723 , 17.90159211, 18.37792041, 35.9752136 ,  
            18.70332375, 20.73245054, 14.26120238, 37.42491304, 33.21332751,  
            15.71597323, 25.43355876, 15.75085825, 32.73059787, 23.50424327,  
            30.53905967, 20.94953693, 14.23724779, 15.62513511, 24.19958196,  
            20.9596028 , 13.96885838, 37.38444735, 30.55691623, 19.2151501 ])
```

```
[39]: root_mean_squared_error(y_test, y_pred_sr2)
```

[39]: 6.414412415181024

```
[40]: r2_score(y_test, y_pred_sr2)
```

[40]: 0.8707180746579599

```
[41]: estimators3 = [  
    ('rfr', rfr),  
    ('ridge', ridge_cv.best_estimator_),  
    ('svr', SVR(C=1.0, kernel='linear')),  
    ('random_forest', RandomForestRegressor())  
]  
  
sr3 = StackingRegressor(  
    estimators=estimators3,  
    final_estimator=Ridge(alpha=1.0)  
)  
  
param_grid_sr = {  
    'random_forest__n_estimators': [50, 100, 250],  
    'svr__C': [0.1, 1.0, 10.0],  
    'final_estimator__alpha': [0.1, 1.0, 10.0]  
}  
  
sr_cv = RandomizedSearchCV(sr3, param_grid_sr, n_iter=5, cv=3,  
    scoring='neg_root_mean_squared_error', n_jobs=-1)  
sr_cv.fit(X_train, y_train)
```

```
[41]: RandomizedSearchCV(cv=3,  
    estimator=StackingRegressor(estimators=[('rfr',  
RandomForestRegressor(random_state=13)),  
                                            ('ridge',  
Ridge(alpha=10)),  
                                            ('svr',  
SVR(kernel='linear')),  
                                            ('random_forest',  
RandomForestRegressor())],  
    final_estimator=Ridge()),  
    n_iter=5, n_jobs=-1,  
    param_distributions={'final_estimator__alpha': [0.1, 1.0,  
10.0],  
                        'random_forest__n_estimators': [50, 100,  
250],  
                        'svr__C': [0.1, 1.0, 10.0]},  
    scoring='neg_mean_squared_error')
```

```
[42]: y_pred_8 = sr_cv.predict(X_test)
      y_pred_8
```

```
[42]: array([24.45505141, 26.08064778, 25.377915 , 15.93300187, 15.24263626,
          25.45247833, 17.30640013, 12.365641 , 18.28087359, 25.39590536,
          17.09670375, 30.33733978, 26.29055207, 16.56544171, 19.26286061,
          27.02489251, 29.40600026, 32.61210673, 15.07380717, 11.0614695 ,
          15.61011049, 13.4696746 , 24.67148938, 23.53833227, 11.43036111,
          31.19528332, 13.48103115, 37.06619408, 12.66684219, 21.79112221,
          34.90552177, 28.92212115, 14.95776505, 25.75974317, 32.5812645 ,
          36.34849374, 24.79244634, 12.77105095, 32.95633545, 14.24149635,
          31.54489825, 24.29351523, 14.88776827, 30.01059792, 26.08003928,
          25.61700765, 21.68020634, 15.65541277, 14.90706114, 13.46447027,
          15.61558932, 18.61179576, 34.77888606, 16.10115523, 34.68386065,
          13.07459897, 31.41977482, 19.17668167, 17.95115591, 34.99862448,
          16.6652014 , 20.35133345, 12.5105491 , 35.28992113, 33.31986247,
          15.64511174, 24.53707771, 15.51719415, 31.25388316, 23.41547812,
          30.63741646, 21.42362772, 14.42679083, 14.19939658, 25.45697562,
          22.11579612, 13.25555966, 37.29781596, 30.58671278, 19.00212692])
```

```
[43]: root_mean_squared_error(y_test, y_pred_8)
```

```
[43]: 5.580300867802961
```

```
[44]: r2_score(y_test, y_pred_8)
```

```
[44]: 0.8875295204795366
```

```
[45]: sr_cv.best_estimator_
```

```
[45]: StackingRegressor(estimators=[('rfr', RandomForestRegressor(random_state=13)),
                                   ('ridge', Ridge(alpha=10)),
                                   ('svr', SVR(C=10.0, kernel='linear')),
                                   ('random_forest', RandomForestRegressor())],
                       final_estimator=Ridge(alpha=0.1))
```

27-Random_Forest_Regressor

October 20, 2024

1 Random Forest Regressor

A **Random Forest Regressor** is an ensemble learning method that combines the predictions of multiple decision trees to improve predictive performance for regression tasks. Unlike a single decision tree, which might overfit or have high variance, a random forest aggregates the outputs of many trees to make the final prediction, thus reducing overfitting and variance while improving generalization.

It works by constructing several decision trees during training and averaging their predictions. This process of averaging multiple trees helps to smooth out the predictions and gives more robust results.

1.0.1 Why is Random Forest Regressor Important?

- **Reduces Overfitting:** Individual decision trees tend to overfit, especially when the trees are deep. By averaging the predictions of many trees, Random Forest mitigates overfitting, making the model more generalizable to unseen data.
- **Handles Non-linear Relationships:** Decision trees naturally handle non-linear relationships. A random forest regressor benefits from this, enabling it to capture complex patterns in the data.
- **Robustness to Noise and Outliers:** Because Random Forest averages the predictions of many trees, it is less sensitive to outliers and noise than individual decision trees.
- **Feature Importance:** Random Forest provides feature importance scores, which can be useful in understanding which features have the most predictive power.

1.0.2 How does Random Forest Regressor work?

Random Forest Regressor is based on an ensemble of decision trees. Here's a step-by-step explanation of how it works:

1. Bagging (Bootstrap Aggregating): The first key idea behind Random Forest is bagging. Random subsets (bootstrapped samples) of the original training data are created. Each decision tree in the forest is trained on a different random sample of the data. This technique helps in reducing the variance and overfitting that can occur with individual trees.

2. Random Feature Selection:** During the construction of each tree, Random Forest also selects a random subset of features to consider at each split. This random selection of features ensures that the trees in the forest are more diverse, further reducing correlation between the trees.

3. Tree Construction: Each tree is built independently by splitting the data at different nodes based on feature values (as in any decision tree). The process is recursive, and the tree grows until a stopping criterion is met (e.g., a minimum number of samples per leaf or a maximum tree depth).

4. Prediction: After all the trees are trained, predictions are made by each tree, and the final prediction for regression is the average of the predictions from all the trees.

1.0.3 When should you use Random Forest Regressor?

- **Non-linear and complex datasets:** Random Forest works well when your dataset has complex, non-linear relationships between features and target values.
- **High-dimensional data:** It handles high-dimensional datasets with many features and can reduce the risk of overfitting by selecting random subsets of features.
- **When interpretability is not the primary goal:** Random Forest is less interpretable than simpler models like linear regression, so it's best used when predictive accuracy is more important than understanding the model's inner workings.
- **When there are many categorical and continuous variables:** It can handle mixed-type data without much preprocessing.

1.0.4 Who uses Random Forest Regressor?

- **Data Scientists:** Random Forest is widely used in industries like finance, healthcare, and marketing to solve regression problems where complex patterns need to be captured.
- **Business Analysts:** It is often used to predict continuous outcomes such as sales forecasts, customer retention rates, and other business metrics.
- **Researchers:** In research fields like genomics or environmental science, Random Forest is applied to make predictions based on numerous features, and it's useful for feature selection.

1.0.5 Key Points to Remember:

- **Ensemble of Decision Trees:** Random Forest is a collection of decision trees, each trained on a different bootstrap sample of the data.
- **Random Feature Selection:** At each split, a random subset of features is chosen, making the trees less correlated with one another.
- **Averaging Predictions:** For regression tasks, the final prediction is the average of the predictions made by all the trees.
- **Feature Importance:** Random Forest naturally provides feature importance scores, which can help identify the most relevant features.

1.0.6 Advantages of Random Forest Regressor:

- **Reduced Overfitting:** By averaging multiple trees, Random Forest smooths out the predictions and reduces overfitting compared to a single decision tree.
- **Handles Missing Data:** Random Forest can handle missing values relatively well by filling in missing data based on correlations.

- **Robust to Outliers:** Due to the averaging process, Random Forest is less sensitive to outliers in the data.
- **No need for feature scaling:** Unlike models like SVM or linear regression, Random Forest doesn't require feature scaling (e.g., normalization or standardization).

```
[1]: import pandas as pd
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

healthexp = sns.load_dataset('healthexp')
healthexp
```

```
[1]:
```

	Year	Country	Spending_USD	Life_Expectancy
0	1970	Germany	252.311	70.6
1	1970	France	192.143	72.2
2	1970	Great Britain	123.993	71.9
3	1970	Japan	150.437	72.0
4	1970	USA	326.961	70.9
..
269	2020	Germany	6938.983	81.1
270	2020	France	5468.418	82.3
271	2020	Great Britain	5018.700	80.4
272	2020	Japan	4665.641	84.7
273	2020	USA	11859.179	77.0

[274 rows x 4 columns]

```
[2]: healthexp = pd.get_dummies(healthexp)
healthexp
```

```
[2]:
```

	Year	Spending_USD	Life_Expectancy	Country_Canada	Country_France	\
0	1970	252.311	70.6	False	False	
1	1970	192.143	72.2	False	True	
2	1970	123.993	71.9	False	False	
3	1970	150.437	72.0	False	False	
4	1970	326.961	70.9	False	False	
..	
269	2020	6938.983	81.1	False	False	
270	2020	5468.418	82.3	False	True	
271	2020	5018.700	80.4	False	False	
272	2020	4665.641	84.7	False	False	
273	2020	11859.179	77.0	False	False	

	Country_Germany	Country_Great Britain	Country_Japan	Country_USA
0	True	False	False	False
1	False	False	False	False
2	False	True	False	False
3	False	False	True	False
4	False	False	False	True
..
269	True	False	False	False
270	False	False	False	False
271	False	True	False	False
272	False	False	True	False
273	False	False	False	True

[274 rows x 9 columns]

```
[3]: X = healthexp.drop(['Life_Expectancy'], axis=1)
y = healthexp['Life_Expectancy']

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
↳random_state=19)

rfr = RandomForestRegressor(random_state=13)
rfr.fit(X_train, y_train)
```

```
[3]: RandomForestRegressor(random_state=13)
```

```
[4]: y_pred = rfr.predict(X_test)
y_pred
```

```
[4]: array([79.705, 81.089, 74.833, 79.357, 71.152, 79.186, 82.417, 81.7 ,
78.368, 71.48 , 76.446, 75.79 , 82.694, 78.601, 76.848, 77.614,
71.838, 78.91 , 72.414, 82.033, 71.816, 78.415, 79.937, 78.513,
78.265, 72.639, 75.759, 81.337, 81.408, 82.269, 74.003, 74.395,
78.677, 74.98 , 71.777, 74.347, 74.581, 76.581, 81.142, 80.895,
74.979, 81.68 , 77.489, 78.219, 80.888, 83.572, 78.392, 73.976,
77.236, 80.114, 74.728, 76.782, 78.031, 78.778, 77.519])
```

```
[5]: mean_absolute_error(y_test, y_pred)
```

```
[5]: 0.25916363636361917
```

```
[6]: root_mean_squared_error(y_test, y_pred)
```

```
[6]: 0.31970520512155615
```

```
[7]: r2_score(y_test, y_pred)
```

```
[7]: 0.9910457602615238
```

```
[8]: param_grid = {
      'n_estimators': [100, 200, 300],
      'max_depth': [10, 20, 30],
      'min_samples_split': [2, 5, 10],
      'min_samples_leaf': [1, 2, 4]
    }

    rfr_cv = GridSearchCV(rfr, param_grid, cv=3, n_jobs=-1)
    rfr_cv.fit(X_train, y_train)
```

```
[8]: GridSearchCV(cv=3, estimator=RandomForestRegressor(random_state=13), n_jobs=-1,
                  param_grid={'max_depth': [10, 20, 30],
                              'min_samples_leaf': [1, 2, 4],
                              'min_samples_split': [2, 5, 10],
                              'n_estimators': [100, 200, 300]})
```

```
[9]: y_pred_cv = rfr_cv.predict(X_test)
      y_pred_cv
```

```
[9]: array([79.71933333, 81.09833333, 74.83733333, 79.40433333, 71.159      ,
           79.232      , 82.41166667, 81.686      , 78.43533333, 71.404      ,
           76.47266667, 75.904      , 82.69633333, 78.64833333, 76.843      ,
           77.58633333, 71.715      , 78.93433333, 72.42533333, 82.016      ,
           71.731      , 78.45666667, 80.         , 78.529      , 78.31966667,
           72.51466667, 75.78433333, 81.29033333, 81.392      , 82.281      ,
           73.972      , 74.404      , 78.67233333, 74.99466667, 71.68766667,
           74.34333333, 74.56433333, 76.604      , 81.14366667, 80.91      ,
           74.97633333, 81.66533333, 77.59366667, 78.17166667, 80.87033333,
           83.52633333, 78.43366667, 73.93433333, 77.149      , 80.104      ,
           74.718      , 76.741      , 78.136      , 78.70566667, 77.53133333])
```

```
[10]: mean_absolute_error(y_test, y_pred_cv)
```

```
[10]: 0.25089696969702713
```

```
[11]: root_mean_squared_error(y_test, y_pred_cv)
```

```
[11]: 0.31042865768991584
```

```
[12]: r2_score(y_test, y_pred_cv)
```

```
[12]: 0.9915578528520343
```


28-Gradient_Boosting_Regressor

October 20, 2024

1 Gradient Boosting Regressor

A **Gradient Boosting Regressor** is a machine learning algorithm that builds an ensemble of decision trees sequentially to minimize the loss function (or error) for regression tasks. Unlike Random Forest, which builds trees independently, Gradient Boosting builds trees sequentially, where each new tree aims to correct the errors made by the previous trees. It uses the gradient descent technique to optimize the model by reducing the difference between the predicted and actual values.

It combines the predictions of many “weak learners” (typically shallow decision trees) to create a more accurate final prediction. The key idea is to reduce bias and variance by focusing on improving areas where the model previously performed poorly.

1.0.1 Why is Gradient Boosting Regressor Important?

- **Accurate Predictions:** Gradient Boosting often produces very accurate models that outperform simpler models like linear regression, and even Random Forest, in certain cases.
- **Focuses on Hard-to-Predict Instances:** By iteratively correcting errors, Gradient Boosting effectively captures complex patterns in the data that may be missed by other models.
- **Balances Bias-Variance Tradeoff:** Gradient Boosting reduces both bias (error due to underfitting) and variance (error due to overfitting), providing a powerful model in various situations.

1.0.2 How does Gradient Boosting Regressor work?

The core idea behind Gradient Boosting is to add new models that improve the residual errors (i.e., the difference between predicted and true values) of previous models. Here’s the step-by-step process:

- 1. Initial Model:** Start with an initial model, typically predicting the mean of the target variable.
- 2. Calculate Residuals*:** For each data point, compute the residual (the difference between the true value and the model’s prediction).
- 3. Fit a Weak Learner:** A decision tree is fitted to these residuals. This new tree learns to predict the errors (or residuals) made by the previous model.
- 4. Update the Model:** The new predictions are added to the previous predictions. This step is scaled by a learning rate to control the contribution of the new tree.

5. Repeat: This process is repeated iteratively, each time fitting a new tree to the residuals and updating the model.

6. Final Prediction: After a certain number of iterations (trees), the final model is the sum of the predictions from all trees.

The key difference between Gradient Boosting and Random Forest is that Gradient Boosting builds trees sequentially, with each new tree correcting the errors made by the previous trees, whereas Random Forest builds trees independently in parallel.

1.0.3 When should you use Gradient Boosting Regressor?

- **When high accuracy is needed:** Gradient Boosting is particularly powerful when you need highly accurate predictions and are willing to trade off some interpretability.
- **Imbalanced or Noisy Data:** Gradient Boosting handles imbalanced and noisy data well, making it suitable for real-world applications where clean, well-balanced datasets are rare.
- **When dealing with complex relationships:** If the data exhibits complex relationships between features and target variables, Gradient Boosting can model these effectively.
- **When feature importance is important:** Like Random Forest, Gradient Boosting also provides feature importance, which can help in understanding which features contribute the most to the predictions.

1.0.4 Who uses Gradient Boosting Regressor?

- **Data Scientists and Machine Learning Engineers:** Gradient Boosting is widely used in industries like finance, healthcare, and marketing for predictive modeling tasks where high accuracy is crucial.
- **Kaggle Competitors:** Gradient Boosting models, especially variants like XGBoost, LightGBM, and CatBoost, are commonly used in data science competitions because of their ability to achieve state-of-the-art performance.
- **Researchers:** In research fields like genetics or climate modeling, Gradient Boosting is used for regression tasks that require capturing non-linear patterns and interactions among features.

1.0.5 Key Points to Remember:

- **Sequential Trees:** Gradient Boosting builds trees one at a time, with each tree focusing on correcting the residual errors of the previous trees.
- **Learning Rate:** The learning rate controls how much the new tree's predictions influence the overall model. A lower learning rate requires more trees to be effective but helps prevent overfitting.
- **Combining Weak Learners:** It combines many "weak learners" (typically small decision trees) to create a strong predictive model.
- **Loss Function:** Gradient Boosting minimizes a loss function (e.g., Mean Squared Error for regression) by performing gradient descent in function space.

1.0.6 Advantages of Gradient Boosting Regressor:

- 1. High Predictive Accuracy:** Gradient Boosting tends to produce highly accurate models that often outperform other models on many datasets.
- 2. Flexibility:** You can specify different loss functions (e.g., least squares for regression) and customize the model's hyperparameters.
- 3. Handles Missing Data:** Gradient Boosting can handle missing data quite well and does not require complex imputation strategies.
- 4. Feature Importance:** It provides feature importance, which can be useful for understanding which features have the most predictive power.
- 5. Robust to Outliers:** It focuses on reducing residual errors, making it more robust to outliers in some cases.

1.0.7 Limitations of Gradient Boosting Regressor:

- **Longer Training Time:** Since Gradient Boosting builds trees sequentially, training can take significantly longer than other algorithms like Random Forest, especially on large datasets.
- **Sensitive to Hyperparameters:** Gradient Boosting can be sensitive to the choice of hyperparameters, such as the number of trees, learning rate, and tree depth. It often requires careful tuning to perform optimally.
- **Can Overfit:** Although Gradient Boosting generally performs well, it can overfit the training data if the number of trees is too high or if the trees are too deep.
- **Not Interpretable:** Like other ensemble methods, Gradient Boosting produces complex models that are difficult to interpret compared to simpler models like linear regression.

```
[1]: import pandas as pd

from sklearn import datasets
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

diabetes = datasets.load_diabetes()
diabetes.data
```

```
[1]: array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
            0.01990749, -0.01764613],
        [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
        -0.06833155, -0.09220405],
        [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
            0.00286131, -0.02593034],
        ...,
        [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
        -0.04688253,  0.01549073],
```

```
[-0.04547248, -0.04464164, 0.03906215, ..., 0.02655962,  
 0.04452873, -0.02593034],  
[-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,  
 -0.00422151, 0.00306441]])
```

```
[2]: diabetes.target
```

```
[2]: array([151., 75., 141., 206., 135., 97., 138., 63., 110., 310., 101.,  
 69., 179., 185., 118., 171., 166., 144., 97., 168., 68., 49.,  
 68., 245., 184., 202., 137., 85., 131., 283., 129., 59., 341.,  
 87., 65., 102., 265., 276., 252., 90., 100., 55., 61., 92.,  
 259., 53., 190., 142., 75., 142., 155., 225., 59., 104., 182.,  
 128., 52., 37., 170., 170., 61., 144., 52., 128., 71., 163.,  
 150., 97., 160., 178., 48., 270., 202., 111., 85., 42., 170.,  
 200., 252., 113., 143., 51., 52., 210., 65., 141., 55., 134.,  
 42., 111., 98., 164., 48., 96., 90., 162., 150., 279., 92.,  
 83., 128., 102., 302., 198., 95., 53., 134., 144., 232., 81.,  
 104., 59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,  
 173., 180., 84., 121., 161., 99., 109., 115., 268., 274., 158.,  
 107., 83., 103., 272., 85., 280., 336., 281., 118., 317., 235.,  
 60., 174., 259., 178., 128., 96., 126., 288., 88., 292., 71.,  
 197., 186., 25., 84., 96., 195., 53., 217., 172., 131., 214.,  
 59., 70., 220., 268., 152., 47., 74., 295., 101., 151., 127.,  
 237., 225., 81., 151., 107., 64., 138., 185., 265., 101., 137.,  
 143., 141., 79., 292., 178., 91., 116., 86., 122., 72., 129.,  
 142., 90., 158., 39., 196., 222., 277., 99., 196., 202., 155.,  
 77., 191., 70., 73., 49., 65., 263., 248., 296., 214., 185.,  
 78., 93., 252., 150., 77., 208., 77., 108., 160., 53., 220.,  
 154., 259., 90., 246., 124., 67., 72., 257., 262., 275., 177.,  
 71., 47., 187., 125., 78., 51., 258., 215., 303., 243., 91.,  
 150., 310., 153., 346., 63., 89., 50., 39., 103., 308., 116.,  
 145., 74., 45., 115., 264., 87., 202., 127., 182., 241., 66.,  
 94., 283., 64., 102., 200., 265., 94., 230., 181., 156., 233.,  
 60., 219., 80., 68., 332., 248., 84., 200., 55., 85., 89.,  
 31., 129., 83., 275., 65., 198., 236., 253., 124., 44., 172.,  
 114., 142., 109., 180., 144., 163., 147., 97., 220., 190., 109.,  
 191., 122., 230., 242., 248., 249., 192., 131., 237., 78., 135.,  
 244., 199., 270., 164., 72., 96., 306., 91., 214., 95., 216.,  
 263., 178., 113., 200., 139., 139., 88., 148., 88., 243., 71.,  
 77., 109., 272., 60., 54., 221., 90., 311., 281., 182., 321.,  
 58., 262., 206., 233., 242., 123., 167., 63., 197., 71., 168.,  
 140., 217., 121., 235., 245., 40., 52., 104., 132., 88., 69.,  
 219., 72., 201., 110., 51., 277., 63., 118., 69., 273., 258.,  
 43., 198., 242., 232., 175., 93., 168., 275., 293., 281., 72.,  
 140., 189., 181., 209., 136., 261., 113., 131., 174., 257., 55.,  
 84., 42., 146., 212., 233., 91., 111., 152., 120., 67., 310.,  
 94., 183., 66., 173., 72., 49., 64., 48., 178., 104., 132.,
```

```
220., 57.]])
```

```
[3]: X = diabetes.data
y = diabetes.target

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
↳random_state=17)

gbr = GradientBoostingRegressor()
gbr.fit(X_train, y_train)
```

```
[3]: GradientBoostingRegressor()
```

```
[4]: y_pred = gbr.predict(X_test)
y_pred
```

```
[4]: array([192.51225586, 200.5750764 , 82.78231894, 101.58510513,
125.14414245, 81.39853646, 93.84628461, 102.45928951,
121.90945671, 74.25891122, 255.91425718, 175.42143255,
93.81164289, 99.4616262 , 217.43499676, 217.43197684,
190.62007333, 172.98697169, 197.2837651 , 175.80982432,
182.09395792, 269.70763529, 196.07382183, 82.25148771,
114.47110984, 185.87258222, 109.04630496, 165.3091288 ,
84.08079017, 241.11508989, 89.77664274, 69.54907641,
108.03496427, 288.56040447, 170.41296018, 164.48100696,
82.93345568, 179.53609042, 91.32177689, 245.58054393,
112.63272196, 90.77107175, 179.75704585, 306.48853493,
154.11016096, 91.65230129, 200.49974318, 107.20195482,
190.62028854, 167.72815926, 191.98143713, 186.27531144,
200.12776418, 156.00397936, 180.02678304, 138.30178657,
96.52611334, 300.64236743, 138.20868823, 162.7712336 ,
176.10010144, 231.10141163, 91.57716704, 143.7471825 ,
87.83391334, 91.52052802, 216.22481281, 68.19976982,
167.75688514, 157.00162733, 103.12604888, 283.27021778,
267.45306696, 196.29244184, 96.48618333, 124.81762663,
199.10839318, 149.25733184, 145.53020157, 178.10088311,
258.17118587, 92.95097597, 149.38905097, 71.33146438,
272.96959869, 311.01641671, 95.577347 , 104.14250789,
83.75796396])
```

```
[5]: mean_absolute_error(y_test, y_pred)
```

```
[5]: 48.792591465108316
```

```
[6]: root_mean_squared_error(y_test, y_pred)
```

```
[6]: 60.90137361126108
```

```
[7]: r2_score(y_test, y_pred)
```

```
[7]: 0.37191596197265364
```

```
[8]: param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3,4,5],
    'min_samples_split': [2,3,4],
    'min_samples_leaf': [1,2,3]
}

gbr_cv = GridSearchCV(gbr, param_grid, cv=5,
    ↪n_jobs=-1,scoring='neg_root_mean_squared_error')
gbr_cv.fit(X_test, y_test)
```

```
[8]: GridSearchCV(cv=5, estimator=GradientBoostingRegressor(), n_jobs=-1,
    param_grid={'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 4, 5], 'min_samples_leaf': [1, 2, 3],
    'min_samples_split': [2, 3, 4],
    'n_estimators': [100, 200, 300]}},
    scoring='neg_root_mean_squared_error')
```

```
[9]: y_pred_gbr_cv = gbr_cv.predict(X_test)
y_pred_gbr_cv
```

```
[9]: array([167.54547384, 196.45689046, 131.99856228, 128.81980861,
    155.38736053, 104.78343157, 104.46878826, 103.86864233,
    172.16237464,  92.877811  , 226.28290431, 151.96615636,
    99.36099017, 184.74380459, 209.7874883 , 169.38985178,
    240.03918274, 210.69910838, 247.98095965, 113.01953415,
    110.64514541, 232.66480304, 109.31434018,  91.94434266,
    96.51531206, 215.20968459, 105.09234396, 138.79638199,
    82.32851947, 215.84795964, 119.4373777 , 116.85111785,
    109.52830044, 199.97027303, 126.65963299, 231.83342139,
    79.28364229, 253.17948242, 107.99860644, 248.12009523,
    95.73091409, 134.91997317, 165.55760765, 241.81722869,
    107.02535237,  78.75843641, 211.98096742,  99.36099017,
    137.05366889, 163.12544504, 164.34451529, 178.11208294,
    182.44224851, 114.1775484 , 201.79323332, 104.77956892,
    115.67104533, 233.66677325, 163.98975366, 221.99947933,
    198.83299795, 172.19115755,  90.89484407, 123.40323468,
    131.31453483, 191.68730783, 167.43521912, 105.88265868,
    187.91161719, 150.97652913, 119.96678341, 241.33644356,
    257.04655152, 164.34451529,  99.89637165, 143.87516951,
    101.74985603, 243.13993649, 152.49661722, 231.19421556,
    233.42395538, 120.44390831, 140.12835633,  81.63035609,
```

```
236.03992088, 239.51517181, 117.57491008, 104.18911766,  
104.31100263])
```

```
[10]: mean_absolute_error(y_test, y_pred_gbr_cv)
```

```
[10]: 27.560044315482546
```

```
[11]: root_mean_squared_error(y_test, y_pred_gbr_cv)
```

```
[11]: 32.99791585313554
```

```
[12]: r2_score(y_test, y_pred_gbr_cv)
```

```
[12]: 0.8156103237854524
```

```
[13]: gbr_cv.best_params_
```

```
[13]: {'learning_rate': 0.01,  
      'max_depth': 3,  
      'min_samples_leaf': 3,  
      'min_samples_split': 3,  
      'n_estimators': 200}
```

29-Optuna_Hypertuning

October 20, 2024

1 Hypertuning with Optuna

Optuna is an open-source, hyperparameter optimization framework designed for machine learning and deep learning models. It automates the search for optimal hyperparameters, helping data scientists and machine learning engineers tune models more efficiently. Optuna is flexible, scalable, and supports both sequential and parallel optimization, allowing it to adapt to different scales of problems.

Optuna's key strength is its define-by-run approach, where hyperparameter configurations are dynamically constructed during each trial, making it more efficient and flexible than traditional grid or random search methods.

1.0.1 Why is Optuna Important?

Tuning hyperparameters is a critical part of improving machine learning models' performance, but it can be extremely time-consuming, especially for complex models like deep neural networks or ensembles. Optuna helps by:

- **Automating hyperparameter search:** Reducing the manual effort required to find optimal settings.
- **Efficient search:** Optuna uses sophisticated algorithms like Tree-structured Parzen Estimator (TPE) and CMA-ES (Covariance Matrix Adaptation Evolution Strategy) to explore the hyperparameter space more intelligently than grid or random search.
- **Early stopping:** It includes pruning to stop trials that are not promising early, saving time and computational resources.
- **Scalability:** Optuna can be easily scaled to large clusters or cloud environments, making it practical for tuning complex models.

1.0.2 How does Optuna Hypertuning Work?

Optuna performs hyperparameter optimization through an iterative process called trials, where each trial represents a single set of hyperparameters and their corresponding evaluation. Optuna tries to minimize (or maximize) the objective function defined by the user. Here's how it works:

1. Define the Objective Function: The objective function evaluates the model's performance using a given set of hyperparameters. It returns a metric (like accuracy, loss, etc.) to minimize (or maximize).

2. Sampling Hyperparameters: Optuna uses advanced optimization algorithms like TPE or CMA-ES to select hyperparameter values for each trial. It doesn't just randomly choose values but uses past trial results to guide the search toward more promising regions of the hyperparameter space.

3. Run Trials: Optuna evaluates the model with the selected hyperparameters in each trial and records the objective value (e.g., validation accuracy or loss).

4. Pruning: If a trial is not performing well, Optuna can prune (i.e., stop) it early to save resources and move on to more promising hyperparameter combinations.

5. Repeat: Optuna repeats the trial process multiple times (as defined by the user) until it finds the best set of hyperparameters or the predefined number of trials is completed.

1.0.3 How does Optuna Compare to Other Hyperparameter Optimization Methods?

Method	Pros	Cons
Grid Search	Exhaustively searches through predefined values.	Computationally expensive. Limited by pre-defined grid.
Random Search	Explores the search space randomly and is simpler than grid.	Inefficient for large search spaces. May miss optimal values.
Bayesian Search	More efficient by modeling the function being optimized.	Requires more complex setup. Slower for high-dimensional spaces.
Optuna	Define-by-run, efficient search with pruning and parallelism.	More complex to set up than simple random or grid search.

1.0.4 When should you use Optuna?

- **Hyperparameter Optimization for Complex Models:** If you are working with complex machine learning models like Gradient Boosting, Deep Neural Networks, or ensembles, where manual tuning is impractical.
- **Large Hyperparameter Search Space:** When the number of hyperparameters and possible values is large, Optuna's efficiency makes it superior to grid or random search.
- **Need for Resource Efficiency:** Optuna's ability to prune underperforming trials helps save computational time and resources, making it ideal for long-running tasks or expensive-to-evaluate models.

1.0.5 Who uses Optuna?

- **Machine Learning Engineers and Data Scientists:** Optuna is widely adopted by professionals working on machine learning projects where hyperparameter optimization is crucial for achieving high performance.

- **Deep Learning Researchers:** Researchers often use Optuna for neural network hyperparameter tuning, as it efficiently searches through learning rates, optimizers, and architectures.
- **Kaggle Competitors:** Competitors in machine learning competitions use Optuna to gain an edge by finding the best possible hyperparameters.

1.0.6 Key Features of Optuna:

1. **Define-by-Run:** Hyperparameter space is defined dynamically during the execution of the trials, allowing for flexibility in how you construct and explore the space.
2. **Pruning:** Automatically stops unpromising trials to save computation time and focus on better-performing hyperparameter combinations.
3. **Parallelism:** Optuna supports running multiple trials in parallel across different CPU or GPU resources, speeding up the search process.
4. **Visualization:** Provides built-in tools for visualizing optimization history, parameter importance, and more, making it easier to analyze the optimization process.

1.0.7 Advantages of Optuna:

- **Efficient Optimization:** Optuna uses state-of-the-art algorithms like TPE for efficient hyperparameter search, making it faster than random or grid search.
- **Dynamic Construction:** Unlike grid search, Optuna builds the hyperparameter space dynamically, making it more flexible.
- **Automatic Pruning:** Unpromising trials are stopped early, saving computation time. Parallelization: Optuna can parallelize trials to speed up the optimization process.
- **Visualization Tools:** Optuna includes tools to visualize the optimization process, which helps in understanding how hyperparameters impact model performance.

```
[34]: import pandas as pd
import seaborn as sns
import optuna
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_score

healthexp = sns.load_dataset('healthexp')
healthexp = pd.get_dummies(healthexp)
healthexp
```

```
[34]:
```

	Year	Spending_USD	Life_Expectancy	Country_Canada	Country_France	\
0	1970	252.311	70.6	False	False	
1	1970	192.143	72.2	False	True	
2	1970	123.993	71.9	False	False	

3	1970	150.437	72.0	False	False
4	1970	326.961	70.9	False	False
..
269	2020	6938.983	81.1	False	False
270	2020	5468.418	82.3	False	True
271	2020	5018.700	80.4	False	False
272	2020	4665.641	84.7	False	False
273	2020	11859.179	77.0	False	False

	Country_Germany	Country_Great Britain	Country_Japan	Country_USA
0	True	False	False	False
1	False	False	False	False
2	False	True	False	False
3	False	False	True	False
4	False	False	False	True
..
269	True	False	False	False
270	False	False	False	False
271	False	True	False	False
272	False	False	True	False
273	False	False	False	True

[274 rows x 9 columns]

```
[27]: X = healthexp.drop(['Life_Expectancy'], axis=1)
y= healthexp['Life_Expectancy']

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,
random_state=54)
rfr = GradientBoostingRegressor(random_state=34)
rfr.fit(X_train, y_train)
```

```
[27]: GradientBoostingRegressor(random_state=34)
```

```
[28]: y_pred = rfr.predict(X_test)
y_pred
```

```
[28]: array([78.39157822, 79.56433878, 77.29122733, 79.10343937, 80.96445194,
77.29122733, 81.03357969, 81.0478605 , 73.89333767, 77.36302545,
81.88363916, 77.31394029, 71.47882512, 74.82462518, 82.95075741,
78.07791025, 71.37073548, 81.34544275, 73.48856096, 80.50574084,
75.02644369, 80.19365241, 73.22967666, 78.54255466, 77.87168267,
76.14733657, 73.14057728, 82.16757383, 74.42529019, 76.12931714,
76.30302906, 81.41008928, 76.11675626, 78.73694658, 78.5494218 ,
82.64753618, 81.09723669, 76.11675626, 77.89083117, 76.6816684 ,
79.75257164, 84.16434938, 74.41614983, 78.72096389, 82.16529143,
84.03429896, 79.14171978, 80.44485511, 81.71684243, 81.2197126 ,
```

```
80.93125484, 79.89895388, 78.76478282, 78.28311191, 79.3784097 ])
```

```
[29]: mean_absolute_error(y_test, y_pred)
```

```
[29]: 0.2709170644443816
```

```
[30]: root_root_mean_squared_error(y_test, y_pred)
```

```
[30]: 0.355971309502974
```

```
[31]: r2_score(y_test, y_pred)
```

```
[31]: 0.9866397423796623
```

```
[32]: def objective(trial):
    n_estimators = trial.suggest_int('n_estimators', 50, 300)
    learning_rate = trial.suggest_float('learning_rate', 0.01, 0.3)
    max_depth = trial.suggest_int('max_depth', 2, 10)
    subsample = trial.suggest_float('subsample', 0.5, 1.0)

    # Define the model
    model = GradientBoostingRegressor(
        n_estimators=n_estimators,
        learning_rate=learning_rate,
        max_depth=max_depth,
        subsample=subsample,
        random_state=42
    )

    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    mse = root_root_mean_squared_error(y_test, y_pred)

    return mse

study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=100)

print(f"Best hyperparameters: {study.best_params}")
```

```
[I 2024-10-20 11:55:27,700] A new study created in memory with name: no-
name-96fbbaf2-3f29-44b6-9cc6-00f7be9dc2b3
```

```
[I 2024-10-20 11:55:27,950] Trial 0 finished with value: 0.32519704881303185 and
parameters: {'n_estimators': 297, 'learning_rate': 0.267593117461795,
'max_depth': 9, 'subsample': 0.7706360903802824}. Best is trial 0 with value:
0.32519704881303185.
```

```
[I 2024-10-20 11:55:28,018] Trial 1 finished with value: 0.32933339451026583 and
```

parameters: {'n_estimators': 95, 'learning_rate': 0.2589727113442577, 'max_depth': 5, 'subsample': 0.8863003094336444}. Best is trial 0 with value: 0.32519704881303185.

[I 2024-10-20 11:55:28,338] Trial 2 finished with value: 0.34532580713847005 and parameters: {'n_estimators': 279, 'learning_rate': 0.21427980221020418, 'max_depth': 9, 'subsample': 0.9493276445660934}. Best is trial 0 with value: 0.32519704881303185.

[I 2024-10-20 11:55:28,597] Trial 3 finished with value: 0.337141751498068 and parameters: {'n_estimators': 288, 'learning_rate': 0.13182683300401327, 'max_depth': 8, 'subsample': 0.8868435801052132}. Best is trial 0 with value: 0.32519704881303185.

[I 2024-10-20 11:55:28,797] Trial 4 finished with value: 0.3124102730275322 and parameters: {'n_estimators': 274, 'learning_rate': 0.20260187408469205, 'max_depth': 5, 'subsample': 0.9062239337295626}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:28,974] Trial 5 finished with value: 0.33315273771793247 and parameters: {'n_estimators': 218, 'learning_rate': 0.028634699851232688, 'max_depth': 7, 'subsample': 0.9546279936392901}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:29,152] Trial 6 finished with value: 0.3531645162647396 and parameters: {'n_estimators': 256, 'learning_rate': 0.2844066049483992, 'max_depth': 6, 'subsample': 0.6782931822279332}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:29,264] Trial 7 finished with value: 0.3257614028015756 and parameters: {'n_estimators': 151, 'learning_rate': 0.19935313467618743, 'max_depth': 6, 'subsample': 0.7203121769063907}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:29,363] Trial 8 finished with value: 0.5949002296375471 and parameters: {'n_estimators': 184, 'learning_rate': 0.023014559915282794, 'max_depth': 2, 'subsample': 0.7777546955416367}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:29,414] Trial 9 finished with value: 0.44110445545375904 and parameters: {'n_estimators': 70, 'learning_rate': 0.04678877022641999, 'max_depth': 4, 'subsample': 0.642435911988464}. Best is trial 4 with value: 0.3124102730275322.

[I 2024-10-20 11:55:29,559] Trial 10 finished with value: 0.283432180253704 and parameters: {'n_estimators': 233, 'learning_rate': 0.1213112051937812, 'max_depth': 3, 'subsample': 0.5246775038673241}. Best is trial 10 with value: 0.283432180253704.

[I 2024-10-20 11:55:29,727] Trial 11 finished with value: 0.2733841950906861 and parameters: {'n_estimators': 229, 'learning_rate': 0.11446729346660696, 'max_depth': 3, 'subsample': 0.5082849682777049}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:29,882] Trial 12 finished with value: 0.33738549688910224 and parameters: {'n_estimators': 221, 'learning_rate': 0.1142162775589909, 'max_depth': 2, 'subsample': 0.512223578308002}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,016] Trial 13 finished with value: 0.2821326667996498 and

parameters: {'n_estimators': 225, 'learning_rate': 0.08900454928810904, 'max_depth': 3, 'subsample': 0.5104345068793334}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,118] Trial 14 finished with value: 0.2936606123578851 and parameters: {'n_estimators': 153, 'learning_rate': 0.07506187156406356, 'max_depth': 4, 'subsample': 0.582380490793184}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,236] Trial 15 finished with value: 0.2971268863998212 and parameters: {'n_estimators': 186, 'learning_rate': 0.08185626842879058, 'max_depth': 3, 'subsample': 0.5830011523959723}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,386] Trial 16 finished with value: 0.2784443164517285 and parameters: {'n_estimators': 244, 'learning_rate': 0.16865472566121278, 'max_depth': 3, 'subsample': 0.5847662153259866}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,549] Trial 17 finished with value: 0.32945547373396866 and parameters: {'n_estimators': 251, 'learning_rate': 0.16414146586098646, 'max_depth': 4, 'subsample': 0.5904421244601817}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,662] Trial 18 finished with value: 0.3243825595212696 and parameters: {'n_estimators': 195, 'learning_rate': 0.15154369262583414, 'max_depth': 2, 'subsample': 0.6578053194559447}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,797] Trial 19 finished with value: 0.352405750294523 and parameters: {'n_estimators': 124, 'learning_rate': 0.17254249831711935, 'max_depth': 10, 'subsample': 0.5615022984441036}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:30,950] Trial 20 finished with value: 0.36472528222929507 and parameters: {'n_estimators': 206, 'learning_rate': 0.22891639396218252, 'max_depth': 5, 'subsample': 0.8154315863946258}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,100] Trial 21 finished with value: 0.3073020016778538 and parameters: {'n_estimators': 250, 'learning_rate': 0.08462566454055234, 'max_depth': 3, 'subsample': 0.5022795336584275}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,240] Trial 22 finished with value: 0.28500727995486796 and parameters: {'n_estimators': 238, 'learning_rate': 0.09996116056277207, 'max_depth': 3, 'subsample': 0.5454805945154205}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,359] Trial 23 finished with value: 0.3241893163539905 and parameters: {'n_estimators': 162, 'learning_rate': 0.06051969627768748, 'max_depth': 4, 'subsample': 0.6320900822024877}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,508] Trial 24 finished with value: 0.3340743120889516 and parameters: {'n_estimators': 269, 'learning_rate': 0.14404392207024466, 'max_depth': 2, 'subsample': 0.6182634714905018}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,646] Trial 25 finished with value: 0.28898722310627706

and parameters: {'n_estimators': 215, 'learning_rate': 0.18325206586845688, 'max_depth': 3, 'subsample': 0.6968492329500976}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,805] Trial 26 finished with value: 0.2859771349934609 and parameters: {'n_estimators': 233, 'learning_rate': 0.11248707230834971, 'max_depth': 5, 'subsample': 0.5452022752235064}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:31,942] Trial 27 finished with value: 0.31601034658811294 and parameters: {'n_estimators': 199, 'learning_rate': 0.09525980881503929, 'max_depth': 4, 'subsample': 0.6028687123419023}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:32,049] Trial 28 finished with value: 0.3411566621028202 and parameters: {'n_estimators': 170, 'learning_rate': 0.1377780919990122, 'max_depth': 2, 'subsample': 0.5441643781159113}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:32,266] Trial 29 finished with value: 0.3058034089127453 and parameters: {'n_estimators': 293, 'learning_rate': 0.06232785721282394, 'max_depth': 7, 'subsample': 0.5007527369849425}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:32,354] Trial 30 finished with value: 0.34296501275487 and parameters: {'n_estimators': 127, 'learning_rate': 0.239036660636862, 'max_depth': 3, 'subsample': 0.750523188775213}. Best is trial 11 with value: 0.2733841950906861.

[I 2024-10-20 11:55:32,498] Trial 31 finished with value: 0.2687702975135457 and parameters: {'n_estimators': 233, 'learning_rate': 0.12255598931314905, 'max_depth': 3, 'subsample': 0.5309087953054625}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:32,650] Trial 32 finished with value: 0.2911019907959925 and parameters: {'n_estimators': 255, 'learning_rate': 0.10341750383000117, 'max_depth': 3, 'subsample': 0.5610610466342697}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:32,813] Trial 33 finished with value: 0.2939611647216501 and parameters: {'n_estimators': 235, 'learning_rate': 0.16223175960762215, 'max_depth': 4, 'subsample': 0.999010386111054}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:32,964] Trial 34 finished with value: 0.344139250148219 and parameters: {'n_estimators': 269, 'learning_rate': 0.13409401544100277, 'max_depth': 2, 'subsample': 0.521813772821312}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:33,109] Trial 35 finished with value: 0.3082259462703164 and parameters: {'n_estimators': 212, 'learning_rate': 0.04311457745659346, 'max_depth': 5, 'subsample': 0.5707021871807052}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:33,287] Trial 36 finished with value: 0.2845700970493313 and parameters: {'n_estimators': 224, 'learning_rate': 0.1307828538799112, 'max_depth': 3, 'subsample': 0.6103923191274128}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:33,487] Trial 37 finished with value: 0.31574363020947704

and parameters: {'n_estimators': 284, 'learning_rate': 0.18707784889203227, 'max_depth': 5, 'subsample': 0.5353396695041222}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:33,724] Trial 38 finished with value: 0.35073747410206124 and parameters: {'n_estimators': 264, 'learning_rate': 0.15184723264528946, 'max_depth': 9, 'subsample': 0.8315364287992428}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:33,877] Trial 39 finished with value: 0.3071835616924356 and parameters: {'n_estimators': 238, 'learning_rate': 0.12409019896753956, 'max_depth': 4, 'subsample': 0.6574995767812493}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,029] Trial 40 finished with value: 0.4677086522347409 and parameters: {'n_estimators': 194, 'learning_rate': 0.012764854978969356, 'max_depth': 6, 'subsample': 0.7053418368479953}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,173] Trial 41 finished with value: 0.28897035646844416 and parameters: {'n_estimators': 243, 'learning_rate': 0.1181888403545856, 'max_depth': 3, 'subsample': 0.5221660145614726}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,301] Trial 42 finished with value: 0.32736784804216373 and parameters: {'n_estimators': 222, 'learning_rate': 0.10679544639401044, 'max_depth': 2, 'subsample': 0.5305061638806614}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,463] Trial 43 finished with value: 0.29762694215456287 and parameters: {'n_estimators': 279, 'learning_rate': 0.06518549601844656, 'max_depth': 3, 'subsample': 0.5031059382923468}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,610] Trial 44 finished with value: 0.26878256364376213 and parameters: {'n_estimators': 228, 'learning_rate': 0.09139112662408969, 'max_depth': 4, 'subsample': 0.5477457735008036}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,751] Trial 45 finished with value: 0.29143930744339114 and parameters: {'n_estimators': 205, 'learning_rate': 0.08965053455365322, 'max_depth': 4, 'subsample': 0.5685473731521724}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,885] Trial 46 finished with value: 0.39313205257246503 and parameters: {'n_estimators': 225, 'learning_rate': 0.046345326038113074, 'max_depth': 2, 'subsample': 0.5938485089230439}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:34,947] Trial 47 finished with value: 0.3396645631610075 and parameters: {'n_estimators': 56, 'learning_rate': 0.0731028791459466, 'max_depth': 6, 'subsample': 0.5526624348425062}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,071] Trial 48 finished with value: 0.312967918752511 and parameters: {'n_estimators': 178, 'learning_rate': 0.20356881669983035, 'max_depth': 4, 'subsample': 0.6369500029479468}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,231] Trial 49 finished with value: 0.2833206390868708 and

parameters: {'n_estimators': 262, 'learning_rate': 0.14482048416305823, 'max_depth': 3, 'subsample': 0.5772907570360203}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,394] Trial 50 finished with value: 0.3254244477042977 and parameters: {'n_estimators': 244, 'learning_rate': 0.29049585313894283, 'max_depth': 5, 'subsample': 0.5184295143589501}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,554] Trial 51 finished with value: 0.28673224445533974 and parameters: {'n_estimators': 262, 'learning_rate': 0.17198550686886205, 'max_depth': 3, 'subsample': 0.5617611851623718}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,709] Trial 52 finished with value: 0.2910215237132724 and parameters: {'n_estimators': 256, 'learning_rate': 0.14777145647614381, 'max_depth': 3, 'subsample': 0.5884989110350857}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:35,841] Trial 53 finished with value: 0.33663834445763846 and parameters: {'n_estimators': 227, 'learning_rate': 0.12351040479543017, 'max_depth': 2, 'subsample': 0.5352066137495815}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,044] Trial 54 finished with value: 0.29805083664401705 and parameters: {'n_estimators': 297, 'learning_rate': 0.09354034571522202, 'max_depth': 4, 'subsample': 0.6228478739475506}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,176] Trial 55 finished with value: 0.3186550196558229 and parameters: {'n_estimators': 212, 'learning_rate': 0.16147731103371912, 'max_depth': 3, 'subsample': 0.5820879890308186}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,323] Trial 56 finished with value: 0.33128018709998036 and parameters: {'n_estimators': 247, 'learning_rate': 0.14010660706178468, 'max_depth': 2, 'subsample': 0.6544211258426351}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,537] Trial 57 finished with value: 0.28831754387250763 and parameters: {'n_estimators': 275, 'learning_rate': 0.07626085785849052, 'max_depth': 4, 'subsample': 0.6799987666663414}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,702] Trial 58 finished with value: 0.320019796192551 and parameters: {'n_estimators': 190, 'learning_rate': 0.10540528737668065, 'max_depth': 8, 'subsample': 0.5130929258940896}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:36,884] Trial 59 finished with value: 0.3431264075744974 and parameters: {'n_estimators': 260, 'learning_rate': 0.1846341094850861, 'max_depth': 3, 'subsample': 0.6059995367785499}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,009] Trial 60 finished with value: 0.36218200880182144 and parameters: {'n_estimators': 206, 'learning_rate': 0.2706652855831479, 'max_depth': 2, 'subsample': 0.5489548662054436}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,151] Trial 61 finished with value: 0.27663318682392724

and parameters: {'n_estimators': 232, 'learning_rate': 0.11744139207969508, 'max_depth': 3, 'subsample': 0.5258295173126227}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,291] Trial 62 finished with value: 0.282327356204379 and parameters: {'n_estimators': 232, 'learning_rate': 0.11441097886689146, 'max_depth': 3, 'subsample': 0.5802111287236595}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,434] Trial 63 finished with value: 0.2846683813869978 and parameters: {'n_estimators': 230, 'learning_rate': 0.10964627012259157, 'max_depth': 3, 'subsample': 0.5343856984399099}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,581] Trial 64 finished with value: 0.29354989451107105 and parameters: {'n_estimators': 216, 'learning_rate': 0.12765983353767738, 'max_depth': 4, 'subsample': 0.5021350240116306}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,720] Trial 65 finished with value: 0.35948926495929134 and parameters: {'n_estimators': 240, 'learning_rate': 0.08687325387700343, 'max_depth': 2, 'subsample': 0.5508074120453469}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:37,876] Trial 66 finished with value: 0.2980607734880508 and parameters: {'n_estimators': 249, 'learning_rate': 0.0971489217202378, 'max_depth': 3, 'subsample': 0.5213821830496556}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,027] Trial 67 finished with value: 0.2757023560730843 and parameters: {'n_estimators': 199, 'learning_rate': 0.11547805321778493, 'max_depth': 4, 'subsample': 0.872277198440613}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,166] Trial 68 finished with value: 0.3212843996999328 and parameters: {'n_estimators': 181, 'learning_rate': 0.08147559495991766, 'max_depth': 5, 'subsample': 0.9404810225597366}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,339] Trial 69 finished with value: 0.2998468920089925 and parameters: {'n_estimators': 199, 'learning_rate': 0.13222337419778468, 'max_depth': 4, 'subsample': 0.8114060461385173}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,511] Trial 70 finished with value: 0.2936321850898209 and parameters: {'n_estimators': 217, 'learning_rate': 0.09959855855692458, 'max_depth': 4, 'subsample': 0.8906174319977933}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,663] Trial 71 finished with value: 0.3031109360473692 and parameters: {'n_estimators': 229, 'learning_rate': 0.11676696430986833, 'max_depth': 3, 'subsample': 0.8394740356203501}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,812] Trial 72 finished with value: 0.30506979109723376 and parameters: {'n_estimators': 235, 'learning_rate': 0.11429279951338023, 'max_depth': 3, 'subsample': 0.7812357657877741}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:38,955] Trial 73 finished with value: 0.299335059932847 and

parameters: {'n_estimators': 204, 'learning_rate': 0.06796442040304482, 'max_depth': 4, 'subsample': 0.5623236135973059}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,103] Trial 74 finished with value: 0.27726456031107777 and parameters: {'n_estimators': 221, 'learning_rate': 0.12350693793514392, 'max_depth': 3, 'subsample': 0.5372381389445163}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,259] Trial 75 finished with value: 0.3197026616693903 and parameters: {'n_estimators': 168, 'learning_rate': 0.15526663355775777, 'max_depth': 7, 'subsample': 0.9211213691341871}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,393] Trial 76 finished with value: 0.3610721016867259 and parameters: {'n_estimators': 221, 'learning_rate': 0.05720600204864207, 'max_depth': 2, 'subsample': 0.8570779651846874}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,529] Trial 77 finished with value: 0.3017856322375178 and parameters: {'n_estimators': 212, 'learning_rate': 0.1742764965575541, 'max_depth': 3, 'subsample': 0.5374811319067782}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,603] Trial 78 finished with value: 0.28771354036692975 and parameters: {'n_estimators': 89, 'learning_rate': 0.1369162710195497, 'max_depth': 4, 'subsample': 0.5089338457782696}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,758] Trial 79 finished with value: 0.36001397984272665 and parameters: {'n_estimators': 251, 'learning_rate': 0.12366063549935136, 'max_depth': 2, 'subsample': 0.7398744705249802}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:39,909] Trial 80 finished with value: 0.28039318485554693 and parameters: {'n_estimators': 241, 'learning_rate': 0.09222678665030536, 'max_depth': 3, 'subsample': 0.5174340187345003}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,055] Trial 81 finished with value: 0.2949761769037094 and parameters: {'n_estimators': 243, 'learning_rate': 0.08205921520156348, 'max_depth': 3, 'subsample': 0.520242399964311}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,195] Trial 82 finished with value: 0.28748086601963063 and parameters: {'n_estimators': 221, 'learning_rate': 0.10383294564680376, 'max_depth': 3, 'subsample': 0.5508761620389389}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,380] Trial 83 finished with value: 0.28988642449517504 and parameters: {'n_estimators': 238, 'learning_rate': 0.09184546850695231, 'max_depth': 4, 'subsample': 0.5319059506207203}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,510] Trial 84 finished with value: 0.330846000770609 and parameters: {'n_estimators': 197, 'learning_rate': 0.11003247880130926, 'max_depth': 3, 'subsample': 0.9959906722867444}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,638] Trial 85 finished with value: 0.2755188597168557 and

parameters: {'n_estimators': 209, 'learning_rate': 0.12065534816826647, 'max_depth': 3, 'subsample': 0.5137576808894949}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,784] Trial 86 finished with value: 0.3151901181133613 and parameters: {'n_estimators': 208, 'learning_rate': 0.12034314215245714, 'max_depth': 5, 'subsample': 0.5105430357013842}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:40,904] Trial 87 finished with value: 0.3394515609516515 and parameters: {'n_estimators': 186, 'learning_rate': 0.13250523083945367, 'max_depth': 2, 'subsample': 0.5689450129169772}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:41,042] Trial 88 finished with value: 0.279441527189532 and parameters: {'n_estimators': 226, 'learning_rate': 0.15520164978259704, 'max_depth': 3, 'subsample': 0.5388918894107891}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:41,247] Trial 89 finished with value: 0.35259678053966315 and parameters: {'n_estimators': 227, 'learning_rate': 0.15566931509182913, 'max_depth': 10, 'subsample': 0.5940160321261414}. Best is trial 31 with value: 0.2687702975135457.

[I 2024-10-20 11:55:41,392] Trial 90 finished with value: 0.26756036736615935 and parameters: {'n_estimators': 216, 'learning_rate': 0.14322018406580803, 'max_depth': 4, 'subsample': 0.5406516005689751}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:41,535] Trial 91 finished with value: 0.2839898731388246 and parameters: {'n_estimators': 218, 'learning_rate': 0.16538656442887342, 'max_depth': 4, 'subsample': 0.5557724344694196}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:41,669] Trial 92 finished with value: 0.27470824373319225 and parameters: {'n_estimators': 202, 'learning_rate': 0.14311660190415712, 'max_depth': 4, 'subsample': 0.5418480086606893}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:41,803] Trial 93 finished with value: 0.29370744684684014 and parameters: {'n_estimators': 202, 'learning_rate': 0.14440344299472405, 'max_depth': 4, 'subsample': 0.5284729481912918}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:41,950] Trial 94 finished with value: 0.28045739215083354 and parameters: {'n_estimators': 209, 'learning_rate': 0.1275406340731873, 'max_depth': 5, 'subsample': 0.5477745710819973}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:42,075] Trial 95 finished with value: 0.29977071364481417 and parameters: {'n_estimators': 191, 'learning_rate': 0.1382706441772703, 'max_depth': 4, 'subsample': 0.5019503850286139}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:42,225] Trial 96 finished with value: 0.29074368761970637 and parameters: {'n_estimators': 212, 'learning_rate': 0.14685314197335717, 'max_depth': 4, 'subsample': 0.5280517474257275}. Best is trial 90 with value: 0.26756036736615935.

[I 2024-10-20 11:55:42,409] Trial 97 finished with value: 0.319604857634209 and

```

parameters: {'n_estimators': 255, 'learning_rate': 0.11867606754400264,
'max_depth': 5, 'subsample': 0.5733044545999941}. Best is trial 90 with value:
0.26756036736615935.
[I 2024-10-20 11:55:42,536] Trial 98 finished with value: 0.2857077027347594 and
parameters: {'n_estimators': 174, 'learning_rate': 0.12736296874721426,
'max_depth': 4, 'subsample': 0.5602314712383062}. Best is trial 90 with value:
0.26756036736615935.
[I 2024-10-20 11:55:42,691] Trial 99 finished with value: 0.2878060321930935 and
parameters: {'n_estimators': 234, 'learning_rate': 0.17614425422064806,
'max_depth': 3, 'subsample': 0.5402212509921593}. Best is trial 90 with value:
0.26756036736615935.

Best hyperparameters: {'n_estimators': 216, 'learning_rate':
0.14322018406580803, 'max_depth': 4, 'subsample': 0.5406516005689751}

```

```
[33]: study.best_params
```

```
[33]: {'n_estimators': 216,
      'learning_rate': 0.14322018406580803,
      'max_depth': 4,
      'subsample': 0.5406516005689751}
```

```
[35]: optuna.visualization.plot_optimization_history(study)
```

```
[36]: optuna.visualization.plot_parallel_coordinate(study)
```

```
[38]: optuna.visualization.plot_slice(study, params=['n_estimators', 'learning_rate',
↪ 'max_depth', 'subsample'])
```

```
[40]: optuna.visualization.plot_param_importances(study)
```

30-K_Means_Clustering

October 20, 2024

1 K-Means Clustering

K-Means Clustering is one of the simplest and most widely used unsupervised learning algorithms for partitioning a dataset into distinct groups or clusters. In K-Means, each cluster is represented by a centroid, and each data point is assigned to the cluster with the closest centroid.

The key idea behind K-Means is to partition data into k clusters, where k is a predefined number. The algorithm works iteratively to assign data points to one of the k clusters based on the similarity (often measured by Euclidean distance).

1.0.1 Why do we use K-Means Clustering?

K-Means Clustering is highly popular due to:

- **Simplicity:** Easy to understand and implement.
- **Speed:** Efficient for clustering large datasets.
- **Applicability:** Useful for a variety of tasks such as customer segmentation, image compression, and document clustering.
- **Interpretability:** Results are easy to interpret, as each data point belongs to exactly one cluster.

1.0.2 How does K-Means Clustering work?

The algorithm works in the following steps:

Initialization:

Choose the number of clusters k . Randomly initialize k centroids (cluster centers) in the feature space. Assignment (Step 1):

For each data point in the dataset, calculate its distance to each of the k centroids. Assign each data point to the nearest centroid (based on minimum distance). Update Centroids (Step 2):

After the assignment step, recompute the centroids by calculating the mean of all data points assigned to each cluster. The new centroid will be the new center of the cluster. Repeat:

Repeat the Assignment and Update steps iteratively until the centroids do not change significantly, or until a predefined number of iterations is reached. This is called convergence. Output:

The algorithm produces k clusters, each with a centroid and a set of assigned data points.

1.0.3 When to use K-Means Clustering?

K-Means is best suited for problems where:

You know the number of clusters (k): K-Means requires you to specify k in advance. Clusters are spherical and equally sized: K-Means assumes clusters are compact, and all of approximately equal size. There are no significant outliers: K-Means is sensitive to outliers, as they can heavily skew the centroid calculation. Common use cases include:

Customer Segmentation: Group customers with similar buying behaviors. Image Compression: Reduce image size by clustering similar pixels. Document Clustering: Group similar documents based on their content.

```
[1]: from sklearn.cluster import KMeans
import numpy as np
import matplotlib.pyplot as plt

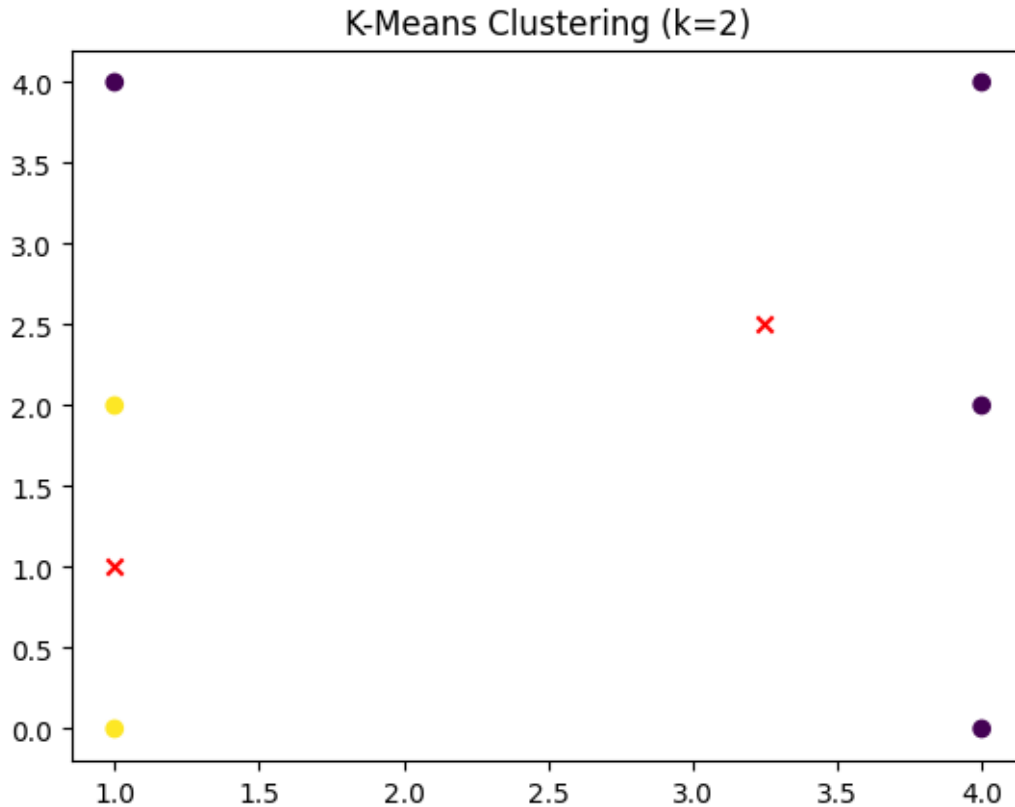
X = np.array([[1, 2], [1, 4], [1, 0],
              [4, 2], [4, 4], [4, 0]])

kmeans = KMeans(n_clusters=2, random_state=0)

kmeans.fit(X)
labels = kmeans.predict(X)

centroids = kmeans.cluster_centers_

# Visualize the clusters
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis')
plt.scatter(centroids[:, 0], centroids[:, 1], c='red', marker='x')
plt.title("K-Means Clustering (k=2)")
plt.show()
```



1.0.4 Advantages of K-Means Clustering:

Scalability: Can handle large datasets efficiently. Easy to implement: Straightforward in both understanding and coding. Fast convergence: The iterative approach usually converges quickly.

1.0.5 Disadvantages of K-Means Clustering:

Choosing k: The number of clusters (k) must be pre-specified, and determining the optimal k can be challenging. Sensitive to initialization: Different random initializations can lead to different clustering results (though you can mitigate this by running the algorithm multiple times or using the k-means++ initialization method). Sensitive to outliers: Outliers can significantly affect cluster centroids. Assumption of spherical clusters: K-Means works best when clusters are of equal size and spherical in shape.

1.0.6 Who uses K-Means Clustering?

- **Marketers:** To segment customers based on purchasing patterns.
- **Biologists:** To classify genes with similar characteristics.
- **Image Analysts:** To compress images or classify regions in an image.
- **Text Miners:** To group similar documents or articles.